

Visual Bias*

Giulia Caprini†

March 1, 2023

Abstract

This paper studies how pictures leading online news pieces influence readers' processing of the news. I document two relevant facts in the US news market. First, the visual language adopted by news outlets is politically partisan and significantly polarized; such visual polarization is comparable in magnitude to the documented verbal polarization of Congress in recent years. For this analysis, I construct a visual vocabulary of graphic features and apply a dictionary method to study the visual language polarization in the leading images published in US news between December 2019 and December 2020. Second, I document that such visual partisanship amounts to an expression of political media bias: in a survey experiment, individuals exposed to identical news pieces but leading pictures with opposite partisanship formulate significantly different opinions, which are slanted towards the images' respective ideological poles. I find that news' *visual bias* causes an increase in the issue polarization of the general public. The slanting effect of images interacts with readers' prior, and audiences on both sides of the political spectrum react more distinctly to pictures aligned with their viewpoint. This pattern implies that the polarizing effect of visual bias is further exacerbated if readers' source their news exclusively from like-minded outlets.

Keywords: Media bias, language polarization, images, visual information processing, news photography
L82 D91 D83

*I am especially grateful to Andrea Mattozzi and Riccardo Salerno for their continued support throughout the project. This work benefited from comments by Zeinab Aboutalebi, Giacomo Calzolari, Flavia Cavallini, Thomas Crossley, David Yanagizawa-Drott, Ruben Durante, Andrea Ichino, Mustafa Kaba, David K. Levine, David Martinez de Lafuente, Ro'ee Levy, Patrizia Mealli, Massimo Morelli, Salvatore Nunnari, Nancy Qian, David Stromberg, Junze Sun, Ekaterina Zhuravskaya, and seminar participants at Alghero (Ibeo), CEPR workshop on Media (EIEF/Tor Vergata), Bocconi (LEAP), Collegio Carlo Alberto, EEA, EUI, Oxford University, University of Milan, Italian Economists Society, and the Royal Economics Society. I gratefully acknowledge financial support from the 2019 EUI Research Council grant, and from the 2020 Early Stage Researchers grant at the European University Institute, where this work was largely developed. The experiment was approved by the IRB ("Ethics Committee") of the European University Institute, and pre-registered in the AEA RCT Registry with Digital Object Identifier: 10.1257/rct.7904-1.0.

†Nuffield College and Department of Economics, Oxford University. Email: Giulia.caprini@economics.ox.ac.uk

I INTRODUCTION

“We don’t see things as they are, we see things as we are”

—Anais Nin (writer), 1961

The starting point of this paper are four facts that today characterize the access to written news: first, people increasingly read their news online, finding news pieces through social media platforms or news apps (Shearer, 2021). Second, news readers first encounter these news pieces through short previews, a format consisting of a headline, a short summary text, and a leading image (Figure I): this format confers leading pictures a prominent position over other news elements. Third, people heavily rely on the content of news previews, for instance when they share the news pieces in their social media feeds without reading the full text (Gabelkov et al., 2016). Fourth, most initiatives tackling online misinformation and fake news are concerned with the analysis of written contents, and pictures largely escape systematic scrutiny. Taken together, these dynamics suggest that leading images gained large strategic importance in the communication strategy of news media, and in particular of ideologically-partisan outlets: not only the unspoken content can reach a broad audience, but the ambiguity of intended meaning in pictures allows to provide controversial hints and cues while limiting the potential backfire.

This study explores the role of leading images in the communication strategy of politically-slanted news producers, investigating the extent of political bias over and above text. The paper documents two complementary instances: on the one hand, news providers with different political leanings use systematically different visual language; on the other hand, partisan visual narratives slant readers’ opinion towards the news outlets’ ideological poles. Jointly, these dynamics prove that the news’ *visual bias* is a tangible expression of political media bias (De Vreese, 2004; Groseclose and Milyo, 2005; McCombs and Reynolds, 2009; DellaVigna and Gentzkow, 2010; Prat and Strömberg, 2013; Strömberg, 2015; Prat, 2018).

The analysis is organized in two Sections. I first study the polarization of the news’ visual language, namely the extent to which the characteristics of the leading images are distinctive of their news outlets’ political leaning. I collect about 300’000 leading images from news published between December 2019 and December 2020 by the main US news outlets, and I exploit computer vision tools to extract information on the images’ content (such as the subjects and objects depicted, their characteristics, and contextual aspects of the image). Drawing from existing studies in photography, semiotics, psychology, and political science, I combine this information into key measures to decode meaning from visual contents. I thus construct a “visual vocabulary”, whose tokens pertain several dimensions relevant to convey political cues through graphic elements. Borrowing from text-analysis methods, I map the images in my dataset to the vector of tokens in my visual vocabulary, using this representation to analyse systematically the portrayal of subjects and compare it across pictures. To measure the partisanship of the news’ visual language I employ the leave-out estimator of phrase partisanship developed by Gentzkow, Shapiro, and Taddy (2019). The results of this first analysis show that leading images chosen by liberal- and conservative- news outlets are systematically different. The news’ visual language is significantly partisan and distinctive of the sources’ political leaning: as a term of comparison, the visual polarization appears close to the verbal polarization that Gentzkow, Shapiro, and Taddy (2019) document in Congress sessions of

recent years. Additionally, I analyse the text of the news summaries to categorize the news pieces into seven relevant topics; I thus re-apply the polarization measure to news within each topic. The measure indicates high visual polarization in news broadly pertaining the following topics: *Politics*, *Security*, *Health & Covid-19* (polarized only from April 2020 onward), *Entertainment* and, to a lesser extent, *Society*.

The second Section of the paper presents a survey experiment conducted on a nationally-representative sample of 2'000 US residents to examine the effect of visual partisanship on news readers' opinion. I test two hypotheses: whether partisan leading images distinctive of Republican/Democrat outlets effectively slant the audience towards their respective party, and whether partisan images increase the polarization of public opinion. The results indicate that individuals exposed to identical news previews but leading pictures with opposite partisan loadings formulate significantly different opinions, with the slant following that of the news outlet. Exploring the factors underlying this impact, I find that readers who ex-ante hold more extreme opinions on a given issue appear more receptive to leading images compared to respondents who hold more intermediate opinions. Instead, neither the perceived issue salience nor readers' prior knowledge appear to be strong predictors of the susceptibility to leading images. Finally, I find that news' visual bias causes a significant increase of issue polarization in the general public: the slanting effect of images interacts with readers' prior, and audiences on both sides of the political spectrum react more distinctly to pictures aligned with their viewpoint. This pattern implies that the polarizing effect of visual bias is further exacerbated if readers' source their news exclusively from like-minded outlets.

This study seeks to contribute to the understanding of media bias. From a methodological viewpoint, I borrow a NLP framework originally developed for text analysis and I apply it to the study of images, designing a data-driven visual vocabulary for the systematic interpretation of pictures. This approach allows the joint study of a large number of characteristics and, as done in text analysis, to explore lexical semantics and syntactic relations among (pictorial) elements. In this respect, the paper relates to a number of works that explore the role of images in political communication by examining the relative incidence of distinct graphic characteristics across sources with opposite political stances (e.g. Peng, 2018; Boxell, 2021; Haim and Jungblut, 2021). I rely on these studies to inform the construction of the visual vocabulary, so to encompass the many graphic aspects relevant for political visual framing. This study also relates to numerous analyses performing automatic bias detection in text through natural language processing methods (see, e.g. Greene and Resnik, 2009; Yano, Resnik, and Smith, 2010; Recasens, Danescu-Niculescu-Mizil, and Jurafsky, 2013; Gentzkow, Shapiro, and Taddy, 2019). Among those, my analysis of visual polarization draws in particular from Demszky et al. (2019), who use the measure of phrase partisanship originally developed by Gentzkow, Shapiro, and Taddy (2019) to study the political polarization in the text of tweets related to mass shootings events in the US.

A second contribution of this paper consists of presenting novel evidence on the impact of leading images on news readers' opinion. I show that visual partisanship indeed constitutes a form of political bias of news media, and I document a significant causal effect of visual bias on public opinion polarization. In this respect, the paper relates to several works that identified the correlation between the increasing polarization of media and the general population's political stance, underscoring the imperative to accurately detect news bias and to understand its nature (e.g. Gentzkow and Shapiro, 2010, 2011; Prior, 2013).



FIGURE (I)
News Previews on Social Media

Notes: The Figure illustrates two examples of news previews in different social media websites. The preview's main elements, marked with letters to the sides, are: the name of the news source (A), the news' leading text (B), the news' leading image (C), and the news' header (D). Leading images occupy the largest share of the previews' area, dominating other elements in terms of visibility. Photo by Brooks Kraft for Getty Images (“The two-story Board Room in the Eccles Building, Washington, DC”). Image registered and available at shorturl.at/hnuAC.

The remained of the paper is organized as follows: I first quantify visual partisanship in US news (Section II). Then, I test the effect of partisan images on general public opinion (Section III). I close with a summary of the findings and a general discussion of their implications (Section IV).

II QUANTIFYING VISUAL PARTISANSHIP IN US NEWS

In this Section (II) I quantify the extent of visual partisanship in US news between December 2019 and December 2020. I find a high degree of polarization across the visual narratives adopted by news sources across the political spectrum. To estimate visual partisanship I begin by applying a dictionary method, which entails creating a visual vocabulary and expressing images as vectors of dictionary entries; I then use a partisanship estimator borrowed from text-analysis to measure language distance.

The Section is organized as follows: I introduce the methodology for the analysis (*II.A*), illustrate the construction of the visual vocabulary (*II.B*), and estimate the news' visual partisanship overall and within different news topics (*II.C*). I close by discussing the perks and pitfalls of the approach (*II.D*).

II.A Method: A dictionary-based approach to the study of pictures

To perform a comprehensive analysis of the leading pictures collected from US news I adapt the dictionary-based methodology originally developed to study texts. This approach transforms the pictures in a convenient format, allowing me to then exploit the existing measures of language distance from text analysis (such as the partisanship estimator illustrated in Subsection *II.C*).

Dictionary methods entail counting words from a predefined lexicon (the dictionary) in a big corpus, with

the intent to explore or test hypotheses about the corpus itself. In practice, the essence of the method consists of transforming a document in a vector of counts or indicators for the presence of given language elements. The reference vocabularies are generally composed of *unigrams*, *bigrams*, and/or *trigrams*, namely series of one or two/three consecutive words (or word roots) that, once combined (and before the removal of stopwords and word suffixes/prefixes) compose the phrases of a text; these elements are commonly referred to as *tokens*.

I adapt this procedure to study the news' visual language and to extract computationally the meaning of the large number of leading images in my dataset. By processing the raw information described in previous section, I draft a vocabulary of graphic and content-related *features* which, once combined, result in the pictures' backbone. Following the parallel with text analysis, these can be considered as my set of "visual tokens". Describing dictionary methods for text analysis, Gentzkow, Kelly, and Taddy (2019) illustrate the simplifications that help reduce raw text to a representation suitable for statistical analysis:

"We typically make three kinds of simplifications: dividing the text into individual documents, reducing the number of language elements we consider, and limiting the extent to which we encode dependence among elements within documents. The result is a mapping from raw text Δ to a numerical array C . A row c_i of C is a numerical vector with each element indicating the presence or count of a particular language token in document i ."

I reduce the pictures in my set to simpler representations through three steps. The first entails dividing the corpus into single documents; in my application, since the attributes of interest are defined at the single image level, I consider each picture as an individual document. The second step entails adapting the number of language elements that are considered. The purpose of my analysis is to study how the visual narrative differs among sources with different political leanings. To this extent, I consider both general graphic elements and politically-relevant cues in the pictures, as described in details in the next Section.

The third step entails encoding the dependence among elements within a document. In text analysis, this is aided by including consecutive words/stems (bigrams and trigrams) in the vocabulary. The study of consecutive words helps the extraction of meaning from a text because words' contiguity proxies pertinence to the same textual object. In images, instead, I model the pertinence of multiple characteristics to the same portrayed object using combinations of co-occurring features of each depicted element (this "features engineering" is described in Subsection II.B.3). The unigrams, bigrams and trigrams in my visual vocabulary are thus represented by single graphic features, feature-pairs, and features-triplets. To create meaningful combinations I exploit the features' semantic categories, organizing them in a syntax as in verbal language. To this extent, I extract image features passing the pictures to the computer vision algorithms from the Azure suite by Microsoft; for each image, these algorithms provide me with an array of *words* describing each detected elements, as well as the pixel-coordinates of the "bounding boxes" the element falls in. Hence, this type of features extraction not only detects the features' presence, but also directly provides annotated meanings.¹ This is not only crucial to create the combinations mentioned above, but also implies that all entries of the vocabulary are interpretable by design.

¹In this way, the cognition (and re-cognition") phase of image reading coincides with features extraction. Conversely, other methods for mapping pictures to vectors of features express visual tokens as patterns of pixels, which may or may not mark an identifiable object and whose meaning is not annotated (as, for instance, the "Bag of Visual Words" method).

II.B Creating a Visual Vocabulary

This Subsection describes the three phases of the vocabulary creation: the retrieval of pictures (*II.B.1*), the features extraction (*II.B.2*), and the features engineering (*II.B.3* and *II.B.4*).

II.B.1 Retrieving Pictures

News sources. I begin by constructing a comprehensive list of the relevant news outlets from a list of the top 50 US news media by digital circulation from Similarweb.com. The circulation metric is based on the number of Unique Visitors per Month (UVM), and it indicates how many people in the U.S. market visit a website in a month.² I discard sources that do not cover political news and are exclusively focused on entertainment, celebrity news, fashion, beauty news, or local news.³ I derive the sources' party affiliation through the political bias ratings from Adfontesmedia.com and Allsides.com, keeping the sources with concordant partisanship attribution.⁴ The final sample consists of 22 sources evenly divided on the two sides of the political spectrum. Table A.1.1 presents a list of the sources as well as their partisanship scores as documented by the above mentioned sources.

News data. From the Twitter accounts of the selected sources, I obtain all the news articles shared on the social media between December 1, 2019 and Dec 13, 2020. I focus exclusively on tweets sharing written news pieces (discarding links to video, voice recordings etc.), and I filter out all news pieces written by an outlet but tweeted by a different source. As sources commonly share their pieces multiple times to maximize audience, I keep only the latest version of each piece. The resulting dataset counts 298'850 news pieces.

From the articles' metadata I retrieve and store the headline, description, publication outlet, publication date, and leading image. I define an article's leading image as the main picture accompanying a news piece: loosely speaking, those correspond to the pictures displayed in the news summary when articles are shared on social media.⁵

II.B.2 Features Extraction

As mentioned in Subsection *II.A*, the visual vocabulary in this paper builds on image features expressed in terms of their annotated meanings. This grants that features can later be combined to decode contents of higher semantic level, and that all vocabulary entries are interpretable by design. I pass each picture to image analysis algorithms from the computer vision API by Microsoft (Azure cognitive services).⁶ The following paragraphs illustrate the algorithms used and their output.

Image analysis, Face detection, Face verification, and emotion recognition. The image analysis algorithm detects the presence of faces and assigns tags to the picture based on the depiction of "iconic", recognizable items (e.g. clothing pieces, natural elements, animals, etc.). I pass images that contain at least one human face to the face detection, description (age, gender, hair colour, eye-nose-mouth landmarks etc) and emotion recognition algorithms. The latter classifies the emotions expressed by a face into happiness, sadness, anger, fear, contempt, disgust and surprise. Using the subset of images containing a human face, I check

²UVM data by SimilarWeb.com accessed on October 27, 2020. See <https://www.similarweb.com>.

³The labels correspond to the tags in the descriptions by Similarweb.com; discarded outlets are mainly small local outlets, with the notable exception of the "Los Angeles Times", the "Chicago Tribune" and the "Arizona Republic".

⁴<https://www.adfontesmedia.com>, <https://www.allsides.com>

⁵See Figure I for an illustration of the news previews on social media.

⁶For a list, see <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision>

whether the depicted persons are members of the US congress or prominent figures of the US recent public scene. To this purpose, I first train the face-verification algorithm on a comprehensive set of images created by manually selecting 9 pictures of each congressmen and congresswomen sitting in the 114, 115, 116, and 117th US congresses.⁷ Then, to record the presence of prominent public figures outside the setting of Congress (e.g. Governors, Supreme Court judges, athletes, actors etc.) I pass the pictures to Microsoft’s “celebrities” API, a face-verification algorithm pre-trained to recognize a wide set of celebrities. In addition, whenever the picture contains a congressperson, I record her relative political leaning as measured through the first dimension of the Common Space DW-NOMINATE score from McCarty, Poole, and Rosenthal (2015).⁸ For each of the above-mentioned extracted elements, a confidence score is returned along with the bounding box coordinates of each detected element; this allows to determine the element’s position within an image.

Finally, the Image Analysis algorithm returns general information on each image. For instance, it identifies and categorizes the pictures using a category taxonomy with parent/child hierarchies (e.g. “indoor_marketstore”); it describes the “type” of image, such as whether it is a drawing or clip art, whether an image is black and white or color and, for color images, dominant and accent colors. The algorithm also produces a list of image tags from a set of thousands of recognizable objects, living things, scenery, and actions. These tags are indicators for noteworthy contextual elements of a picture, such as natural elements (e.g. fire, water, etc.), transportation means (e.g. cars, ambulances, etc.), architectural elements (e.g. skyscrapers, castles, etc.), or text content (e.g. banners, signals, etc.).⁹ Importantly, in the case of general image information, the API does not return bounding box coordinates. Hence, these elements can be used as general image descriptors but do not possess information on location (this implies that they can hardly be “attributed” to other elements to characterise their depiction).

Overall, the final information set extracted from each image through the computer vision suite includes the following: detection of people and recognition of politicians and celebrities; details on people position (coordinates), head poses (pitch, yaw, roll), facial expressions, position of landmarks (nose, eyes, mouth, etc); detection and recognition of objects, details on colors; detection of places and background elements through tags; details on the image category, type, color scheme, tags.

In the rest of the paper, I refer to this set as to the “raw information” on leading pictures given by the algorithms.

II.B.3 Engineering: Single Features

This subsection describes how the “raw information” is then processed to obtain a vocabulary of valid tokens to decode visual language. I do so in two steps: first, I derive meaningful individual features (for instance, using the coordinates of a face to derive its size), organizing them in syntactic classes; second, I combine individual features in couples and triplets. Jointly, single features and combinations compose the visual vocabulary.

Features Class “Subjects” (S):

“Subject features” are attributes capturing characteristics that are constant for a subject across all pictures

⁷When a person’s portraits did not cover a wide range of angles, I added a 10th picture to her set. Portraits were chosen so to include different camera angles for each person.

⁸I attribute Donald Trump (who didn’t seat in congress before the presidency) the same DW-NOMINATE score as the most partisan Republican congressperson (Tommy Tuberville, with score 0.916). I attribute Joe Biden the same partisanship he had as congressman in 2008 (-0.314).

⁹For a complete list of the tags classes available, see <https://docs.microsoft.com/en-us/azure/cognitive-services/computer-vision/category-taxonomy>.

in the sample (such as a given person’s name or political party affiliation). This syntax class encompasses indicators for whether or not a person is well-known to the public (a “celebrity” status), whether the depicted person is a man or a woman, the subjects’ names, and the subjects’ relative position in the political partisanship distribution (measured through the first dimension of the Common Space DW-NOMINATE score from Poole and Rosenthal, 1985).¹⁰ This partisanship score originally ranges from -1 to +1 and is scaled so that the lowest scores are those of liberal Democrats and the highest scores are those of conservative Republicans; dividing the index domain into 10 equally spaced bins, I produce four vocabulary indicators to mark pro-Democratic partisanship, four to mark the relative pro-Republican leaning of a congress member, and one indicator to mark the central bin. The unique identification of people within an image is essential to correctly attribute the adjectives features through features’ combinations. For this reason, the “Subjects” class also includes within-picture unique identifiers for all the persons portrayed. Those are indicators for their saliency rank within a picture, obtained from the weighted average of their face area share (70%) and centrality in the picture (30%) (the higher the rank, the more salient the Subject).¹¹ Appendix Section A.1.2 provides a summary of all Subject subclasses.

Features Class “Adjectives” (A):

With the term “Adjective features”, I refer to other features defined at face-level, like Subjects. However, Adjectives indicate variable attributes that subjects may possess in given pictures and not in others. I organize Adjective features by their pertinence to three dimensions: *Size*, *Centrality*, and *Kinesics*. The first two pertain the pictures’ proxemics, i.e. the way the space is used in the portrayal, while the third concerns the dynamics and gestures portrayed.

-Size. Subjects’ size is relevant to image analysis primarily because higher graphic dimension induces higher visibility. As humans do not receive a picture’s content through a single glance but rather through separate scans, the longer a person looks at a picture the higher the chances of marginal details to be integrated in such mental map and thus be “seen”¹². As a consequence, bigger objects are more likely to be grasped by a viewer. In this sense, we can interpret the relative size of depicted objects as informative of the illustrative intent behind the choice of a picture: if an element occupies a large portion of the image, the person who chose the illustration meant to highlight that element to the viewers. Therefore, objects’ size proxies a criterion of precedence among the objects portrayed in the picture. The visual vocabulary includes three individual features for subjects’ size: a “close-up” indicator for faces whose area is equal or greater than 1/6 of the total image area, a “mid range” indicator, for faces from 1/6 to 1/24 of the total image area, and a “long shot” indicator for faces with size below 1/24 of the total image area.

-Centrality. I define an object’s *centrality* in a picture as its ability to attract the viewer’s attention. It is measured in terms of proximity to the two vertical and two horizontal parallel lines that divide a picture in three equal sections, vertically and horizontally, following the “rule of thirds” (illustrated in Figure II). Such “attention lines” have been shown to attract and guide viewers’ attention within a picture (see, e.g. Koliska and

¹⁰Data from voterview.com, by Lewis et al., 2021

¹¹The maximum number of individuals I contemplate in the data is 10 persons in the same image.

¹²Human vision is sharp only in a small central area of the visual field (the fovea), while on the retina of the eye acuity falls off rapidly from this area. Because detailed discriminations are possible only on the fovea, eyes need to scan pictures to take in all the details. Such a scanning does not occur in smooth sweeps but, rather, as a series of very rapid jumps (“saccades”) and stops (“fixations”).

Oh, 2021) and are often marked in cameras' viewfinders to aid photographers' frame choice.¹³ For every face in a picture, its centrality is inversely proportional to the distance between the eyes-midpoint and the closest of 5 focal points (either the center of the image, or one of the four intersections of the attention lines determined through the rule of thirds). Formally, it is expressed as:

$$c^{ROTC}(x, y) = \text{argmax}_i e^{-\left(\sqrt{\left(\frac{x-x_i}{W}\right)^2 + \left(\frac{y-y_i}{H}\right)^2}\right)} \quad (1)$$

where i indicates the focal point, x and y are the coordinates of the eyes' midpoint, x_i and y_i are the coordinates of point i , and W and H express the total width and height of the image. The distances in (1) are measured in pixels, with the top-left angle of the images always marking the origin of coordinates axes. The distance between focal point i and the eyes-midpoint is normalized with respect to the image dimensions to ensure cross-pictures comparability. The centrality index therefore ranges in 0-1, with higher values indicating higher proximity to a focal point.

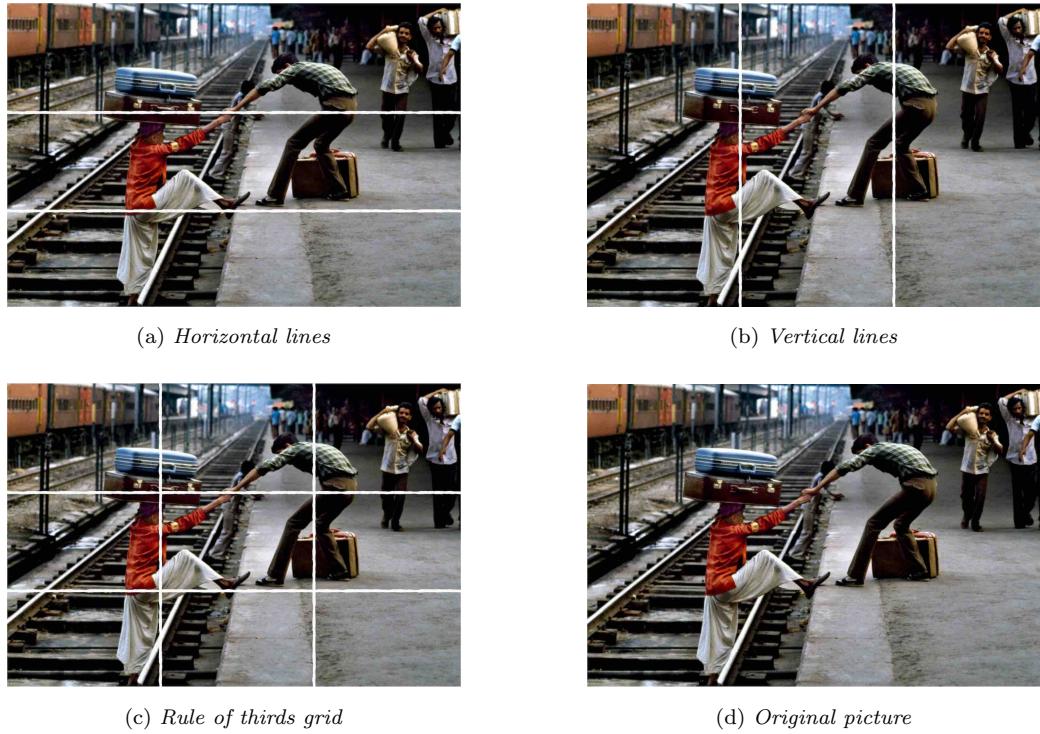


FIGURE (II)
The rule of thirds in “Hand in hand” by Steve McCurry

Notes: The Figure illustrates an application of the “rule of thirds”, which the photographer uses to guide the viewer’s attention towards the elements of interest (placed along the two vertical and two horizontal lines that divide the image in equal thirds).
Picture: *India, “Hand in hand”* gallery by Steve McCurry..

The visual vocabulary includes four main individual features related to subjects’ centrality: an indicator for high centrality (index higher than .95), medium-high centrality (index ranging in .85 and .95), medium-low

¹³This is only one of the ways to measure viewer’s attention towards a subject. Other aspects that are relevant to this extent include color vividness, body posture etc.; however, they do not form part of my “raw information” set on the images, and thus nevertheless necessarily fall outside the scope of the present work.

centrality (index ranging in .85 and .95), and low centrality (index strictly lower than .75).

-Kinesics. Kinesics is broadly defined as the study of body movements (Bowden, 2015; Furnham and Petrova, 2010; Walters, 2011). It entails any body dynamics such as gestures, facial expressions, eye behavior, or touching. Some gestures, such as facial expressions or eye movements, have been shown to be important markers of the emotional and cognitive inner state of a person. The particular look on a person’s face, for instance, provides reliable cues as to approval, disapproval, or disbelief (Bailenson et al., 2008; Grabe and Bucy, 2009; Lunenburg, 2010). Importantly, those elements are also relevant political cues in pictures’ visual narratives: during the 2016 US election campaign, for instance, news websites of varying ideologies portrayed the two candidates displaying more positive and less negative emotions of the candidate they supported (Peng, 2018; Boxell, 2021). In line with these findings, I include in the visual vocabulary a feature for each of the seven emotions detected by the emotion recognition algorithm (happiness, sadness, anger, fear, contempt, disgust and surprise).¹⁴ I additionally create a “negative emotion” indicator taking value 1 whenever the subject expresses either sadness, anger, fear, contempt or disgust. Correspondingly, I create a “positive emotion” indicator taking value 1 when the subject’s face expresses happiness. The above negative- and positive-emotion indicators exclude surprise: as leading pictures often portray subjects during speeches or public appearances, the depicted persons often have their mouth open, a pose the algorithm often mistakenly associates to surprise. Whenever none of the emotion variables (including surprise) takes value 1, I code the emotion as “neutral”.

While single-subject, close-up portraits focus the viewer’s attention on the portrayed person’s characteristics (and in this context, her emotion), shots with many subjects tend to shift the attention from an individual to a group. For this reason, the interaction of multiple persons shapes the overall emotion expressed in a picture. For the visual vocabulary to capture this aspect, and to deepen the understanding of images’ emotional loading, I produce a “triggered emotion” feature: for each portrayed person, it is the average emotion of the subjects whose glance is directed towards the person.¹⁵ For each subject in a picture I additionally include an indicator for the number of other people who are looking towards her, as well as one indicating whether she is looking towards someone. The vocabulary also captures a number of other kinesics dimensions. It encompasses indicators for each of the individual head pose angles used to construct the triggered emotionality (namely positive, neutral, or negative yaw/pitch/roll), as such angles by themselves can affect a viewer’s judgement of the social relations among the depicted subjects (Ekman, 2009); Gawronski and Payne, 2011). It captures two types of color-related factors that can be used to augment or decrease the salience of a subject in a picture: the light exposure of a subject’s face (measured through indicators for “normal exposure”, “bright exposure”, and “dark exposure”) and its blur level (measured through indicators for low blurring, medium blurring, and high blurring).¹⁶ The vocabulary also contains indicators for non-verbal cues related to the mode of dress: those capture whether the portrayed subjects are wearing clothes or accessories with particular distinctive features,

¹⁴I consider the emotion as correctly identified when the algorithm expresses a confidence level of 80% or higher (as in Peng, 2018).

¹⁵If persons X , Y and Z are portrayed in a picture, with X and Y looking towards Z , the latter’s triggered emotion is the average of X and Y ’s expressions. I use the subjects’ head pose angles as proxy for their glance direction and to compute their sight regions, borrowing this approach from studies on intelligent vehicle systems. Appendix Section A.1.1 describes the details of the method.

¹⁶Importantly, there exist a number of other color-related dimensions that are known to be relevant devices in visual communication (such as the overall color dominance in a picture, or the presence of evocative color patterns); unfortunately, those dimensions are outside the scope of the algorithms employed.

such as uniforms, formal dresses, suits, ties, hats, face masks, scarves, or glasses. From a visual communication perspective, individual clothing pieces –e.g. a suit– can be used to prime specific evaluations – such as success and wealth (Ekman, 2009). In addition, and particularly in the period under analysis, the presence of specific clothing or accessories –such as face masks– can *per se* constitute a politically and socially important signal. Appendix Section A.1.2 provides a summary of all Adjective subclasses, with *AS* labelling features pertaining size, *AC* indicating centrality, and *AK* indicating kinesics.

Features Class “Context elements” (C):

The third and last Syntax class of features encompasses indicators for the presence of specific contextual elements, varying at the image level. Previous research on political candidates’ imagery has shown the communicative relevance of contextual features –such as the portrayal of many individuals together– or of structural characteristics –such as the camera angle– (see, e.g., Sutherland et al., 2013; Abele et al., 2016; Haim and Jungblut 2021). In light of the existing evidence, the visual vocabulary includes indicators for the following elements varying at image level and defining its context: the presence of any human face; the total number of faces; the presence of any well-known person and the total number of recognized individuals; the joint presence of any two individuals from a “common subjects” lists encompassing all well-known individuals who appear in at least 50 images; the presence and total number of members of the Republican (Democratic) party sitting in at least one of the 114th-117th Congresses; the presence and number of politicians in each range of the politicians’ partisanship distribution; the presence and total number of men or women in images featuring at least one human face; the relative presence of the two sexes (majority women/men/equal number); the average emotionality –neutral, positive or negative– for pictures featuring at least one human face, measured through the average emotion expressed among all faces; the number of faces expressing each of the seven emotions, in pictures with at least one recognized individual; the presence, total number, and relative frequency (majority/minority/equal number) of individuals wearing a facial mask; the overall camera angle of the picture, measured through the average head poses of the portrayed individuals.

Other individual features pertaining context are conveyed (without transformations) by image tags. As mentioned, tags are indicators for the general presence (that is, without coordinates details) of particular elements or objects within the image. These mark the presence of natural elements (e.g. fire, water, etc.), transportation means (e.g. cars, ambulances, etc.), architectural elements (e.g. skyscrapers, castles, etc.), or text content (e.g. banners, signals, etc.).¹⁷ Because tags information doesn’t include location details (i.e., there are no bounding boxes, as mentioned), from a list of all possible tags in my dataset I manually classify them into subclasses, to expand the list of tags through meaningful tag mixes. Appendix Section A.1.2 provides a summary of Context tokens, with *CNtagmix* indicating features derived from tags, and *CNtxt* indicating other general context features.

II.B.4 Engineering: features Combinations

As mentioned, text-analysis vocabularies include bigrams and trigrams (couples and triplets of consecutive words). This allows to decode the pertinence of two language elements to the same textual object through

¹⁷An example of image tagging is available at <https://learn.microsoft.com/en-us/azure/cognitive-services/computer-vision/concept-tagging-images?tabs=3-2>.

words proximity. As noted, visual language doesn't have a clear order of words; for this reason, I rely on features-combinations to model the pertinence of multiple characteristics to a portrayed object, exploiting overlapping locations or simple co-occurrence.

The bigrams and trigrams in my visual vocabulary are represented by features pairs and triplets. To combine features in a meaningful way, I exploit the syntactic role of the different features classes that distinguish among represented subjects, characteristics of their depiction, and characteristics of the context and interactions between subjects. This structure is intended to help decode the meaning of images the same way words syntax helps to analyse a text: individual features convey the lexical composition of an image, while features' combinations inform about their "visual syntax". In the remainder of this work, I refer to vocabulary entries as *tokens*, a term that indicates either an individual feature or a combination of features.

The structure of syntax classes combinations is summarized in Table I. Table II presents summary statistics for the three syntax classes within the visual vocabulary, distinguishing between tokens made of single features (upper panel) and of feature-combinations (lower panel).¹⁸ The vocabulary encompasses a total of 4'541 different individual subjects (row 1), 40 different individual adjectives (row 2), and 541'488 individual elements of context (row 3), for a total of 519'069 distinct tokens derived from single features. This adds to more than 2.5M distinct combinations, of which 1.2M involve a Subject (row 4), 0.7M involve an Adjective (row 5), and 1.8M involve a Context element (row 6). Column 6 of Table II summarizes the occurrences of vocabulary tokens in the images analyzed images in this study. Overall, the pictures display 0.9M individual Subjects, 2.2M individual Adjectives, and 6.1M individual context elements. In combinations, Subjects score an additional 28.4M appearances, Adjectives 26.2M, and Context elements 16.9M. Appendix Tables A.1.5, A.1.6, and A.1.7 provide additionl summary statistics for each of the subclasses within the Subjects, Adjective, and Context groups, respectively.

¹⁸Note that these statistics necessarily only encompass visual words that are actually part of the visual language of analysed images. For instance, suppose the feature extraction algorithms are trained to detect the presence of Johnny Cash: the vocabulary would however only include a token with his name if the singer was present at least once across the images analysed.

TABLE (I)
FEATURES ENGINEERING: SUMMARY OF COMBINATIONS.

	-	SR	SN	SC	SP	SG	SRxSC
S	SR	✓					
	SN	✓	✓				
	SC	✓	✓				
	SP	✓	✓		✓		✓
	SG	✓	✓		✓		✓
A	AK	✓	✓	✓	✓	✓	✓
	AC	✓	✓	✓	✓	✓	✓
	AS	✓	✓	✓	✓	✓	✓
	ASxAC	✓	✓	✓	✓	✓	✓
	AKxACxAS	✓	✓	✓	✓	✓	✓
	AKxAS	✓	✓	✓	✓	✓	✓
	AKxAC	✓	✓	✓	✓	✓	✓
C	CNtagmix	✓	✓	✓			
	CNtxt	✓	✓	✓			
	CNtxt×CNtagmix	✓					

Notes: The Table summarizes the combinations of features included in the vocabulary, marked by “✓”. The tokens’ syntax classes are: **Group S**= Subject features: either defining subjects’ constant characteristics (*SN*= person name, *SP* = political partisanship, *SG*= sex, *SC* celebrity status) or varying ones (*SR* = Saliency Rank). **Group A**= Adjective features: those encompass aspects related to Kinesics (*AK*), Size (*AS*), and Centrality (*AC*). **Group C**= Context features: those include general context features (*CNtxt*), and features derived from tags and their mix (*CNtagmix*).

TABLE (II)
VOCABULARY SUMMARY STATISTICS BY SYNTAX CLASS

Syntax classes:	N distinct tokens in syntax class	Class share of total tokens	Total tokens across classes	Total class presence in all pictures
<i>Single features:</i>				
Subject (“S”)	4'541	.00875	519'069	913'716
Adjective (“A”)	40	.0000771	519'069	2'227'114
Context (“C”)	514'488	.991	519'069	6'129'115
<i>Combinations of features:</i>				
Subject (“S”)	1'225'751	.461	2'661'417	28'485'577
Adjective (“A”)	753'168	.283	2'661'417	26'177'849
Context (“C”)	1'885'382	.708	2'661'417	16'973'802
<i>of which:</i>				
Subject×Adjectives	753'168			348'664
Subject×Context	449'716			249'599
Subject×Subject	22'867			10'431
Context×Context	1'435'666			640'583

Notes: The Table presents summary statistics for the three syntax classes within the Visual Vocabulary, distinguishing between tokens that consist of single features (upper panel) and of feature-combinations (lower panel). Column 3 presents the number of distinct tokens for each class, i.e., the distinct values a class can assume. Column 4 indicates the proportion of tokens in the class relative to the total number of (single or combined) vocabulary features, which is reported in Column 5. Column 6 lists the total number of occurrences of tokens belonging to the class within the analyzed images. Note: all statistics encompass only tokens that appear at least once in the images analyzed.

II.C Measuring Visual Partisanship in US news

II.C.1 Pre-processing

Following in parallel the text analysis by Demszky et al. (2019), I restrict my attention to features used at least 10 times in at least one of the 2-week periods, used in at least 10 different periods, and used at least 50 times across all periods. Following Gentzkow and Shapiro, 2010, I remove features that appear too frequently because their use is likely not informative about the inter-party differences I wish to measure. Following Demszky et al. (2019), I also remove infrequently used features to economize on computation.¹⁹ The resulting vocabulary contains 3.18M unique visual tokens used a total of 53.70M times in 298'850 leading images. I analyze data at the image level and within time periods of two weeks, for a total of 26 periods between Dec 2019 and December 2020.

II.C.2 Estimating Visual Partisanship

I study the *visual* partisanship in leading images by adapting the leave-out estimator of phrase partisanship by Gentzkow, Shapiro, and Taddy (2019).²⁰ Like these authors, I define partisanship as the expected posterior probability that an observer with a neutral prior would correctly guess a picture's political leaning (i.e. whether it was published by a Republican-leaning or Democrat-leaning source) after observing a single token randomly drawn from the image. If the token is used equally in images published by Republican- and Democrat-leaning news sources, then this probability is .5 and the token is uninformative of the image's political leaning.

The leave-out estimator solves the problem of finite-sample bias, which arises because the features an image could contain are many more than those present in any image leading the news. As a consequence, many pictures' features are used mostly by one party or the other purely by chance; however, naive estimators interpret such differences as evidence of partisanship, leading to a bias estimate that is much larger than the true signal in the data (see Gentzkow, Shapiro, and Taddy (*ibid.*)).

The leave-out estimate of partisanship π^{LO} between images from Democrat-leaning sources, $i \in D$, and images from Republican leaning sources, $i \in R$, is

$$\pi^{LO} = \frac{1}{2} \left(\frac{1}{|D|} \sum_{i \in D} \hat{\mathbf{q}}_i \hat{\rho}_{-i} + \frac{1}{|R|} \sum_{i \in R} \hat{\mathbf{q}}_i (1 - \hat{\rho}_{-i}) \right) \quad (2)$$

where $\hat{\mathbf{q}}_i = \mathbf{c}_i / m_i$ is the vector of empirical token frequency for image i , with \mathbf{c}_i being the vector of token counts for image i and m_i being the sum of token counts for image i ; $\hat{\rho}_{-i} = (\hat{\mathbf{q}}^{D \setminus i} \oslash (\hat{\mathbf{q}}^{D \setminus i} + \hat{\mathbf{q}}^{R \setminus i}))$ is a vector of posterior probabilities, excluding image i and any token that is not present in least two images. Following the notation in Demszky et al. (2019), I let \oslash denote element-wise division and $\hat{\mathbf{q}}^G = \sum_{i \in G} \mathbf{c}_i / \sum_{i \in G} m_i$ denote the empirical token frequency of images in group G .

As described by these authors, the LO estimator captures two components of image partisanship: the difference between groups (posterior probability for each feature) and the similarity within a group (dot-product between the feature vector of each speaker and that of their group).

¹⁹I cut tokens at the bottom 0.0025 of the tf-idf score distribution. The score is computed within event and at the news source level (given that, differently from text analysis, no image can contain a given feature more than once).

²⁰The analysis in this subsection greatly benefited from the availability of programming codes by Demszky et al. (2019) made available at <https://github.com/ddemszky/framing-twitter>.

II.C.3 Overall polarization

Figure III shows that in the entire period between December 2019 and Dec 2020, with one estimate every week. The visual language of leading images is highly polarized, with estimates ranging between .514 and .524, and maintaining a constant level around .518. These estimates are robust to the use of different event frequencies, such as monthly or weekly.²¹ Following Demszky et al. (2019), I quantify the noise by calculating the leave-out estimates after randomly assigning images to parties: the values resulting from random assignment are close to .5, suggesting that the actual values capture a true signal in the data.

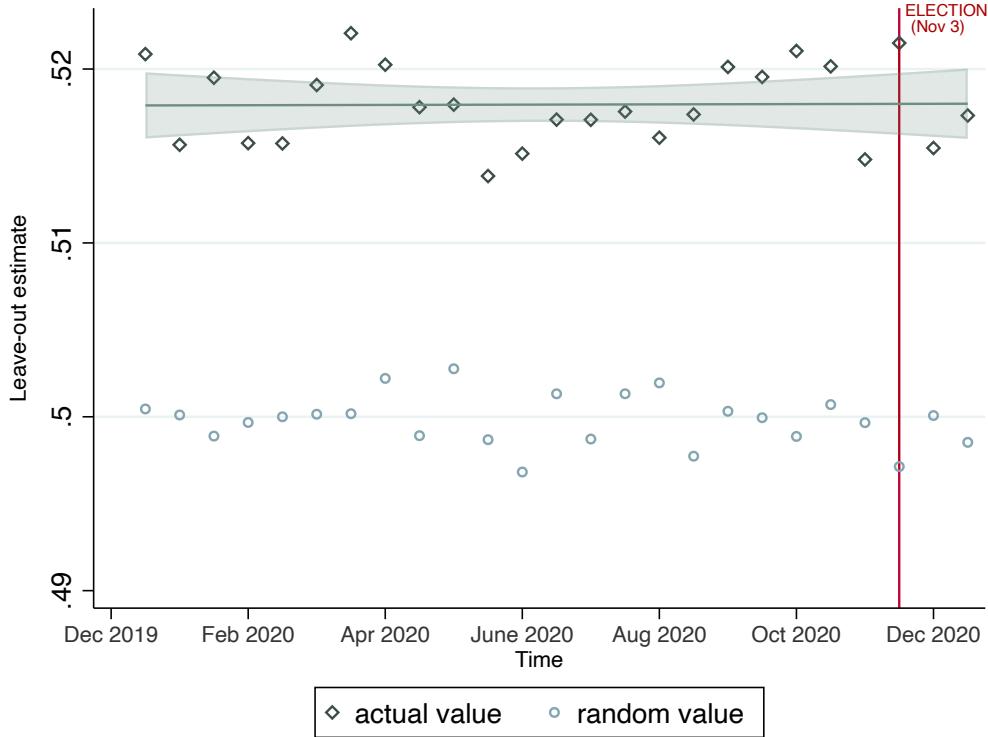


FIGURE (III)
Visual polarization

Notes: The Figure shows the partisanship of leading images estimated through the leave-out estimator (Gentzkow, Shapiro, and Taddy, 2019). The results indicate that the *visual* language of US news sources is highly polarized. The shaded region represents the 95% confidence interval of the linear regression fit to the actual values. The Figure replicates Figure 2 in Demszky et al. (2019). Like these authors, I quantify noise by calculating the leave-out estimates after randomly assigning images to parties: the values resulting from random assignment are all close to .5, suggesting that the actual values are not a result of noise.

As term of comparison for the magnitude of the estimated visual polarization, Gentzkow, Shapiro, and Taddy (2019) estimated the verbal polarization of US congress members from 1800s to present time. Around the year 2000 verbal polarization is .52, a figure close to the present estimate of visual polarization in US news. In another study, Demszky et al. (2019) use the same polarization measure to estimate the verbal polarization of Twitter discussions on mass-shootings in the US; their verbal polarization ranges between .517 and .547, again a magnitude comparable to that of the visual partisanship here estimated.

²¹See Appendix Figure A.1.1 for weekly estimates.

II.C.4 Polarization by topic

I further explore the extent of visual polarization dividing the images by the topic of the news pieces they lead. I model the news' topics by analysing the text of the tweets that linked to the articles.²² To this extent I use BERTopic, a topic modelling approach that operates through sentence-transformers to create embeddings, and exploits a class-based tf-idf for clustering.²³ The algorithm creates the tweets embeddings using a pre-trained BERT-based model for tasks of semantic similarity in English (“Paraphrase-MiniLM-L6-v2”). It lowers the dimensionality of the tweets embeddings with UMAP, to then cluster the reduced embeddings in semantically similar groups to define the topics.²⁴ To get a sense of the words composing each topic, I extract the most important words within each cluster through their within-cluster tf-idf score (“class-based tf-idf”, or c-tf-idf). The c-tf-idf score of a word is a proxy of information density: the higher the score of a word, the more representative it should be of its topic. Hence, the list of words with the highest scores provide for each topic an easily interpretable description. The unsupervised model identifies 75 granular topics²⁵ I manually inspect their descriptions to reduce their number to 8 macro-topics. The granular topics, their descriptions, and this hierarchical clustering are summarized in Appendix Section A.1.2. The 7 macro topics roughly pertain the following categories: environment (grouping related news including on natural events, animals, and climate), politics (grouping all news pertaining to domestic or foreign politics), health and covid (grouping all news broadly concerning healthcare, and those related to the pandemic from a medical perspective), economy (grouping all news pertaining to finance, economic policy, businesses and management), security (including news related to reform, social movements/protests, and crime), society (grouping news pertaining to education, the judicial system, and lawmaking), and entertainment (including movies, sports, and celebrity news). In addition, about half of the tweets eligible for the analysis by topic are assigned to a miscellaneous category: those are tweets sharing news pieces that pertain multiple topics equally, or whose topic is otherwise difficult to assign.²⁶ To preserve the internal coherence of the other main topics I consider the miscellaneous category separately.

The results in Figure IV indicate that the documented visual polarization during the entire period under analysis and *across* news topics is evident even *within* topics. Visually-partisan narratives stably populate the news in almost all of the above mentioned topics. Within the news on society (i.e. those related to education, the judicial system, and lawmaking) estimates are lower than in other domains, and for the first part of the year often indistinguishable from random values. The estimates obtained from “environment” news exhibit the highest variation. The overlapping confidence bounds in this topic should not be interpreted as evidence of a “shared” visual language by opposite sources: if anything, the frequent departures of random values from .5 indicate the poor reliability of the estimates in this topic. This suggests the visual vocabulary may contain too few of the elements relevant to characterize the partisan narratives of media sources for environment-related news. I discuss these aspects more in details in Subsection II.D.

²²This part of the analysis necessarily excludes all tweets that contain no other text than the url of the shared articles.

²³For more details on BERTopic, see [https://maartengr.github.io/BERTTopic](https://maartengr.github.io/BERTopic).

²⁴Using HDBSCAN for clustering.

²⁵I run a sensitivity analysis with a topic number ranging from 5 to 150 clusters. The optimal number of topics was identified both by comparing the BERTopic coherence score and through manual inspection of the topics' descriptions.

²⁶For example, the following is a news piece equally relevant to the “security” and “entertainment” topics, and more specifically to the granular topics “police use of force” and “football”: “Police shot tear gas and rubber bullets into a massive crowd that lined the streets of Argentina’s capital city to pay their respects to soccer legend Diego Maradona, who died on Wednesday at the age of

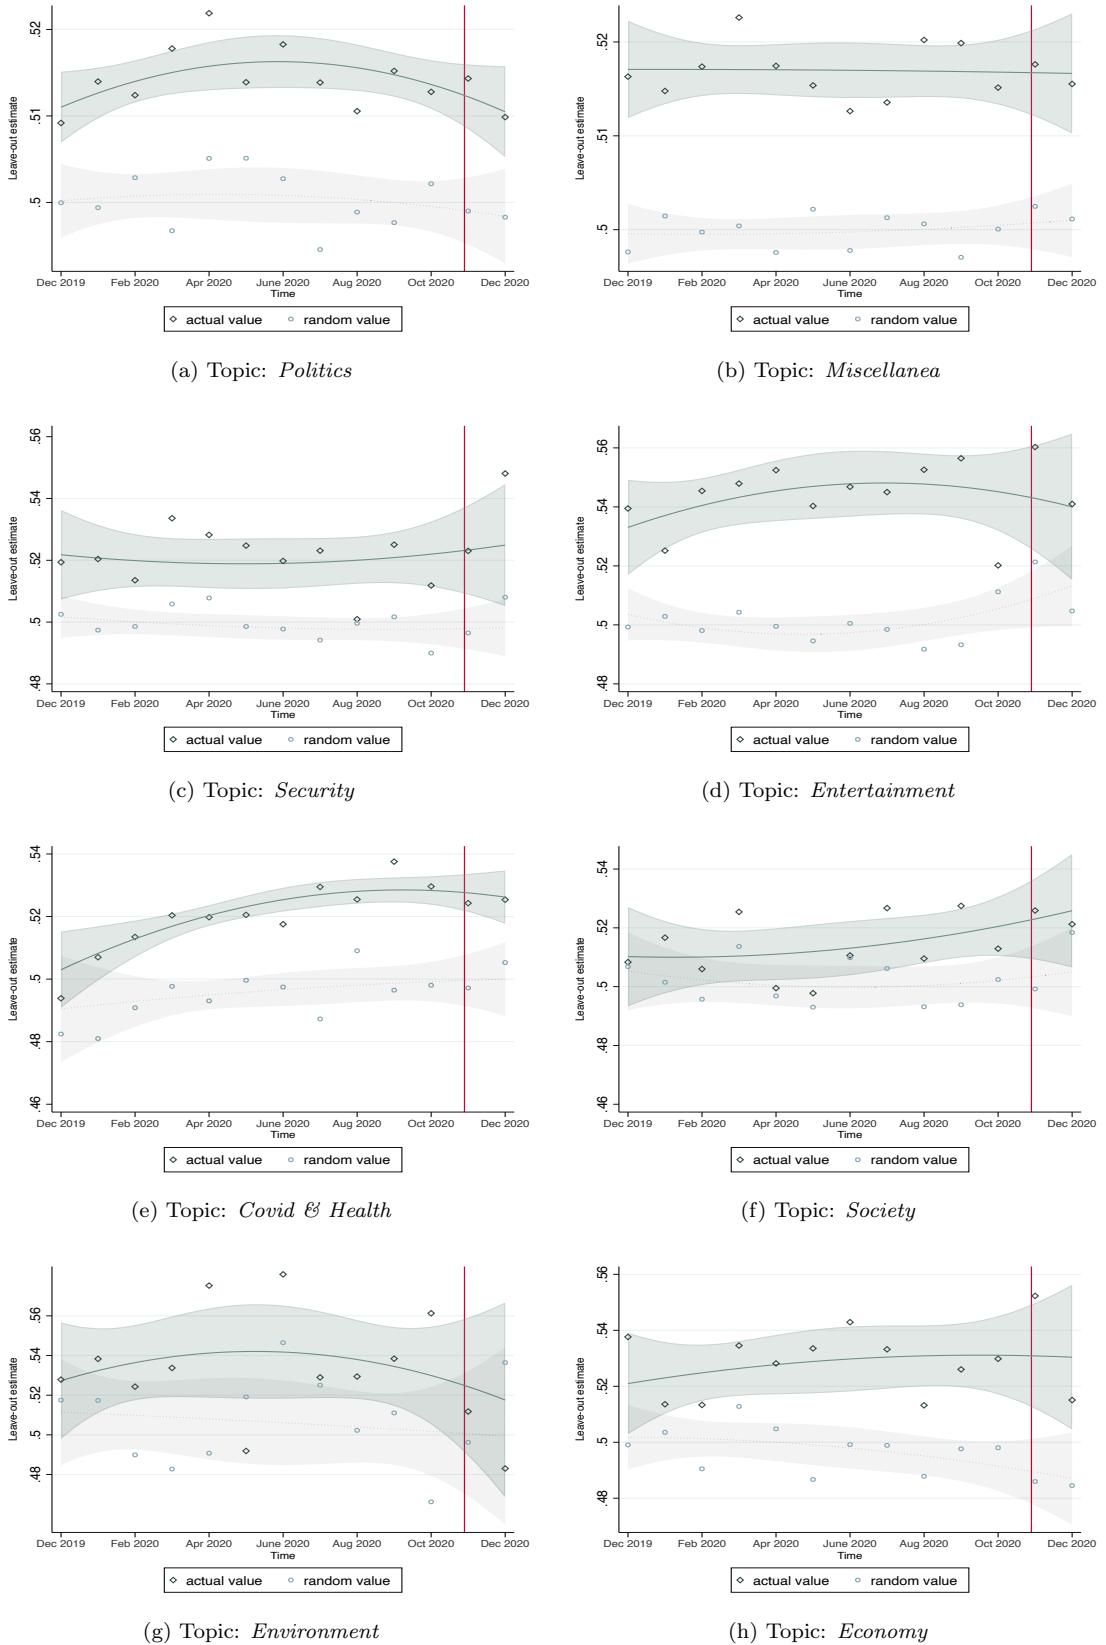


FIGURE (IV)
Visual Polarization By News Topic

Notes: The Figure shows the monthly partisanship of leading images estimated through the leave-out estimator by Gentzkow, Shapiro, and Taddy (2019). News are divided by topic: Politics (a); Society (b); Economy (c); Covid & Health (d); Entertainment (e); Environment (f); Security (g); Misellanea (h). The shaded regions represents the 95% confidence interval of a second order polynomial fit to the actual values. Random values correspond to leave-out estimates obtained after randomly assigning images to parties. The vertical red line marks the date of the 2020 Presidential election (Nov 3).

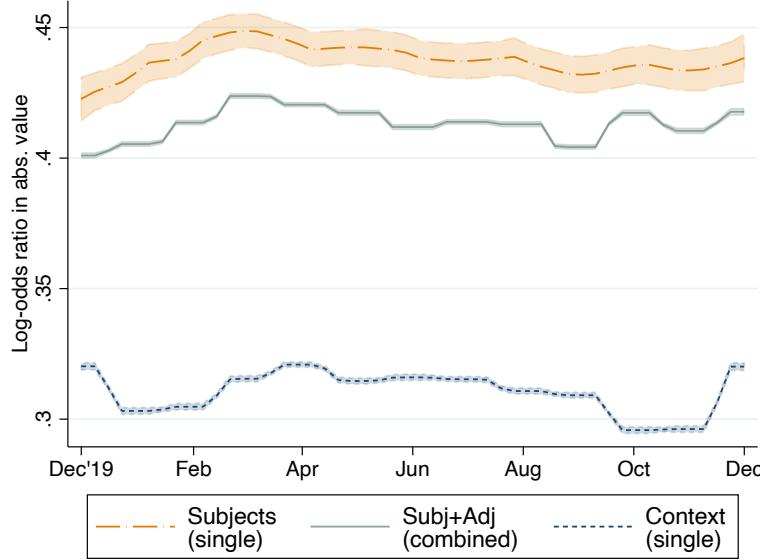


FIGURE (V)
SYNTAX ANALYSIS OF VISUAL PARTISANSHIP, BY NEWS TOPIC

Notes: The Figure shows the absolute value of log-odds ratios over time, for three syntax groups. The more the absolute value deviates from 0, the higher the average partisanship of tokens in the group (disregarding the direction of partisanship). Tokens are organized in three syntax groups: subject features (not combined), features combinations of adjectives and subjects, and context features (not combined).

II.C.5 Syntactic and Lexical Analysis of Visual Partisanship

This subsection explores how the visual languages of Republican-leaning and Democrat-leaning sources differ, and which visual features types are distinctive ideological markers. I study the partisanship of the visual tokens via the tokens' log-odds ratio of Democrats relative to Republicans, within each topic and across all periods. I distinguish between three syntax categories that offer distinct framing devices: represented human subjects ("subjects"), characterizations of subjects ("subject+adjectives"), and characterization of the context ("context").²⁷.

Figure V shows the evolution over time of the average log-odds ratios in absolute value for each of the groups. The more this value deviates from 0, the higher the partisanship of the tokens in the syntax group (disregarding the direction of partisanship). Over the entire period under analysis, tokens in the "subject" class had the highest differential use by news sources of opposite political leaning. Loosely speaking, this indicates that images from opposite sources tended first and foremost to depict different subjects. To a lesser extent, these sources' narratives also differed in terms of *how* (with which visual adjectives) subjects were represented. Context elements were stably the least relevant partisan device: net of who is represented and of her characterization, visual narratives of news sources with opposite leaning differ less in their use terms of the places or situations depicted.

The evidence presented in Figure V is particularly relevant from the perspective of news readers, because audiences are plausibly less aware of "subject-driven" partisan visual narratives than of "adjective-driven" ones.

^{60".}

²⁷This set includes inanimate objects, animals, physical spaces or social events as specified in the previous subsection and in Appendix Section A.1.2

In fact, personal intuition is sometimes sufficient to unmask very partisan characterizations. Instead, news readers can only unveil subject-driven visual partisanship if they actively expose themselves to news coverage from multiple sources.

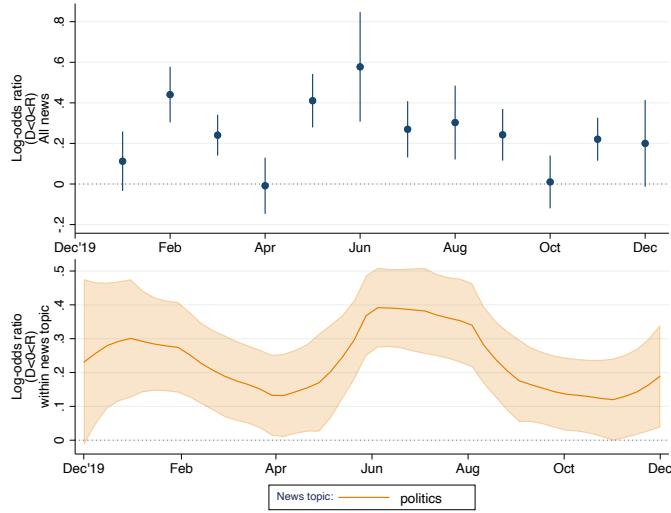
I then turn to analyse the “lexicon” of visual partisanship. In text documents, this type of analysis has proven useful to detect clear partisan markers (such as the expression “death tax”), whose partisanship is constant across periods and topics. Hence, repeated exposure to documents with these tokens is likely to develop readers’ intuition, and enable them to become aware of a latent partisanship in a document.

In the visual language of US media I identify three types of patterns. Some visual tokens are clear partisan markers: they display the same partisanship across all time periods and topics. Other tokens are partisan depending on the context in which they are used, as their partisanship changes across topics. Finally, for a third type of tokens partisanship evolves both over time and across topics. Figure VI presents an example from each type. The upper panel provides a case of clear visually-partisan marker: the characterization of President Biden with negative emotionality. The Panel shows the log odds ratios of the token over the entire period under analysis: the consistently positive score indicates systematically higher frequency of use by Republican-leaning news sources, persisting over time across all news topics (top half of the panel). The pattern is even more marked in the news pertaining to politics (lower half of the panel). The middle Panel in Figure VI presents an example of a token whose partisanship changes across topics. This is the context token “fire”. The top half of the panel displays log odds across all topics: the tokens is generally more likely to appear in images from Republican-leaning sources. The bottom half of the panel displays the log odds ratios in the two topics in which the token is used more. This part of the graph reveals that while “fire” is generally a Republican marker in news pertaining to the security topic, it is instead a Democratic marker in news pertaining to environment. Finally, the bottom panel of Figure VI plots the log odds ratios of a token resulting from a feature combination of a human subject (police) with a context element (a weapon). While overall this appears to be a more Republican marker, in the months of May and June 2020 the token is more likely to be part of the visual narratives of Democrats. This corresponds to the period following the homicide of George Floyd, and the resulting protests. The coverage of these particular news by the Democrats involved the representation of armed police forces, as confirmed in the lower half of the panel, where the log odds ratios pertain just the “security” topic.

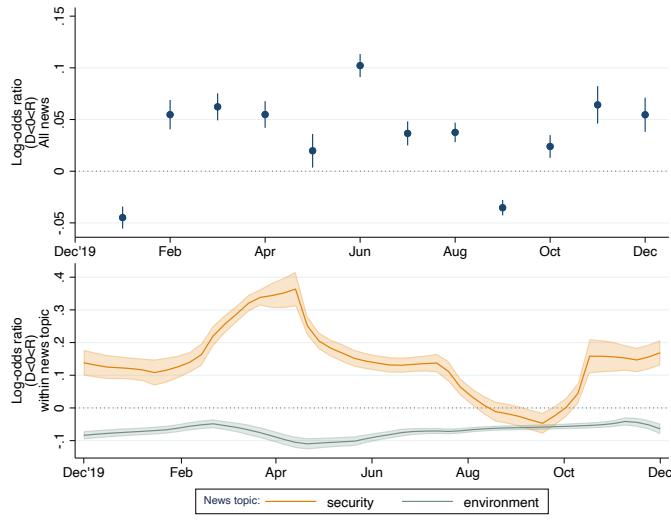
In sum, visual partisanship appears to be articulated in terms of tokens whose variability may change over time and contexts. This makes it plausibly more difficult for readers to unveil the partisanship of an image than of a text. The lexical analysis indicates that to detect partisanship in visual language readers need to know both the moment in time to which pictures refer to, and the topic they illustrate.

II.D Evaluating the Method: Perks and Pitfalls

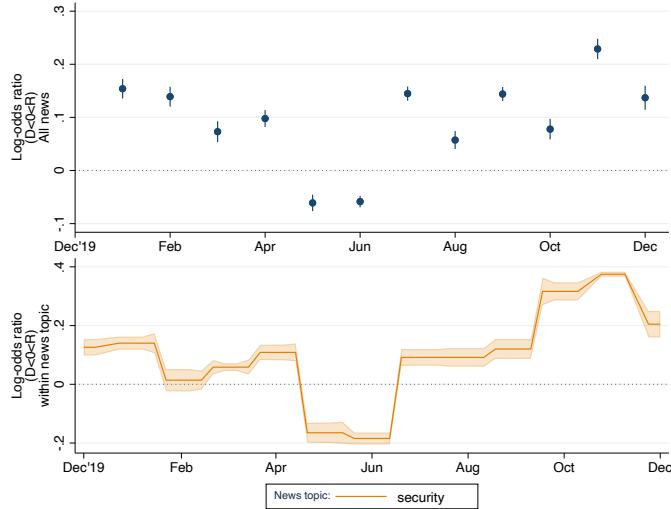
I here discuss some pitfalls and perks of the method proposed in this Section, that is, of adopting a dictionary approach to study pictures. The ability of dictionary-based procedures to extract meaning from a corpus is always tightly linked to the design of the underlying vocabulary. This is particularly relevant in the study of pictures, where the link between symbols and meaning is less express and more context-dependent than in words, as a vast semiotics literature suggests (Peirce,1931 Cassirer,1944; Morris,1946; Knowlton,1964 and



(a) Token: “*Biden & negative emotionality*”



(b) Token: “*Fire*”



(c) Token: “*Police & Weapon*”

FIGURE (VI)
PARTISAN LOG ODDS RATIOS OF TOKENS

Notes: The Figure shows the log-odds ratios for tokens in news overall (upper half of panels) and within topic (bottom half of panels).

1966 Veltrusky, 1976; Eco, 1979; Hołownia, 1981; Cassidy, 1982; Sebeok, 1985; Langer, 2009). For the application in this paper, the vocabulary's design aims to interpret the political nuances of the news' visual language; hence some nuances of the same visual expressions may remain undetected. For instance, the visual language of environment-related news could actually be very polarized, just in dimensions not captured by the dictionary used in this study (see Figure IV). It's hence worth remarking two observations on the method proposed in this Section. First, that inference is necessarily bounded within the "terminology" present in the vocabulary, as in all dictionary methods. The vocabulary hereby presented can easily be modified and adapted to serve different goals, through changes in lexical richness or syntactic breadth. Second, the aim of this Section is to establish a metric for detecting differences in visual language across the political spectrum. The analysis identifies language differences using the same vocabulary for all the images and all the sources: by doing so, irrespective of the vocabulary's completeness, it produces "internally valid" results within the dictionary employed. This is also relevant to the discussion over "algorithmic bias": since the approach in this paper rests on the same set of algorithms to extract features from all the images, any detected difference in visual language is *net* of algorithmic biases concerning the included visual words. Finally, constructing a visual vocabulary as described in the previous Sections grants the interpretability of the analysis results. Interpretable visual tokens allow to capture both "lexical" and "semantic" language traits, a goal otherwise difficult to pursue.

III THE EFFECT OF VISUAL PARTISANSHIP ON OPINION

The evidence presented so far in Section II indicates that US news visual language is significantly partisan. Nevertheless, to legitimately consider visual partisanship as an expression of media bias (a "*visual bias*"), leading images need not only be distinctive of Republican- or Democrat-leaning sources, but also to favor the corresponding party.

Given this, in this Section I test and verify that distant partisan visual narratives also have significantly different impacts on news readers' opinion. In particular, I introduce a survey experiment to test two hypotheses: first, whether partisan leading images distinctive of Republican/Democrat outlets slant the audience towards their respective party; second, whether partisan images, by interacting with the audience priors, increase the polarization of public opinion.

III.A Experimental strategy

I conduct a survey experiment on a nationally representative sample of the US population consisting of 2'000 respondents. The eligible population of the study consists of US citizens between 18 and 65 years of age. I recruit survey respondents on IPSOS's survey panel between July 2, 2021 and July 22, 2021. The experiment was approved by the European University Institute's Ethics Committee (the EUI's IRS board), and informed consent was obtained from the respondents at the beginning of the survey. The experiment and its pilot were pre-registered in the AEA RCT Registry.²⁸

Each respondent is exposed to news on five news issues, displayed sequentially. An issue is introduced through the following steps:

²⁸With respective digital object identifiers (DOIs): 10.1257/rct.7904-1.0 and 10.1257/rct.7247-1.0.

1. The respondent reads a short summary of the news.
2. She is asked to evaluate her knowledge of the issue and express the issue's relevance to her personal life/experience (i.e. its perceived salience), and her viewpoint on the issue.
3. The respondent accesses a page containing one piece of news on the issue. The piece appears in the same compact format of news previews when pieces are shared on social media, as described in the introductory Section. This preview's main elements are a brief summary-text on top (what is widely called "lead statement", an introductory text that summarizes the key details of a news piece), a leading image, a header (i.e. the title), and a byline (the line of text below the title generally providing context details). Respondents can read or skim the news pieces through a scroll-down movement, as in social media. The overall news look is that of the news previews featured on Facebook, and already illustrated in the right panel of Figure I.
4. After being exposed to the news, respondents express their opinion on the issue shifting a graphic slider to provide a numeric answer.

Steps 1-4 above defined repeat five times, one for each issue.²⁹ The experiment consists of exogenously varying the images leading the news in step 3 among three alternatives: non partisan, distinctive of Democrat-leaning sources (hereinafter: "Democrat-leaning") or distinctive of Republican-leaning sources (hereinafter: "Republican-leaning"). All the other aspects of the news previews (texts, headlines, bylines and graphic look) are held constant. Treatment assignment is randomized at individual level, and respondents are equally likely to be exposed to either of the three treatment branches (with treatment status for each issue being orthogonal to the status in others).

The text in the news pieces is non-partisan, namely it depicts facts covered by both liberal and conservative news sources without using partisan narrative frames or language.³⁰ Democratic- (Republican-) leaning images contain Democratic- (Republican-) visual features with high partisanship score (measured following the method described in Section II), hence they depict issues in a manner that is distinctive of Democrats (Republicans) news outlets. Vice versa, non-partisan images (hereinafter "neutral" images) contain features with low partisanship scores. The pictures' partisan loadings are homogeneous to ensure cross-issues comparability. Images are congruent with the true coverage on the same issues from outlets on both political sides (news pieces sourced from www.allsides.com). Appendix section A.2.1 displays and describes the chosen images.

The treatment news issues pertain five news topics characterised by significant and comparable visual partisanship as described in Section *II.C.4*, i.e. Security, Politics, Economy, Covid & Health, and Society. I select one recent news issue from each of those broader topics, and respectively: the debate on police budget cuts (hereinafter: "Police funds" issue); Biden's efforts to renew the 2015 US-Iran nuclear deal ("Iran deal"); the FED forecasts on inflation ("Inflation"); the anti-Covid measures implemented in March 2020 in the US ("Covid measures"); the institution of Juneteenth as Federal holiday ("Juneteenth").

I collect respondents' opinions on these issues through the following questions:

²⁹The order of issues is randomized.

³⁰For all the issues the text is comparable to that in news pieces on the same issues rated "non-partisan" on www.allsides.com.

- Police funds: *The total state and local government spending on police is currently about \$119 billion a year. If you were to decide the police budget, how much would you set it to?* [Answers readjusted to range in -100%/+100%]
- Iran deal: *From 0 to 100, in your opinion what is the probability for Biden to succeed in reviving the 2015 nuclear deal with Iran?*
- Inflation: *From 0 to 100, in your opinion what is the probability of inflation returning to pre-pandemic levels by July 2022?*
- Covid measures: *From 0 to 100, how much do you approve of the pandemic handling by public health experts in March last year?*
- Juneteenth: *From 0 to 100, how much do you support the creation of a new federal holiday for Juneteenth?*

The outcome variable of interest is the respondent's opinion on an issue after being exposed to the news.

For each news issue, I estimate the following specification through OLS:

$$Y_i = \beta_0 N_i + \beta_1 D_i + \beta_2 R_i + \beta_3 X_i + \epsilon_i \quad (3)$$

where Y_i is the post-treatment opinion expressed by respondent i on a given issue, N_i is an indicator for exposure to news led by neutral-leaning images, and D_i and R_i are similar indicators for exposure to news led by Democrat-leaning or Republican-leaning leading images, for the given issue. Finally X is a vector of demeaned control variables uncorrelated with the treatment indicators, to aid the precision of the estimates.³¹

Appendix Table A.2.1 summarizes the main variables of the study, and Appendix Tables A.2.3 to A.2.7 display the balance of observables characteristics across treatment branches for the 5 treatment news issues. As expected by virtue of randomization, for all the treatment news issues the respondents in the three treatment branches are balanced in terms of observable characteristics, and the standardized difference is always below the critical threshold of 0.25 suggested by Imbens and Rubin (2015).

III.B Do partisan images affect public opinion?

I investigate the extent to which Democrat- and Republican- leaning images shift public opinion by testing, for each news issue, the significance and the equality of treatment coefficients β_1 and β_2 in (3). The analysis presented in this Section excludes survey respondents who do not pass an attention check placed at half survey; it also discards single answers given after treatment exposures of strictly less than 5 seconds.³² Both exclusion criteria are pre-registered.

Table III reports the estimated treatment effects on the respondents' opinion on each issue. News issues are ordered by the distinctiveness of Democrats' and Republicans' prior ideological positions (most similar for

³¹These treatment-independent controls are indicators for: 4 age groups, ethnicity (White, Black, Latinx, Asian, Native American), literacy, political opinion (liberal-conservative), level of interest for politics, party preference, previous knowledge on the issue, perceived salience of the issue, opinion on the issue prior to treatment exposure, main type of information outlet (Radio, TV, Social networks, Newspapers), frequency of use for 6 media outlets (Fox News, Breitbart, New York Post, MSNBC, New York Times, CNN), technical aspects of the survey filling (indicator for low screen resolution, total number of clicks in the survey introduction), and State of residence fixed effects. Appendix Table A.2.2 reports the corresponding estimates from specifications omitting controls.

³²This is the time just sufficient to load the news page and immediately scroll down to the "next page" button.

Police funding, and least similar for Juneteenth, as indicated by the distribution of respondents' first opinions (Appendix Figure A.2.6). All dependent variables are readjusted to range between -50 and +50, with the exception of the "Defund Police" issue, whose opinion ranges between -100% and +100% of the true Police budget.³³ All the dependent variables have been adjusted so that higher and lower values correspond respectively to Democrats' and Republicans' ideological positions relative to an "indifference", intermediate position (marked with value 0); hence positive coefficients indicate a relative pro-Democratic opinion stance, and vice versa. Round parentheses contain robust standard errors, while square brackets contain the p-values for two-sided tests of equality between coefficients (with tested pairs indicated on the left) using heteroskedasticity-robust standard errors.

TABLE (III)
Impact of Leading Images On News-Readers' Opinion

Dependent variable:	(1) Opinion on "Defund Police"	(2) Opinion on "Iran deal"	(3) Opinion on "Inflation"	(4) Opinion on "Covid measures"	(5) Opinion on "Juneteenth"
	(Budget cut in -100 +100)	(Confidence, in -50+50)	(Confidence, in -50+50)	(Blame in -50+50)	(Policy support, in -50+50)
Neutral images (N)	-0.594 (0.745)	-0.421 (0.692)	1.084 (0.414)	-7.260 (0.223)	8.687 (0.457)
Democrat images (D-N)	1.376 (1.097) [0.210]	-0.283 (0.987) [0.775]	-0.548 (1.071) [0.609]	1.364 (1.291) [0.291]	-0.495 (0.836) [0.554]
Republican images (R-N)	-2.394 (1.309) [0.068]	-2.329 (0.898) [0.010]	-2.941 (1.118) [0.009]	-0.903 (1.352) [0.505]	-0.229 (0.830) [0.782]
Democrat-Republican (D-R)	3.771 (1.240) [0.002]	2.047 (0.910) [0.025]	2.392 (1.161) [0.039]	2.267 (1.366) [0.097]	-0.266 (0.840) [0.751]
Observations	1565	1599	1615	1584	1542
Controls:	Y	Y	Y	Y	Y

Notes: The Table presents OLS estimates of the effect of the Democrat-leaning (D), neutral (N), and Republican-leaning (R) news-leading images on respondents' opinion after exposure to the news. Column headers indicate the relevant news issue. The dependent variable for the "Defund Police" issue ranges in [-100,+100], while all others range in [-50+50]. Dependent variables are adjusted so that the maximum value corresponds to Democrats' ideological position (thus, positive coefficients indicate a pro-Democratic opinion, and vice versa). In the Table, round parentheses present robust standard errors and square brackets contain the p-values for two-sided tests of equality between coefficients (tested pairs are noted on the left). Treatment-independent controls are indicators for: 4 age groups, ethnicity (White, Black, Latinx, Asian, Native American), literacy, political opinion (liberal-conservative), level of interest for politics, party preference, previous knowledge on the issue, perceived salience of the issue, opinion on the issue prior to treatment exposure, main type of information outlet (Radio, TV, Social networks, Newspapers), frequency of use for 6 media outlets (Fox News, Breitbart, New York Post, MSNBC, New York Times, CNN), technical aspects of the survey filling (indicator for low screen resolution, total number of clicks in the survey introduction), and State of residence fixed effects.

Police funding. For this issue, the dependent variable is the answer to the question: "*If you were to decide the police budget, how much would you set it to? [Relative to the current budget of \$119 billion]*". Coefficients represent the deviation, in percentage points, from the indifference opinion (0). To ease the comparison with

³³Respondents are given as reference the State and Local total Police expenditure in 2018, amounting to \$ 118,800,032,000 (rounded to 119 billion). Data accessed on January 29, 2021 from: <https://state-local-finance-data.taxpolicycenter.org/pages.cfm>

other issues, the answers to this question (ranging between -100 and +100) have been adapted so that positive values indicate a budget decrease (i.e. a positive budget cut). Relative to the news piece led by a Republican-leaning image, the same news piece led by a Democratic-leaning picture significantly increases the desired budget cut by an additional 3.77 percentage points – equivalent to about \$ 4.5 billion in monetary terms (st. error = 1.240, p-value = .002). A comparison of the maximum opinion spread produced by image variation (that is, the difference between the largest and the smallest treatment coefficients) and the smallest effect exerted by news exposure (that is, the smallest coefficient in absolute value) provides an indication of the effect of visual partisanship relative to the more general effect of news previews. The rationale is the following: as all treatment branches display the same text content, all coefficients capture the effect of exposure to the constant elements (headline, summary, byline, etc.). Given this, any difference in opinion across treatment branches identifies the additional effect that image partisanship can exert on top of the overall effect of news previews.³⁴ In this first news issue, image variation can increase the desired Police budget cuts by up to 3.77 percentage points, (from a minimum of -2.39 to a maximum of 1.37), that is more than 6 times the increase produced from overall exposure to news previews (amounting to .54 percentage points, as indicated by the smallest coefficient in absolute value, that of the Neutral treatment).³⁵

Iran deal. For this issue, the dependent variable is the answer to the question: “*From 0 to 100, in your opinion what is the probability for Biden to succeed in reviving the 2015 nuclear deal with Iran?*”. The answers to this question have been adapted to range in -50 + 50 and the coefficients can once again be interpreted as deviation from the indifference position (0). All the coefficients are negative, indicating that after the exposure to the news piece all respondents perceive the deal success as less likely. However, compared to respondents exposed to Republican-leaning leading images, those exposed to Democratic-leaning images judge the deal success as significantly more likely, with a margin of 2.05 percentage points (st. error= .910, p-value = .025). Similarly, respondents’ exposed to neutral images report a higher perceived likelihood of the deal success (with a margin of 2.33 percentage points, estimated with st. error = .898 and p-value = .010).

Using the criterion above mentioned, namely comparing the coefficient range to the smallest treatment coefficient in absolute value, I conclude that while the deal news always produce a loss in confidence of Biden’s success, the variation in the news leading images can exacerbate such an effect, producing an additional confidence loss 5.54 times as big.³⁶

Inflation. For this issue, the dependent variable is the answer to the question: “*From 0 to 100, in your opinion what is the probability of inflation returning to pre-pandemic levels by July 2022?*”. Once again, the answers to this question have been adapted to range in -50 + 50, and coefficients represent deviation from an indifference stance. Compared to respondents exposed to Republican-leaning images, those who see Democratic-leaning images report a higher perceived likelihood of Biden’s success, with a 2.39 percentage points difference (st. error = 1.161, p-value = .039); the smallest effect is obtained by news exposure with Democrat-leaning images, with a .54 p.p. coefficient. Hence, the variation in images attains about 4.4 times the opinion change of the

³⁴Note that that the measured effects of images add to that of neutral (non politically partisan) text elements. The present experiment does not contemplate the impact of partisan text, neither by itself nor in combination with images.

³⁵More precisely, image variation produces an opinion change whose magnitude is 634% that produced by the news preview with a neutral image.

³⁶To illustrate, image variation in this news issue produces a change in opinion by up to 2.329 p.p.; exposure to the news attains a minimum opinion change of .42 p.p. (negative). The former is 554 % the latter.

news preview overall.

Covid measures. For this issue, the dependent variable is the answer to the question: “*From 0 to 100, how much do you approve of the pandemic handling by public health experts in March last year?*”. Also in this case the answers have been adapted to range in -50 + 50, and coefficients represent deviation from the indifferent opinion. Moreover, to ease the comparison with the other issues, a value of 50 marks Democrats’ ideological pole (i.e. the strongest blame for the pandemic handling).

Compared to Republican-leaning leading images, Democratic-leaning ones increase (i.e. decrease by less) the dissatisfaction for the pandemic management, with a 2.27 percentage points gap (st. error = 1.366, p-value = .097). As above, I compare the coefficients’ range to the smallest treatment coefficient (-5.89); image variation produces an additional approval increase of more than a third the increase from overall exposure to the news (+38%).

Juneteenth. For this issue, the dependent variable is the answer to the question “*From 0 to 100, how much do you support the creation of a new federal holiday for Juneteenth?*”. Again, the answers have been adapted to range in -50 + 50. For this issue however treatments lead to negligible differences and imprecise estimates: the opinion margin between Republican-leaning and Democratic-leaning images amounts to .266 percentage points, and the effect is statistically indistinguishable from 0 (st. error= .840, p-value .751).

Overall, the results in Table III indicate that leading images have a non-negligible impact on news readers’ opinion. Pictures distinctive of Democrats/Republican news outlets pull the audience towards their respective parties’ ideological poles. These results indicate that the visual partisanship above documented effectively favors outlets’ preferred political factions, and in so doing it translates to a proper “visual bias”, another tangible expression of political bias in the media.

Moreover, when images produce a significant impact on opinion, its magnitude ranges between 38% and 634% of the overall effect from exposure to news previews; in 3 of the 4 precisely estimated impacts, the “slanting effect” of pictures dominates that of other elements of the news previews, and notably of written content.³⁷ One implication of these findings is that a news piece rated as “non partisan” through a text-based analysis could still exert a partisan influence on readers. Hence, any measure of political bias in written news shall take into account the eventual proximity of text and images, and interpret them jointly.

Another pattern highlighted by the results in Table III concerns the decrease of the effect of images relative to text in the distance between parties’ ideological positions. As above mentioned, parties’ stances are most similar for the “Defund police” issue, and least similar for the “Juneteenth” issue.³⁸ The effect of images relative to text shrinks for two reasons: first, due to a decrease in the numerator (i.e. the maximum distance across treatment branches, which becomes smaller across columns from left to right); second, and more evidently, due to an increase in the denominator (the smallest coefficient in absolute value, which in the last column is more than 15 times bigger than in the first). Independently of the images leading the news, readers seem to react more to news previews covering issues for which the ideological positions across parties are more distinct. These patterns are suggestive but should not be taken as conclusive evidence: a formal assessment of such relationships

³⁷The statement refers to the written news summary featured in news previews like those encountered by scrolling social media feeds. It does not apply to a piece full-text, since survey participants could never access a full-length article.

³⁸See Appendix Figure A.2.6.

would require testing a large number of issues, hence falls outside the scope of the present work.

III.C Does visual partisanship cause opinion polarization?

I test whether the exposure to partisan images causes polarization to increase in the general public. As pre-registered, I study the heterogeneity of the treatment effects along respondents' political affiliation (Republicans, Democratic or Independents), examining the extent of opinion polarization within-party and between-parties.

I find that the polarizing effect of partisan images across parties dominates the polarizing effect within each party. While Republicans and Democrats hold significantly different opinions already *ex ante*, their exposure to partisan aligned leading images further exacerbates the opinion gap, increasing overall polarization.

III.C.1 Polarization within-party

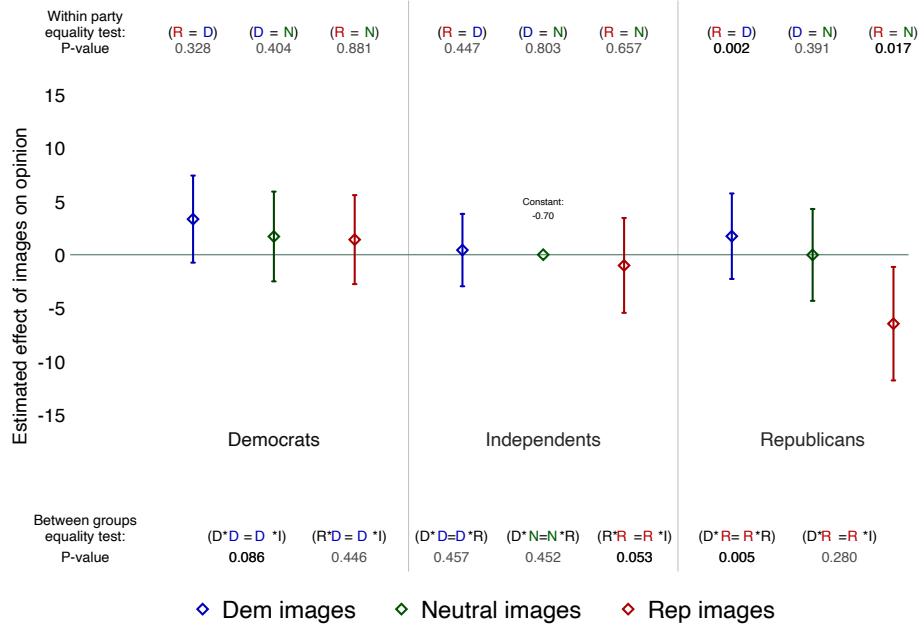
Within-party polarization occurs if the exposure to different leading images brings members of the same group apart. Figures VII, and VIII show the heterogeneous impact of leading images on individuals from different political affiliations separately for the *Defund Police*, *Covid measures*, *Iran deal*, and *Inflation* news issues.³⁹ The equality tests on top of each panel compare the treatment coefficients within the same party (e.g. whether Dem-leaning and Rep-leaning images have the same effect on Democrats); the p-values therefore measure the within-group polarization induced by visual bias for a given news issue. The equality tests at the bottom of each Figure instead compare the effect of exposure to the same leading image across different parties (e.g. whether Dem-leading images have the same effect on Democrats and on Republicans); the corresponding p-values indicate the extent to which one's party affiliation affects the impact of leading images.⁴⁰

The symmetry of treatment effects in all political groups is apparent already at a first glance: for a given news issue, the rank of coefficients for Republican-leaning, neutral, and Democrat- leaning images is nearly constant, suggesting the “direction” of each image slant is universal (i.e. individuals with opposite prior do not react in opposite ways to a given image). Vice versa, the relative distances among the three treatment coefficients varies with the political party, indicating that the effect of an image interacts with readers' prior. For Democrats, partisan and non-partisan images never have statistically distinguishable effects, with the exception of the difference between Republican-leaning and neutral images leading the news on Iran (p-value: 0.065). For Independents (that is all individuals not affiliated to either Democrats or Republicans), the only significant difference is once again between Republican-leaning and neutral images leading the news on Iran (p-value: 0.011). Among Republicans, the effect of Republican-leaning vs. Democrat leaning images differs for news on police funding (p-value: 0.002) and on inflation (p-value: 0.061); similarly, the effect of Republican-leaning images differs from that of neutral images for police funding (p-value: 0.016) and inflation news (p-value: 0.050).

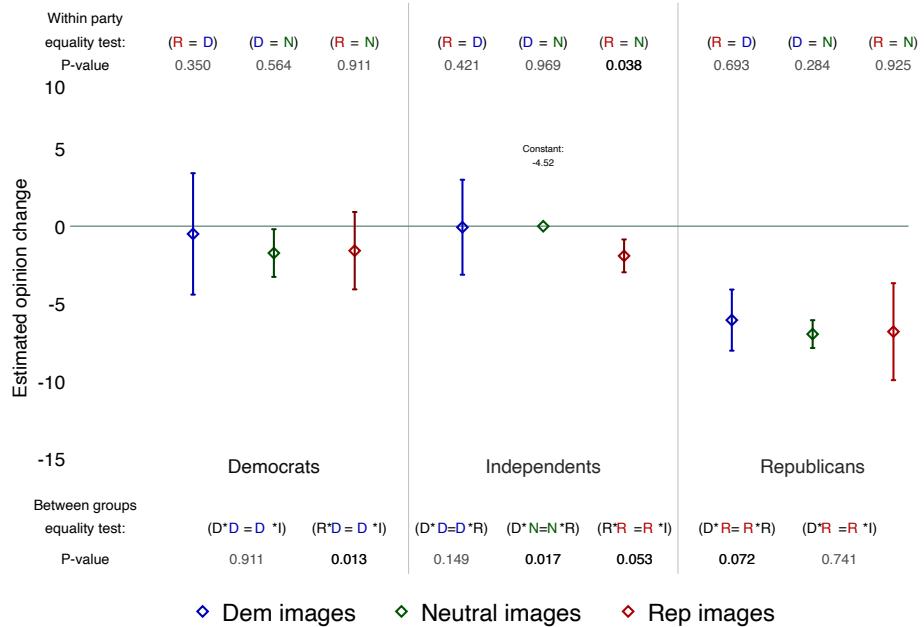
Does the exposure to the same partisan image produce equal effect on members of the Democratic and Republican party? On the one hand, the coefficient for Democrats exposed to Rep-leaning images is always statistically different from the coefficient for Republicans exposed to the same images (see the third coefficient from the left, marked by a red line, in each Figure). The equality between coefficients is rejected with confidence no less than 90% in all four news issues. On the other hand, the equality of effects for Democrats

³⁹ Appendix Table A.2.8 reports the point estimates for all the issues, including *Juneteenth*.

⁴⁰ The test label indicates respondents' party affiliation through the external letter, and the images while the internal letters. For instance, test (D*R=R*I) compares coefficients for Democrats and Independents both exposed to Republican-leaning images.



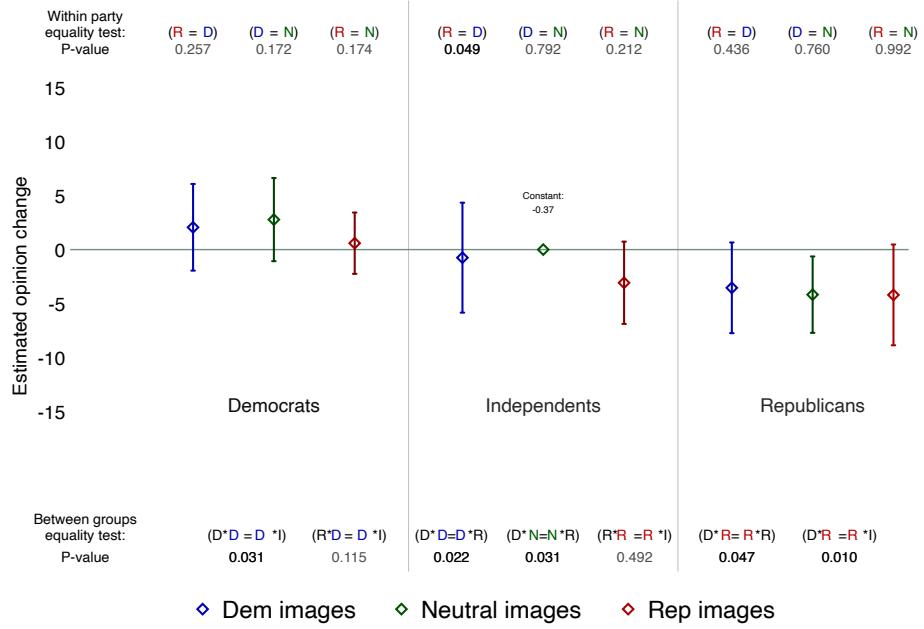
(a) News issue: *Defund Police*



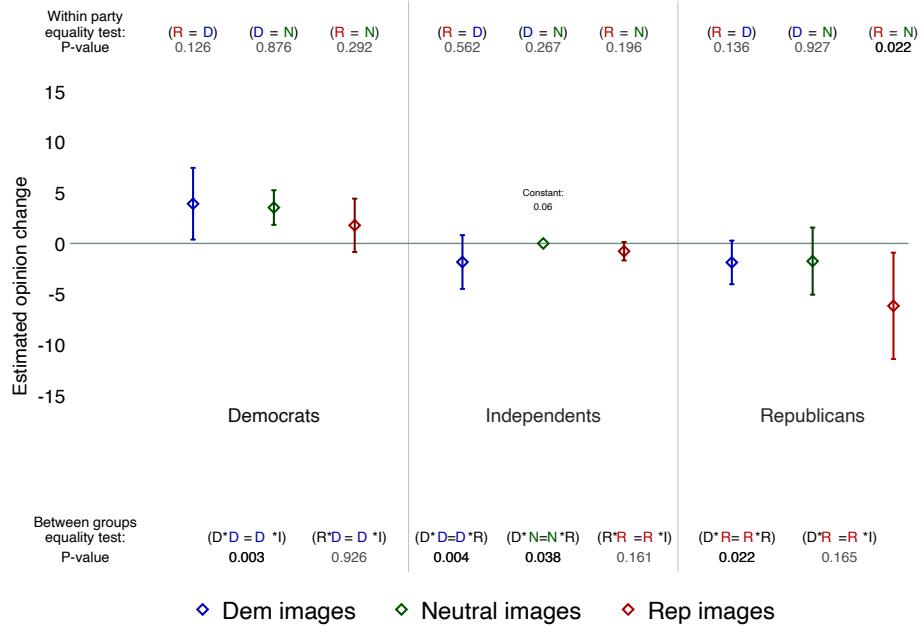
(b) News issue: *Covid Measures*

FIGURE (VII)
Heterogeneous effects of images on opinion,
by respondents' political party affiliation

Notes: The Figure shows OLS estimates of opinion changes after news exposure (news issues indicated below each panel). Treatments are interacted with respondent's party affiliation. Omitted regression category: Republicans exposed to Rep-leaning images. Lines indicate 95% CI (heteroskedasticity-robust st. errors). Equality tests on top of each Figure compare coefficients within each party; those at the bottom compare coefficients across parties (tested coefficients indicated in parentheses). All p-values are for two-sided tests of equality, with bold font marking statistical significance at 10 percent level or higher.



(a) News issue: *Nuclear deal with Iran*



(b) News issue: *Inflation*

FIGURE (VIII)
Heterogeneous effects of images on opinion,
by respondents' political party affiliation

Notes: The Figure shows OLS estimates of opinion changes after news exposure (news issues indicated below each panel). Treatments are interacted with respondent's party affiliation. Omitted regression category: Republicans exposed to Rep-leaning images. Lines indicate 95% CI (heteroskedasticity-robust st. errors). Equality tests on top of each Figure compare coefficients within each party; those at the bottom compare coefficients across parties (tested coefficients indicated in parentheses). All p-values are for two-sided tests of equality, with bold font marking statistical significance at 10 percent level or higher.

and Republicans exposed to Dem-leaning images is rejected with $< 90\%$ confidence in two of the four news issues, as indicated by the tests at the bottom of Figures VII and VIII (test label: $D^*D=D^*R$; p-values: 0.004 for Covid measures; 0.098 for Iran deal).⁴¹ Similarly, the equality test comparing Democrats and Republicans exposed to neutral images (test label: $D^*N=N^*R$) rejects the equality in the same two issues (p-values: 0.037 for Covid measures; 0.003 for Iran deal).

Overall, partisan images increase polarization within the Republican party more than within other parties. In general, Democrats and Republicans tend to react differently to the same pictures. However, part of their opinion gap accrues to their different political stances and it is independent of the image partisanship (it persists across all images). In the following section I explore how the effect of partisan images interacts with readers' political stance by widening or closing the opinion divide.

III.C.2 Polarization across parties

The equality tests at the bottom of Figures VII, and VIII indicate that part of the opinion gap across parties is, as above mentioned, unrelated to which images lead the news. However, the Figures also show that when Republicans (Democrats) hold the highest opinion on a given issue, Republican- (Democrat-) leaning images are also the ones that increase the opinion the most (the same is true with the opposite sign). This reinforcing effect has an important implication: the –always opposite– slants of Dem-leaning and Rep-leaning images, albeit mildly significant within each party, become significant if added up; that is, if readers are selectively exposed to partisan images aligned with their party's viewpoint (hereinafter: “partisan aligned images”).

The polarizing effects of partisan aligned images are verified by testing the equality between the maximum and minimum opinion distance images can instil between Democrats and Republicans. Formally, this amounts to testing whether the distance between Democrats and Republicans exposed to aligned images is greater than the distance between the same groups exposed to images that do not reinforce their prior.⁴² If members of different parties had homogeneous reactions to the partisan images, the tested quantities would be equal. Appendix Table A.2.8 reports the p-values of the corresponding one sided-tests (lower Panel, first line), showing that the null hypothesis is rejected in all four issues.

This indicates that whenever images affect the opinion, the direction and the magnitude of the effects are such that the selective exposure to aligned images relative to opposite images causes polarization to increase significantly. In the presence of “information echo chambers” (readers’ tendency to source news from like-minded outlets), these results imply that visual bias materially translates to a significantly more polarized public opinion.

To summarize, one could argue the causal effect of visual bias on public opinion polarization has three mutually reinforcing components: first, as documented in previous literature, readers’ tendency to source news from like-minded outlets; second, the news visual bias, whereby news sources predominantly use images that pull readers towards their ideological position; third, the interaction between the slanting effect of images and readers’ prior, with readers on both sides of the political spectrum reacting more distinctly to pictures that are more aligned with their viewpoint.

⁴¹The coefficients and the test for the Juneteenth issue are displayed in Appendix Figure ???. For this topic, the p-value of this test is 0.052.

⁴²This nets out the gap accruing to the political stance and independent of the images.

III.C.3 Which implications?

A large literature documents that recently –and in particular during the first months of the Covid-19 pandemic– partisan divisions significantly shaped health behavior, support to specific policies, attributions of responsibility, and general beliefs (e.g., Allcott et al., 2020a, 2020b; Druckman et al., 2020; Gadarian, Goodman, and Pepinsky, 2021); Gollwitzer et al., 2020; Romer and Jamieson, 2020). There is additional evidence suggesting that issue polarization is rising (see e.g. Doherty, Kiley, and Asheer, 2019), to which the present paper adds by demonstrating the causal effect of news visual bias in this direction.

Given that partisan aligned images significantly polarize public opinion, a legitimate question is whether, to the contrary, the opinion gap across parties would close if readers were hypothetically exposed only to images opposite to their prior (hereinafter: opposite partisan images).⁴³ The hypothesis of a gap closure can be tested verifying the equality of opinions of Republicans and Democrats exposed to opposite partisan images. Overall, the partisan divide shrinks, but the tests fails to reject the null in all issues.⁴⁴ Hence, while the data suggest a shrink in the opinion divide, the evidence is otherwise inconclusive.

III.D Underlying Mechanisms: When Does Visual Bias Bite?

This last Section explores the mechanisms underlying the effect of images on opinion. I study three possibly moderating factors: the ex-ante opinion held on the issue (readers' prior on the news issue), the perceived issue salience, and the respondents' self-reported prior knowledge on the issue.

In general, respondents who ex-ante hold more extreme opinions (in either direction) appear more receptive to leading images compared to respondents who hold more intermediate opinions. Neither issue salience nor prior knowledge on the issue appear to be strong predictors of respondent' sensibility to the images leading the news, and no neat patterns arise.

III.D.1 Images and ex-ante opinion

Do images exert a different effect on readers who previously expressed intermediate or extreme opinions? To respond, I explore the heterogeneity of treatment effects across terciles of respondents' prior opinion on each issue. Overall, the results show that respondents who ex-ante hold more intermediate opinions are less affected by leading images, compared to respondents who belong to the first/third terciles. There are no other cross-issue systematic patterns between the level of prior opinion and the magnitude of treatment effects.

Appendix Table A.2.9 reports the estimates from an OLS regression of the respondents' updated opinion on treatments interacted with terciles of prior opinion distribution, and Figure A.2.7 displays coefficients plots and equality tests. The following paragraphs discuss the heterogeneity of treatment effect separately for each news issue.

Police funding. For this issue, respondents who ex ante chose the lowest police budget update their response by further lowering the budget. Within this group, those who were exposed to Dem-leaning images chose an even lower budget than both those exposed to Rep-leaning images ($p = 0.067$) and those exposed to neutral images ($p = 0.015$).⁴⁵ Respondents in the intermediate tercile of prior opinion, who expressed the mildest variations

⁴³This exercise doesn't mean to suggest that a policy of this intent would be socially desirable.

⁴⁴The p-values for this (two-sided) test of equality are displayed in the second line of tests at the bottom of Table A.2.8.

⁴⁵P-values from equality tests reported on top of the panels in Figure A.2.7.

to the Police budget (in either direction), do not exhibit statistically different reactions to treatments. Finally, respondents who ex-ante were choosing the highest Police funding reduce the budget significantly more if exposed to news with Dem-leaning images as opposed to Rep-leaning ones ($p\text{-value} = 0.036$), and even more so if exposed to neutral images as opposed to Rep-leaning ones ($p\text{-value} = 0.006$).

Covid measures. For this issue, Rep-leaning and Dem-leaning images exert statistically different effects only among respondents who ex ante express the lowest judgement on the adequacy of anti-covid measures implemented in March 2020. Among those, people exposed to Republican-leaning images have a significantly more positive opinion on the Government's measures ($p\text{-value} = 0.034$) than individuals exposed to Democrat-leaning images. In the same group of respondents, there is no significant difference between those exposed to neutral images and the others. No differences across treatment branches exist in the middle and higher terciles of prior opinion.

Iran deal. Respondents who ex ante express the lowest belief in the success of a US-Iran nuclear deal decrease their judgement on the likelihood of success significantly more if exposed to Republican-leaning images than if exposed to either neutral images ($p\text{-value} = 0.007$) or Dem-leaning images ($p\text{-value} = 0.086$). In the intermediate tercile of prior opinion there is a difference between Rep-leaning and Dem-leaning images ($p\text{-value} = 0.043$), and no other difference across treatment branches. Once again, no differences across treatment branches exist in the higher tercile of prior opinion.

Inflation. Respondents who ex ante express the highest belief in the regress of inflation by June 2022 exhibit the largest upward opinion update if exposed to neutral images as opposed to Rep-leaning ones ($p\text{-value} = 0.018$). Otherwise, there are no other significant differences across treatment branches in either of the terciles of prior opinion.

III.D.2 Images and issue salience

Does the effect of images depend on how relevant is the news issue to the individual respondent? I investigate the relationship between issue salience and treatments by interacting the treatment indicators with the distribution terciles of perceived issue salience. Overall, respondents in the lowest and highest tercile of the perceived issue salience appear mildly more susceptible to the effect of leading images, relative to the respondents in the intermediate tercile. However, the evidence is inconclusive as to whether issue salience is a strong predictor of respondent' sensibility to the images leading the news.

Appendix Table A.2.10 reports the estimates from an OLS regression of the respondents' updated opinion on these interaction terms, and Figure A.2.8 displays coefficients plots and equality tests. The following paragraphs discuss the heterogeneity of treatment effect separately for each news issue.

Police funding. Respondents who are in the lowest tercile for perceived relevance of the police funding issue update their response by lowering the budget comparatively more if exposed to Dem-leaning images than to neutral images, albeit the difference is only modestly significant ($p\text{-value} = 0.099$). Respondents who perceive the issue as more relevant to them (i.e. those in the highest tercile of perceived issue salience) decrease the desired police budget by comparatively more if exposed to Dem-leaning images as opposed to Rep-leaning images ($p\text{-value} = 0.013$). No other significant differences exists across treatment branches in either groups. Similarly, individuals in the intermediate salience tercile do not exhibit significantly different opinion updates across any

of the treatment branches.

Covid measures. For this issue, none of the terciles of issue salience display significant differences in the effects across treatment branches.

Iran deal. Rep-leaning images and neutral images have significantly different effects both in the first and in the third salience tercile, with Republican-leaning images producing a relatively lower perceived likelihood of success of a US-Iran nuclear deal (p -values = 0.057 in the lowest salience group, and 0.070 in the highest salience group). No significant effect exist between these two treatment branches in the intermediate tercile; in this group, instead, the effect of Republican-leaning and Democrat-leaning images is significantly different, with the latter eliciting a higher perceived likelihood of success of the deal (p -value = 0.044).

Inflation. For this issue, respondents in the highest salience tercile display a significantly different response to Republican-leaning and neutral images. In fact, the latter induce a relatively higher perceived likelihood of inflation to return to pre-pandemic levels by June 2022 (p -value = 0.076). No other statistically significant differences exist across treatment branches in any of the salience terciles.

III.D.3 Images and opinion development

Does the effect of images depend on news readers' stage of opinion development? Does it depend on the knowledge about the issue? I explore whether the effect of images varies between respondents whose prior knowledge and opinion are more vs. less consolidated. Those are directly measured with a question before treatment takes place. While no neat patterns arise, image variation seems to affect highly knowledgeable respondents more often than the others (3 news issues displaying significant differences across branches, vs. 1 news issue for least knowledgeable respondents). The evidence on whether prior issue knowledge is a determinant factor is however inconclusive.

Appendix Table A.2.11 reports the estimates from an OLS regression of the respondents' updated opinion interacted with high and low levels of prior knowledge on the news issue, and Figure A.2.9 displays coefficients plots and equality tests. The following paragraphs discuss the heterogeneity of treatment effect separately for each news issue.

Police funding. Respondents who consider themselves not very knowledgeable about the issue do not update their response differently across the treatment branches. Vice versa, respondents who consider themselves highly knowledgeable about the issue update their response by lowering the desired Police budget comparatively more if exposed to Dem-leaning images than to Rep-leaning images (p -value = 0.006), and if exposed to neutral images than to Rep-leaning ones (p -value = 0.073). The effect of any image type never differs between the least and the most knowledgeable respondents (tests at the bottom of panels of Figure A.2.9).

Covid measures. For this issue, neither the most knowledgeable nor the least knowledgeable respondents' update their response differently across the treatment branches. Moreover, the effect of any image type never differs between the least and the most knowledgeable respondents.

Iran deal. Respondents who consider themselves not very knowledgeable about the issue update their response by increasing the perceived likelihood of success of a US-Iran deal relatively more if exposed to neutral images than to Rep-leaning images (p -value = 0.049). Vice versa, respondents who consider themselves highly aware about the US-Iran deal update their response by increasing the perceived likelihood of success comparatively

more if exposed to Dem-leaning images than to Rep-leaning images (p -value = 0.053). No other significant differences exist among treatment coefficients within either knowledge groups. Moreover, the effect of an image never differs between the least and the most knowledgeable respondents.

Inflation. Respondents who consider themselves not very knowledgeable about the issue do not update their response differently across the treatment branches. Vice versa, respondents who consider themselves knowledgeable about the issue update their response by increasing the perceived likelihood of inflation to return to pre-pandemic levels by June 2022 comparatively more if exposed to neutral images than to Rep-leaning images (p -value = 0.054). No other significant differences exist among treatment coefficients within either knowledge groups. Moreover, the effect Democrat-leaning images is mildly different between the least and the most knowledgeable respondents, with a slightly smaller positive effect in the latter group ($p= 0.092$).

III.E Limits of Inference from the Experiment: Within and Beyond

This Subsection discusses the limits of inference from the experiment, summarising the conclusions that can be drawn and the patterns that remain as speculation, hinting at future studies.

The main experimental results demonstrate that images can have a significant effect on opinion and that such effect directly causes higher polarization in the general public. The experiment was tailored to answer the research question of whether such effect existed, and its robustness across five tested topics. The estimates' external validity must however be considered with care. For instance, testing different images or news issues could lead to higher/lower estimates. In this sense, the estimates should not be interpreted as universal measure of the effect of leading images, but rather as evidence demonstrating that such impact exist and is non-negligible. Further research is needed to explore the sensitivity of the estimates to variation of different news elements, including graphic rendering (testing different formats of online news previews), news topics, or political slant in text.

The experiment also finds that in 3 out of the 4 news issues in which leading images had a significant impact, the opinion variation accruing to images was larger than the general effect of news previews, including the (politically neutral) text elements. This proves that images can be more relevant than text in affecting the opinion of online-news readers. It also demonstrates that newscasts can bypass text-based fact checking and still slant readers' opinion. However, it would be factually incorrect to generalize the result and conclude, for instance, that pictures dominate text in 75% of cases. The experiment measured the opinion of a large sample of individuals on a limited number of news issues (5): hence, any pattern arising from comparisons across issues offers, at best, suggestive evidence. Further research is needed to test simultaneously several issues and allow solid inference *across* news types.

IV CONCLUSION

This study shows the images leading written news pieces in the US exert a political influence on news readers. The first part of this paper quantifies the extent of visual partisanship in US news, finding a high degree of polarization in the visual narratives adopted by news sources across the political spectrum. The second part of the paper tests the direct effect of visually-partisan images on public opinion. It finds that partisan visual

narratives slant readers' opinion towards the outlets' ideological poles, hence that visual partisanship is an expression of political media bias. The experimental results also show that news visual bias has a positive causal effect on issue polarization, as the slanting effect of images interacts with news readers' priors: readers on both sides of the political spectrum react more distinctly to pictures aligned with their viewpoint. This pattern implies that the polarizing effect of visual bias is further exacerbated if readers' source their news exclusively from like-minded outlets.

NUFFIELD COLLEGE AND OXFORD UNIVERSITY

References

- Abele, A. E., N. Hauke, K. Peters, E. Louvet, A. Szymkow, and Y. Duan (2016). "Facets of the fundamental content dimensions: Agency with competence and assertiveness-Communion with warmth and morality". *Frontiers in psychology* 7, 1810.
- Allcott, H., L. Boxell, J. C. Conway, B. A. Ferguson, M. Gentzkow, and B. Goldman (2020a). *What explains temporal and geographic variation in the early US coronavirus pandemic?* Tech. rep. National Bureau of Economic Research.
- Allcott, H., L. Boxell, J. Conway, M. Gentzkow, M. Thaler, and D. Yang (2020b). "Polarization and public health: Partisan differences in social distancing during the coronavirus pandemic". *Journal of Public Economics* 191, 104254.
- Bailenson, J. N., S. Iyengar, N. Yee, and N. A. Collins (2008). "Facial similarity between voters and candidates causes influence". *Public opinion quarterly* 72.5, 935–961.
- Bowden, M. (2015). *Winning body language: Control the conversation, command attention, and convey the right message without saying a word.*
- Boxell, L. (2021). "Slanted images: Measuring nonverbal media bias during the 2016 election". *Political Science Research and Methods*.
- Cassidy, M. F. (1982). "Toward integration: Education, instructional technology, and semiotics". *ECTJ* 30.2, 75–89.
- Cassirer, E. (1944). "An Essay on Man (New Haven and London)". *Yale University*.
- Dari, S., N. Kadrilev, and E. Hullermeier (2020). "A Neural Network-Based Driver Gaze Classification System with Vehicle Signals", 1–7.
- De Vreese, C. H. (2004). "The effects of frames in political television news on issue interpretation and frame salience". *Journalism & Mass Communication Quarterly* 81.1, 36–52.
- DellaVigna, S. and M. Gentzkow (2010). "Persuasion: empirical evidence". *Annu. Rev. Econ.* 2.1, 643–669.
- Demszky, D., N. Garg, R. Voigt, J. Zou, M. Gentzkow, J. Shapiro, and D. Jurafsky (2019). "Analyzing polarization in social media: Method and application to tweets on 21 mass shootings". *arXiv preprint arXiv:1904.01596*.
- Doherty, C., J. Kiley, and N. Asheer (2019). "Partisan antipathy: More intense, more personal". *Pew Research Center*.

- Druckman, J. N., S. Klar, Y. Krupnikov, M. Levendusky, and J. B. Ryan (2020). "How affective polarization shapes Americans' political beliefs: A study of response to the COVID-19 pandemic". *Journal of Experimental Political Science*, 1–12.
- Eco, U. (1979). *A theory of semiotics*. Vol. 217. Indiana University Press.
- Ekman, P. (2009). *Telling lies: Clues to deceit in the marketplace, politics, and marriage (revised edition)*. WW Norton & Company.
- Furnham, A. and E. Petrova (2010). *Body language in business: Decoding the signals*. Palgrave Macmillan.
- Gabielkov, M., A. Ramachandran, A. Chaintreau, and A. Legout (2016). "Social clicks: What and who gets read on Twitter?" In: *Proceedings of the 2016 ACM SIGMETRICS international conference on measurement and modeling of computer science*, 179–192.
- Gadarian, S. K., S. W. Goodman, and T. B. Pepinsky (2021). "Partisanship, health behavior, and policy attitudes in the early stages of the COVID-19 pandemic". *Plos one* 16.4, e0249596.
- Gawronski, B. and B. K. Payne (2011). *Handbook of implicit social cognition: Measurement, theory, and applications*. Guilford Press.
- Gentzkow, M., B. Kelly, and M. Taddy (2019). "Text as data". *Journal of Economic Literature* 57.3, 535–74.
- Gentzkow, M. and J. M. Shapiro (2010). "What drives media slant? Evidence from US daily newspapers". *Econometrica* 78.1, 35–71.
- (2011). "Ideological segregation online and offline". *The Quarterly Journal of Economics* 126.4, 1799–1839.
- Gentzkow, M., J. M. Shapiro, and M. Taddy (2019). "Measuring group differences in high-dimensional choices: method and application to congressional speech". *Econometrica* 87.4, 1307–1340.
- Gollwitzer, A., C. Martel, W. J. Brady, P. Pärnamets, I. G. Freedman, E. D. Knowles, and J. J. Van Bavel (2020). "Partisan differences in physical distancing are linked to health outcomes during the COVID-19 pandemic". *Nature human behaviour* 4.11, 1186–1197.
- Grabe, M. E. and E. P. Bucy (2009). *Image bite politics: News and the visual framing of elections*. Oxford University Press.
- Greene, S. and P. Resnik (2009). "More than words: Syntactic packaging and implicit sentiment". In: *Proceedings of human language technologies: The 2009 annual conference of the north american chapter of the association for computational linguistics*, 503–511.
- Groseclose, T. and J. Milyo (2005). "A measure of media bias". *The Quarterly Journal of Economics* 120.4, 1191–1237.
- Haim, M. and M. Jungblut (2021). "Politicians' self-depiction and their news portrayal: Evidence from 28 countries using visual computational analysis". *Political Communication* 38.1-2, 55–74.
- Hołownka, T. (1981). "On conventionality of signs".
- Imbens, G. W. and D. B. Rubin (2015). *Causal inference in statistics, social, and biomedical sciences*. Cambridge University Press.
- Jha, S. and C. Busso (2020). "Estimation of Driver's Gaze Region from Head Position and Orientation using Probabilistic Confidence Regions". *arXiv preprint arXiv:2012.12754*.
- Knowlton, J. Q. (1964). "A socio-and psycho-linguistic theory of pictorial communication."

- Knowlton, J. Q. (1966). "On the definition of "picture"". *AV communication review* 14.2, 157–183.
- Koliska, M. and K. Oh (2021). "Guided by the Grid: Raising Attention with the Rule of Thirds". *Journalism Practice*, 1–20.
- Langer, S. K. (2009). *Philosophy in a new key: A study in the symbolism of reason, rite, and art*. Harvard University Press.
- Lee, J., M. Muñoz, L. Fridman, T. Victor, B. Reimer, and B. Mehler (2018). "Investigating the correspondence between driver head position and glance location". *PeerJ Computer Science* 4, e146.
- Lewis, J. B., K. Poole, H. Rosenthal, A. Boche, A. Rudkin, and L. Sonnet (2021). "Voteview: Congressional Roll-Call Votes Database". See <https://voteview.com/>.
- Long, J. S. and L. H. Ervin (2000). "Using heteroscedasticity consistent standard errors in the linear regression model". *The American Statistician* 54.3, 217–224.
- Lunenburg, F. C. (2010). "Louder than words: The hidden power of nonverbal communication in the workplace". *International Journal of Scholarly Academic Intellectual Diversity* 12.1, 1–5.
- McCombs, M. and A. Reynolds (2009). "How the news shapes our civic agenda". In: *Media effects*. Routledge, 17–32.
- Morris, C. (1946). "Signs, language and behavior."
- Parks, D., A. Borji, and L. Itti (2015). "Augmented saliency model using automatic 3d head pose detection and learned gaze following in natural scenes". *Vision research* 116, 113–126.
- Peirce, C. S. (1931). "Collected Papers, Cambridge". *Harvard University Press* 1958.2, 643.
- Peng, Y. (2018). "Same candidates, different faces: Uncovering media bias in visual portrayals of presidential candidates with computer vision". *Journal of Communication* 68.5, 920–941.
- Poole, K. T. and H. Rosenthal (1985). "A spatial model for legislative roll call analysis". *American journal of political science*, 357–384.
- Prat, A. (2018). "Media power". *Journal of Political Economy* 126.4, 1747–1783.
- Prat, A. and D. Strömborg (2013). "The political economy of mass media". *Advances in economics and econometrics* 2, 135.
- Prior, M. (2013). "Media and political polarization". *Annual Review of Political Science* 16, 101–127.
- Recasens, M., C. Danescu-Niculescu-Mizil, and D. Jurafsky (2013). "Linguistic models for analyzing and detecting biased language". In: *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 1650–1659.
- Romer, D. and K. H. Jamieson (2020). "Conspiracy theories as barriers to controlling the spread of COVID-19 in the US". *Social science & medicine* 263, 113356.
- Sebeok, T. A. (1985). *Contributions to the Doctrine of Signs*. University Press of America.
- Shearer, E. (2021). "More than eight-in-ten Americans get news from digital devices". *Pew Research Center* 12.
- Strömborg, D. (2015). "Media and politics". *economics* 7.1, 173–205.
- Sutherland, C. A., J. A. Oldmeadow, I. M. Santos, J. Towler, D. M. Burt, and A. W. Young (2013). "Social inferences from faces: Ambient images generate a three-dimensional model". *Cognition* 127.1, 105–118.
- Veltrusky, J. (1976). "Some aspects of the pictorial sign". *Semiotics of Art. Ladislav*.

- Walters, S. B. (2011). *Principles of kinesic interview and interrogation*. Boca Raton, FL: CRC Press.
- Yano, T., P. Resnik, and N. A. Smith (2010). “Shedding (a thousand points of) light on biased language”. In: *Proceedings of the NAACL HLT 2010 Workshop on Creating Speech and Language Data with Amazon’s Mechanical Turk*, 152–158.

A.1 Appendix 1 (for the Analysis of Visual Partisanship)

A.1.1 Establishing gaze regions in pictures

This section describes the details of the methods used to determine the gaze regions of the subjects in a picture. I borrow this approach from studies on Intelligent Vehicle Systems, in which a driver’s head pose is used to predict the attention patterns to the road. See, among others, Parks, Borji, and Itti, 2015; Lee et al., 2018; Dari, Kadrilev, and Hullermeier, 2020; Jha and Busso, 2020).

Given two subjects in a picture, A and B, I determine subject A’s “gaze region”, and measure whether B falls in that gaze region; if so, then I consider B as seen by A. I use this measure to construct the triggered emotionality measure described in the main text. The raw data from Microsoft’s API include the measurement of the following head poses: pitch (i.e. whether the chin is up or down), yaw (i.e. the horizontal rotation of the head, towards the left or the right), and roll (i.e. the head’s inclination to the sides, namely bringing the ear closer to the shoulder).

First, I determine an area of the picture that is “compatible” with an individual’s gaze region, that is, points in the picture that are possibly seen by the individual. I approximate this region through the information on the head’s yaw and pitch. The accurate determination of a subject’s gaze region in a 2-dimensional picture presents two main challenges. First, the head’s position is expressed in degrees (yaw, pitch, roll), and the conversion of an angle to a length requires knowledge of the distance between the viewer and an object. In fact, the sight region flattened in a 2-dimensional space appears as a triangle whose base (i.e. the side most distant from the viewer) is proportional to the triangle’s “height” (namely, the distance between viewer and object). This implies that for a given angle of a visual region, its section is wider the furthest is the observer. Actual distances between subjects in a picture can hardly be measured⁴⁶ and are thus often approximated. Another problem originates in the fact that the sight angle γ between A and B can result from multiple combinations of A’s yaw and pitch, as we ignore the distance between subject and cannot exactly determine the relative contributions of a head’s pitch and yaw in producing γ . To illustrate, imagine a viewer in the center of the picture and consider the picture’s bottom-left corner: such point could be visible both if the person had yaw= -90 (i.e. her head was completely turned to the right) and pitch < 0 (i.e. looking downward), and if the person had pitch=0 (gaze at own eyes’ level) and head turned more toward the camera (e.g. yaw 45). In particular, the more distant the person from the camera, the closer to 0 could be her head’s yaw while maintaining sight of the point at the bottom-left corner of the picture. The ambiguity is once again due to the lack of knowledge of distances between subjects, and the flattening of the scene on a 2-dimensional surface.

To work around the difficulty, I determine each person’s “plausible” sight region using rather ample criteria,

⁴⁶This would require information on focal lengths and the presence of a known object whose dimension is known (e.g. a 1 euro coin).

and then imposing further requirements to increase the precision. First, I consider a margin to the left and to the right of the head’s yaw. Now, the eyes’ main focus region is 30 degrees to each side, but 30 degrees is much less than the actual natural sight region as we also have 30 more degrees of near-peripheral area. Objects in this area would be more comfortably seen by turning the head more, however pictures often capture moments in which the individuals are reacting quickly to a visual stimulus, to which eyes naturally respond before head movements. Therefore, I take an intermediate length between the focus region and the near-peripheral region, and consider a margin of 45 degrees to each side of the yaw. Then, I consider the sign of the head’s pitch, to pin down in which direction (upper or lower) to orient the area determined by the yaw.

For every observer (A) and other subject (B) in the picture, I consider B as falling within A’s sight region if both of the following conditions are verified:

1. The angle γ generated by the line connecting A and B falls within a range around A’s yaw equal to $3 * \sqrt{|yaw|}$.
2. The vertical distance between A and B (i.e. the distance in coordinates $y_b - y_a$) and A’s pitch have the same sign. Formally, the product of the two shall be non-negative: this indicates that A’s head vertical inclination (upwards or downwards) is in B’s direction. Given that for sufficiently small vertical distances or for pitches close to 0 the product may happen to be negative even if B is visible to A, I include a tolerance level considering as 0 values between -15 and +15 for both vertical distance and pitch, so to obtain a vertical vision span of 30 degrees. I also set vertical distance to 0 any distance between -0.9 and +0.9 between Y_a and Y_b .

If two or more persons fall within A’s gaze region, I consider A to be looking only at the person that is closest to her. In this sense, since in images with at least three individuals about 94% of the persons have head yaw between -45 and 45 degrees (indicating a relatively frontal head pose), I rank subjects in a person’s gaze region considering first image depth (namely distance from the camera), then breaking potential ties using horizontal distances (to the left and to the right of the viewer). I establish the relative distance of subjects from the camera using the faces’ dimensions, considering two subjects with the same face size as equally distant, and allowing for a 5% tolerance in face area differences. I then exclude from a person’s gaze region all the subjects who are behind her (and hence cannot be in sight). Finally, I exclude all subjects from the gaze regions of persons whose eyes are occluded (either covered or closed).

Having so approximated the focus of the persons’ gaze (i.e. what they “see”), I compute the triggered emotion of observed individuals as the weighted average of their observers’ emotions. The weights are proportional to the depth-distance of the observer: as stated in the previous section, I assume the picture to confer more visibility to the subjects whose features are meant to matter more.

The triggered emotionality measure rests on the assumption that glances can be used to transfer the observer’s emotion to the observed person, thus that a person’s facial expression is informative of the emotional evaluation of what she sees. The method is clearly limited in cases such as when individuals glance away from an emotionally triggering sight (instance plausibly more frequent with negative emotions). Nevertheless, it allows to go beyond

the mere emotion-labelling of single faces, and to capture the deeper emotional loading of images with multiple individuals. To limit the method's possible flaws, I only measure triggered emotions in images with up to 3 persons: this safeguards the accuracy of the method (the more people are portrayed, the higher number of possible glance-interactions and emotion attributions), while at the same time includes the vast majority of pictures in my sample.

TABLE (A.1.1)
List of US News sources

Source's twitter handle:	Bias Rating on Adfontesmedia:	Bias rating on allsides.com	Included in analysis:
AlterNet	-30.33	LL	✓
TheAtlantic	-19.66	L	✓
Salon	-19.35	LL	✓
politicususa	-16.62	LL	✓
theintercept	-16.5	LL	✓
MSNBC	-13.76	LL	✓
Newsweek	-12.96	L	✓
CNN	-12.15	LL	✓
voxdotcom	-11.93	LL	✓
GuardianUS	-10.35	L	✓
TIME	-10.22	L	✓
thehill	0.1	R	✓
PittsburghPG	0.1	R	✓
WSJ	4.95	R	✓
TPIInsidr	7.67	R	✓
RealClearNews	13.07	R	✓
nypost	14.2	RR	✓
FreeBeacon	15.9	RR	✓
WashTimes	16.12	R	✓
FoxNews	17.19	R	✓
realDailyWire	18.63	RR	✓
BreitbartNews	25.67	RR	✓

Notes: The Table presents the list of the main News sources extrapolated from *Similarweb.com*. The partisanship scores from Adfontes media and Allsides.com are listed in the respective columns (for Adfontesmedia, positive numbers indicate relatively Republican leaning).

TABLE (A.1.2)
Topic-reduction scheme: part1

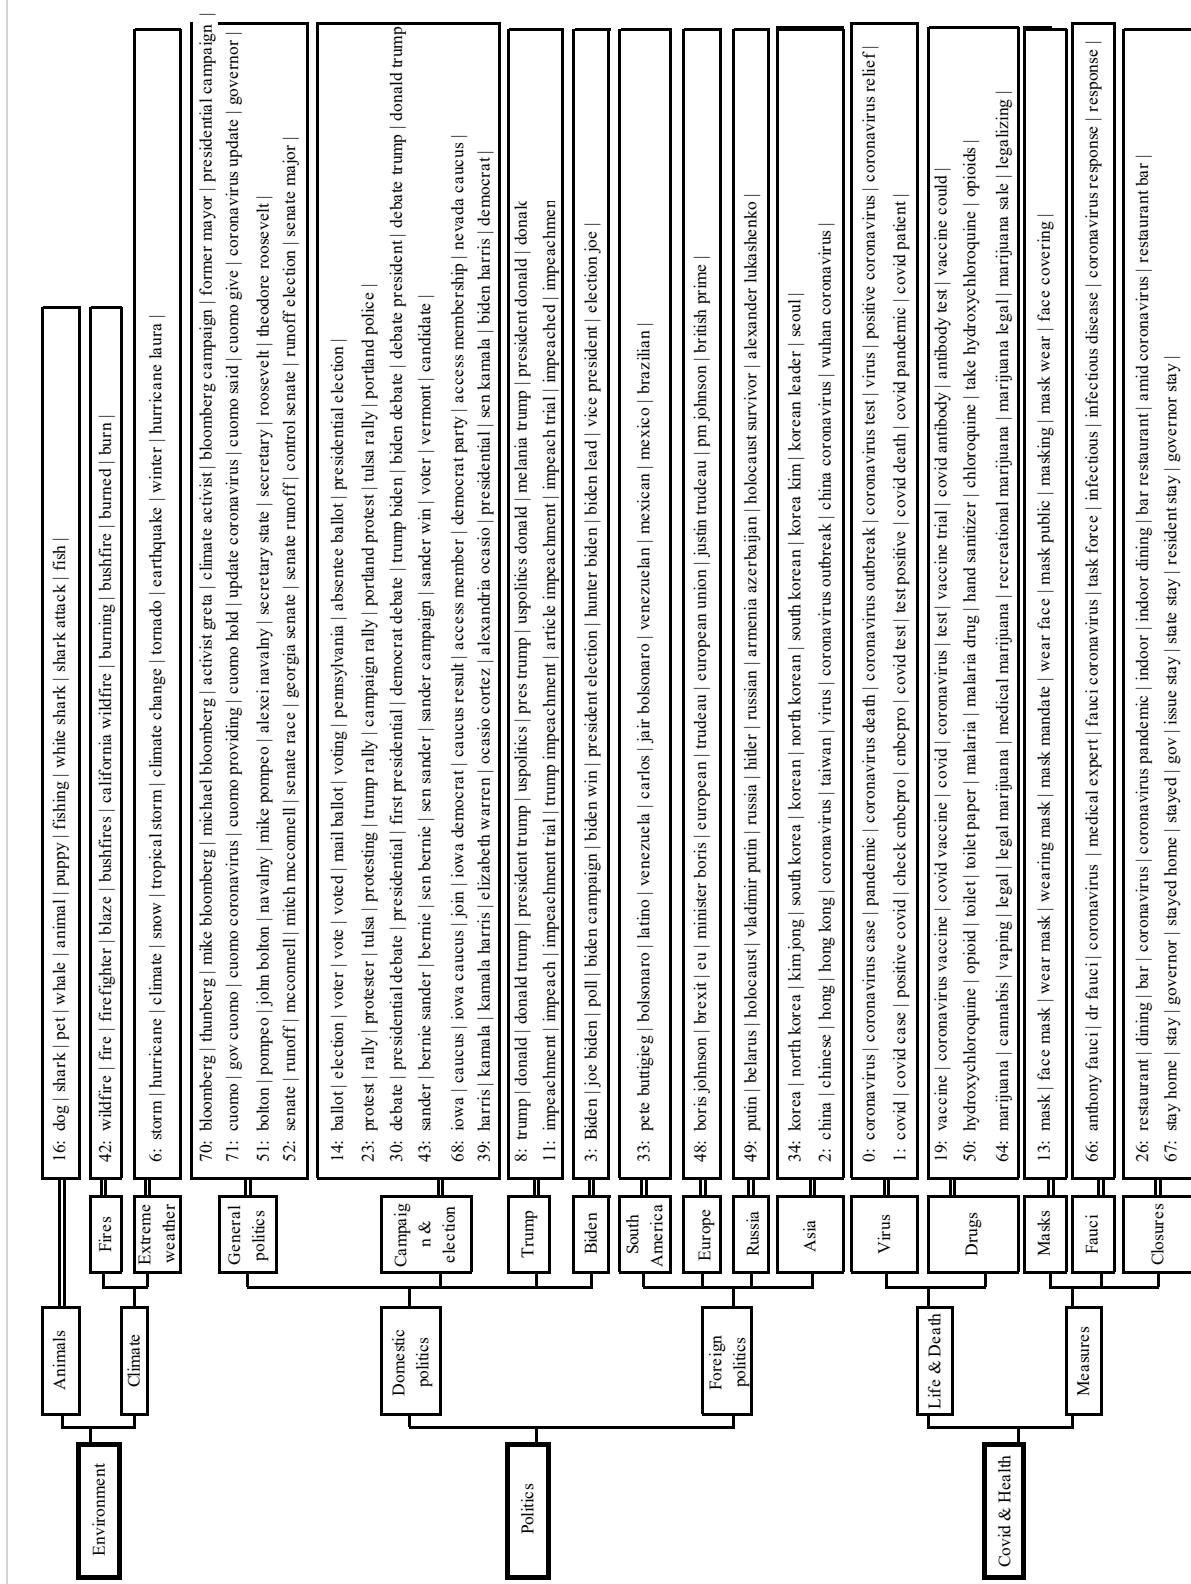


TABLE (A.1.3)
Topic-reduction scheme: part 2



TABLE (A.1.4)
Topic-reducion scheme: part 3

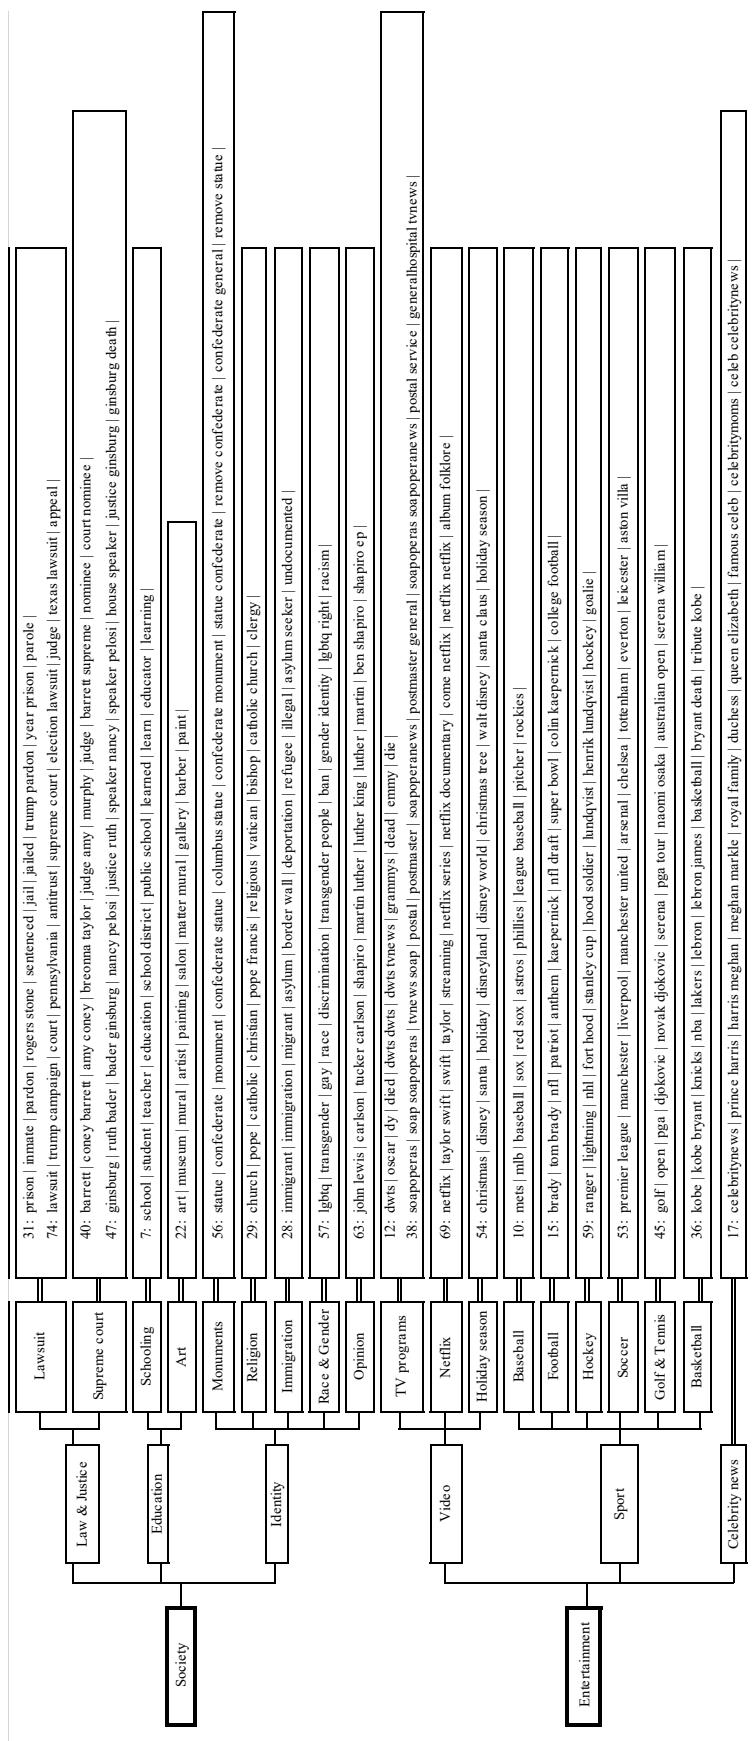


TABLE (A.1.5)
VOCABULARY SUMMARY STATISTICS: “SUBJECT” FEATURES AND FEATURE-COMBINATIONS INVOLVING “SUBJECTS”

“Subjects” subclasses:	Label in syntax index	N distinct tokens in subclass	Share of total “Subject” tokens	Total “Subject” tokens	Total presence in pictures
<i>Single features:</i>					
Celebrity status (y/n)	“SCele”	2	0.00044	4'541	256'677
Name	“SNam”	4'518	0.995	4'541	94'016
Rank in terms of centrality	“SRnk”	10	0.0022	4'541	255'947
Political leaning decile	“SPoL”	9	0.00198	4'541	50'399
Gender	“SGen”	2	0.00044	4'541	256'677
<i>Combinations of features:</i>					
Celebrity status (y/n)	“SCele & ...”	1266	0.00103	1'225'751	8'093'764
Name	“SNam & ...”	1'216'337	0.992	1'225'751	4'233'008
Rank in terms of centrality	“SRnk & ...”	26143	0.0213	1'225'751	8'192'869
Political leaning decile	“SPoL & ...”	3811	0.00311	1'225'751	1'563'848
Gender	“SGen & ...”	1182	0.000964	1'225'751	7'993'290

Notes: The table provides summary statistics for the “Subjects” syntax class. The upper panel provides summary statistics for tokens consisting of single features only (excluding combinations). The bottom panel provides summary statistics for tokens consisting of features combinations. Note: all statistics encompass only tokens that appear at least once in analysed images. Column 1 lists the subclasses in the “Subjects” class, marked in the data using the labels in Column 2. Column 3 indicates the number of distinct tokens for each subclass (denoting the different values a subclass can take). Column 4 indicates the percentage of tokens in the subclass out of all “Subject” tokens, whose total is listed in Column 5. Finally, Column 6 lists the total occurrence of subclass tokens in images in the dataset.

TABLE (A.1.6)
VOCABULARY SUMMARY STATISTICS: “ADJECTIVE” FEATURES AND FEATURE-COMBINATIONS INVOLVING “ADJECTIVES”

“Adjectives” subclasses:	Label in syntax index	N distinct tokens in subclass	Share of total “Adjective” tokens	Total “Adjective” tokens	Total presence in pictures
<i>Single features:</i>					
Centrality	“ACent”	4	.1	40	255'947
Face size	“ASize”	3	.075	40	256'677
Use of Face Mask	“AKmsk”	2	.05	40	252'714
Blur level	“AKblr”	3	.075	40	256'677
Light exposure	“AKexp”	3	.075	40	256'677
Head yaw	“AKyaw”	3	.075	40	256'677
Head pitch	“AKpit”	3	.075	40	256'677
Facial emotion	“AKem”	9	.225	40	260'144
Triggered emotion	“AKtrem”	3	.075	40	54'437
Observed by # people	“AKseen”	7	.175	40	120'487
<i>Combinations of features:</i>					
Centrality	“ACent”	449'576	.597	753'168	13'538'520
Face size	“ASize”	435'916	.579	753'168	13'553'996
Use of Face Mask	“AKmsk”	88'548	.118	753'168	3'666'172
Blur level	“AKblr”	117523	.156	753'168	3'716'344
Light exposure	“AKexp”	99'681	.132	753'168	3'716'344
Head yaw	“AKyaw”	108'401	.144	753'168	3'716'344
Head pitch	“AKpit”	101'968	.135	753'168	3'716'344
Facial emotion	“AKem”	120'013	.159	753'168	3'963'644
Triggered emotion	“AKtrem”	35'984	.0478	753'168	774'544
Observed by # people	“AKseen”	8'193	.0109	753'168	121'912

Notes: The table provides summary statistics for the “Adjectives” syntax class. The upper panel provides summary statistics for tokens consisting of single features only (excluding combinations). The bottom panel provides summary statistics for tokens consisting of features combinations. Note: all statistics encompass only tokens that appear at least once in analysed images. Column 1 lists the subclasses in the “Adjectives” class, marked in the data using the labels in Column 2. Column 3 indicates the number of distinct tokens for each subclass (denoting the different values a subclass can take). Column 4 indicates the percentage of tokens in the subclass out of all “Adjective” tokens, whose total is listed in Column 5. Finally, Column 6 lists the total occurrence of subclass tokens in images in the dataset.

TABLE (A.1.7)
VOCABULARY SUMMARY STATISTICS: “CONTEXT” FEATURES AND FEATURE-COMBINATIONS INVOLVING “CONTEXT”

“Context” subclasses:	Label in syntax index	N distinct tokens in subclass	Share of total “Context” tokens	Total “Context” tokens	Total presence in pictures
<i>Single features:</i>					
Tags and tags combinations	“CNtagmix”	487'925	.948	514'488	4'870'758
Context elements	“CNtxt”	26'563	.0516	514'488	1'258'357
<i>Combinations of features:</i>					
Tags and tags combinations	“CNtagmix”	1'885'382	1	1'885'382	16'973'802
Context elements	“CNtxt”	1'435'666	.761	1'885'382	15'951'254

Notes: The table provides summary statistics for the “Context” syntax class. The upper panel provides summary statistics for tokens consisting of single features only (excluding combinations). The bottom panel provides summary statistics for tokens consisting of features combinations. Note: all statistics encompass only tokens that appear at least once in analysed images. Column 1 lists the subclasses in the “Context” class, marked in the data using the labels in Column 2. Column 3 indicates the number of distinct tokens for each subclass (denoting the different values a subclass can take). Column 4 indicates the percentage of tokens in the subclass out of all “Context” tokens, whose total is listed in Column 5. Finally, Column 6 lists the total occurrence of subclass tokens in images in the dataset.

A.1.2 Visual tokens in the Vocabulary

This section summarizes the structure of the visual vocabulary by listing the individual features (not combinations) contained in each syntax class and subclass (categories are in bold, token names in italic).

Class S: **Subject** identity traits (Features constant across portrayals of the same individual, but varying across different individuals).

- **Name (SN)** (es. *“Kate Blanchett”*, etc.: a token per each known person ever portrayed);
- **Celebrity Status (SC)** (token indicating a person is a celebrity);
- **Sex (SG)** (*“Male”* or *“Female”*);
- **Party bin (SP)** (9 tokens for politicians’ partisanship decile; data by Poole and Rosenthal, 1985);
- **Face Salience Rank (SR)** (i.e. *“Person 1”*, *“Person 2”*, ..., *“Person 10”*) Face unique identifier within a picture. The number is a rank based on a person’s relative salience within the image, measured as a weighted average of face size and centrality (weighting respectively 70% and 30%); rank=1 indicates the most salient person in the picture.

Class A: **Adjectives**, modality of a person’s representation (*Features that vary across individuals and across representations of a given individual*).

- **Face Size (AS)** (*“Large”*: face area share $F > 1/6$ image; *“Medium”*: $F \in [1/6, 1/24]$; *“Small”*: $F < 1/24$);
- **Face Centrality (AC)** (*“Very High”*: centrality $C \in [.95, 1]$; *“M-High”*: $C \in [.85, .95]$; *“M-Low”*: $C \in [.75, .85]$; *“Very Low”*: $C < .75$);
- **Facial emotion (AKem)** (*“Anger”*; *“Contempt”*; *“Disgust”*; *“Fear”*; *“Happiness”*; *“Sadness”*; *“Surprise”*);

- **Emotion triggered in portrayed observers (AKtrem)** (Mean emotion: “Positive” if happiness; “Neutral” if no prominent emotion among observers; “Negative” if Anger, Contempt, Disgust, Fear, or Sadness);
- **Head pitch (AKpit)** (“Negative”: pitch $P < -15^\circ$; “Neutral”: $P \in [-15^\circ, +15^\circ]$; “Positive”: $P > 15^\circ$);
- **Head yaw (AKyaw)** (“Right profile”: yaw $Y < -30^\circ$; “Frontal”: $Y \in [-30^\circ, +30^\circ]$; “Left profile”: $Y > 30^\circ$);
- **Mask (AKmsk)** (Indicator for person wearing a mask);
- **Face blur level (AKblr)** (“High”, “Medium”, “Low”);
- **Face light exposure (AKexp)** (“Overexposed”- bright, “Regular exposure”, “Underexposed”- dark);
- **Number of observers (AKseen)** (Indicator for people observing a person summing to 1-9);

Class C (describing **Context** attributes):

- **Within class C, General image descriptors focused on persons (CNtxt):**
 - **Presence of persons, celebrities and congresspeople** (11 indicators for total number of persons, from none to 10; *Celebrity*: presence of at least one well-known person; *Congresspeople*: presence of at least one congress member; *People but no celebrity*: presence of people but no celebrities; *Celebrity but no congressperson*: presence of celebrities but no congresspeople);
 - **Couples of names** (es. *Donald Trump & Kate Blanchett*: a token per each pair of well-known persons ever portrayed jointly);
 - **Triplets of names** (es. *Donald Trump & Kate Blanchett & Johnny Cash*: a token per each triplet of well-known persons ever portrayed jointly);
 - **Men and women representation patterns** (Tokens indicating the presence of: *men*; *women*; *men only*; *women only*; *majority of men*; *majority of women*; *same number of women and men*);
 - **Republicans, Democrats, and Independents representation patterns** (Tokens indicating the presence of: *Democrats*; *Republicans*; *Independents*; *Dems only*; *Reps only*; *Majority of Dems*; *Majority of Reps*; *Same number of Dems and Reps*);
 - **Mask wearing patterns** (indicators for image portraying wearing a mask: *At least one person*, *All the persons*; *The majority of people*; *The minority of people*);
 - **Facial emotion patterns** (indicators for average emotion: *Positive*; *Neutral*; *negative*);
 - **Image shot angle** (indicators for shot angle: *From above*, *From below*, *From front*; it is derived from average camera angle of person’s face portrayals);
- **Within class C, General image descriptors from image tags, and tags mix (CNtagmix):**
 - **Tags for Person** (tags characterizing the persons portrayed in the picture, e.g. “*Policeman*”);

- **Tags for Animals** (tags indicating the presence of animals or an animal group, e.g. “*Flock*”);
- **Tags for Things** (tags indicating the presence of specific objects, e.g. “*Knife*”);
- **Tags for Verbs** (tags indicating an action a subject in the picture is executing, e.g. “*Walking*”);
- **Tag for Qualifiers** (tags for characteristic of portrayed subjects/objects, e.g. “*Yellow*”);
- **Tag for Bodyparts** (tags for clearly visible body parts of humans or animals, e.g. “*Waist*”);
- **Tag for Places** (tags indicating in which place the picture is taken, e.g. a “*Station*”)
- **Tag: Background** (Single background tags, couple and triplets; Background tags indicate items in the image background, e.g. “*Tower*”);
- **Tag: Setting** (single setting tags, and in combination with background tags; setting tags refer to the pictures setting, situation or theme, e.g. “*Volleyball*”);
- **Tag mix: Thing(wearable) + Thing(wearable)**
- **Tag mix: Thing(wearable) + Thing(wearable) + Thing(wearable)**
- **Tag mix: Person + Thing(wearable)**
- **Tag mix: Person + Thing(wearable) + Thing(wearable)**
- **Tag mix: Verb(drinking) + Thing(beverage)**
- **Tag mix: Person + Verb(transitive;object compatible:animals) + Animal**
- **Tag mix: Person + Verb(transitive; object compatible:animals) + Animal**
- **Tag mix: Person + Verb(transitive; object compatible:things) + Thing**
- **Tag mix: Person + Verb(intransitive; subject compatible: person)**
- **Tag mix: Animal + Verb(subject compatible: animal)**
- **Tag mix: Thing + Qualifier(color)**
- **Tag mix: Thing + Qualifier(material)**
- **Tag mix: Person + Verb(subject compatible: person; object compatible: transports) + Thing(transportation mean)**
- **Tag mix: Thing + Thing**
- **Tag mix: Thing + Thing + Thing**
- **Tag mix: Place + Background + Background + Background**
- **Tag mix: Background + Background**
- **Tag mix: Background + Background + Background**
- **Tag mix: Setting + Setting**
- **Tag mix: Setting + Setting + Setting**
- **Tag mix: Setting(event) + Place**

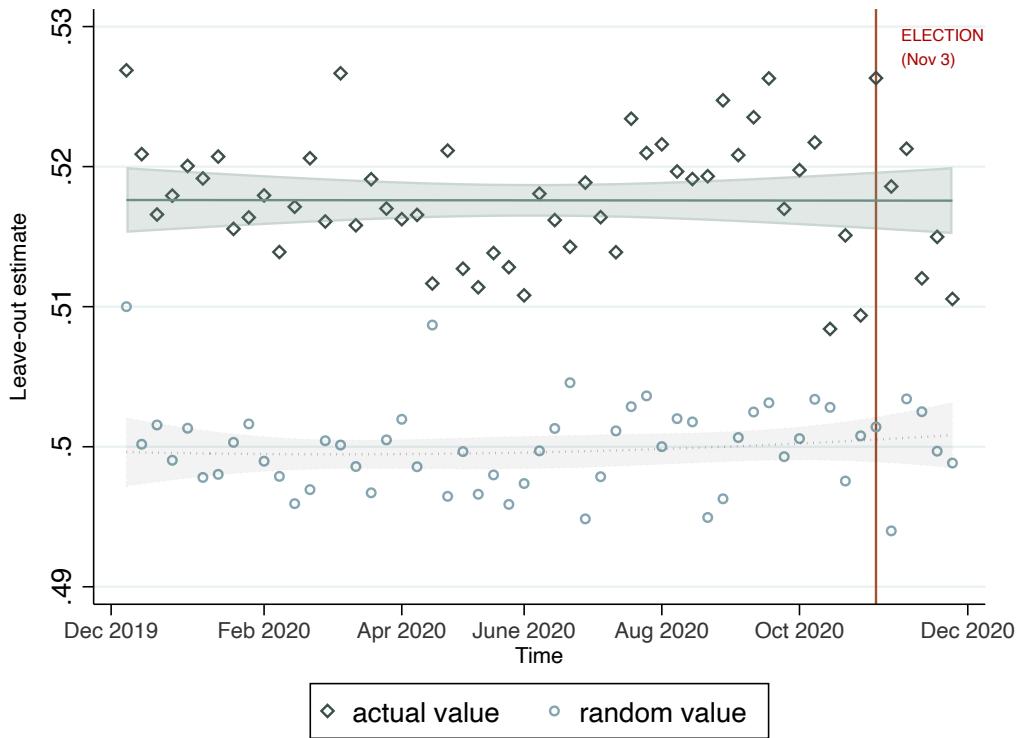


FIGURE (A.1.1)
Visual polarization estimates, weekly frequency

Notes: The Figure shows the partisanship of leading images estimated through the leave-out estimator, with weekly frequency. The shaded region represents the 95% confidence interval of the linear regression fit to the actual values. I quantify noise by calculating the leave-out estimates after randomly assigning images to parties: the values resulting from random assignment do not systematically depart from .5, suggesting that the actual values are not a result of noise.

A.2 Appendix 2: Survey Experiment

A.2.1 News Issues, Leading text, and Leading images

This Section illustrates the choice of treatments' news issues, texts, and pictures. As illustrated in Section II.C.4, five news topics have comparable and significant visual partisanship: Politics, Covid & Health, Economy, Security, and Society. To identify news issues pertaining each of these topics, I rely on the list of relevant news issues drafted by [allsides.com](https://www.allsides.com), which tags issues by their topic. This website compares news issues from sources with opposite slants, and it periodically publishes "Headline Roundups" (syntheses of the main news issues within a given period) to highlight the different takes of the Democrat-leaning and Republican-leaning news sources.⁴⁷ I use these roundups to identify valid issues within each news topic (as listed in the following paragraphs) and to ensure the coherence between treatments' framing and actual media coverage. Based on the roundups, for each news issue I draft a headline, a byline, and a leading text coherent with the neutral tone of non-partisan coverage (i.e. that of news sources rated as "Centre", neither Democrat- nor Republican- leaning, on [allsides.com](https://www.allsides.com)).

Having identified the treatment news issues, I proceed to select three treatment leading images (Dem-leaning, Rep-leaning and neutral) for each issue. A fundamental aspect to consider in the choice of images is that the partisanship scores determined through the method in Section II pertained to a period preceding that of the experiment (Dec 2019- Dec 2020, vs. July 2021). Some notable events intervened in between, such as the switch from a Republican to a Democrat US presidency. This implies that the partisan news outlets' narratives (including the visual one) in use during the experiment could potentially differ from the ones tested in Section II.⁴⁸ I therefore adopt a two-steps process to select experimental pictures that can be both "partisan" according to the method introduced in Section II and aligned with the partisan visual narratives in place at the time of the experiment (2021). I first identify a set of leading pictures actually published in news pieces on the same issue and rated as "strongly Democratic", "Strongly Republican", or "center" on [allsides.com](https://www.allsides.com). Then, I select the ones whose visual features have "partisan loading" in the event-topic vocabularies.

⁴⁷ Roundups are available at: <https://www.allsides.com/story/admin>.

⁴⁸ For instance, Republican outlets previously framing the governments' economic policy through positive affect could later have adopted the opposite stance (to criticize the new government), and the same for Democrat-leaning outlets.

A.2.1.1 Topic: ECONOMY.

Issue: FED's forecasts on inflation.

Headlines Roundup:⁴⁹

"Amid rising consumer prices and debates over large federal spending bills, the government's role in fuelling or relieving Americans' economic burdens is a subject of debate. Many Republicans say current Democratic policies and high-cost spending bills will make current inflation worsen in coming years, hurting consumers, workers and families. The Federal Reserve maintains that current inflation will only be temporary, a stance that President Joe Biden and other prominent Democrats have echoed while advocating for spending packages they say will better the lives of average Americans. Right-rated voices have covered inflation fears more prominently, with some accusing Democrats of dismissing inflation fears while supporting harmful economic policy. Left- and center-rated voices have been less accusatory, often exploring the likelihood of inflation worsening and financially-sustainable legislation being agreed upon in Congress."

Treatments:

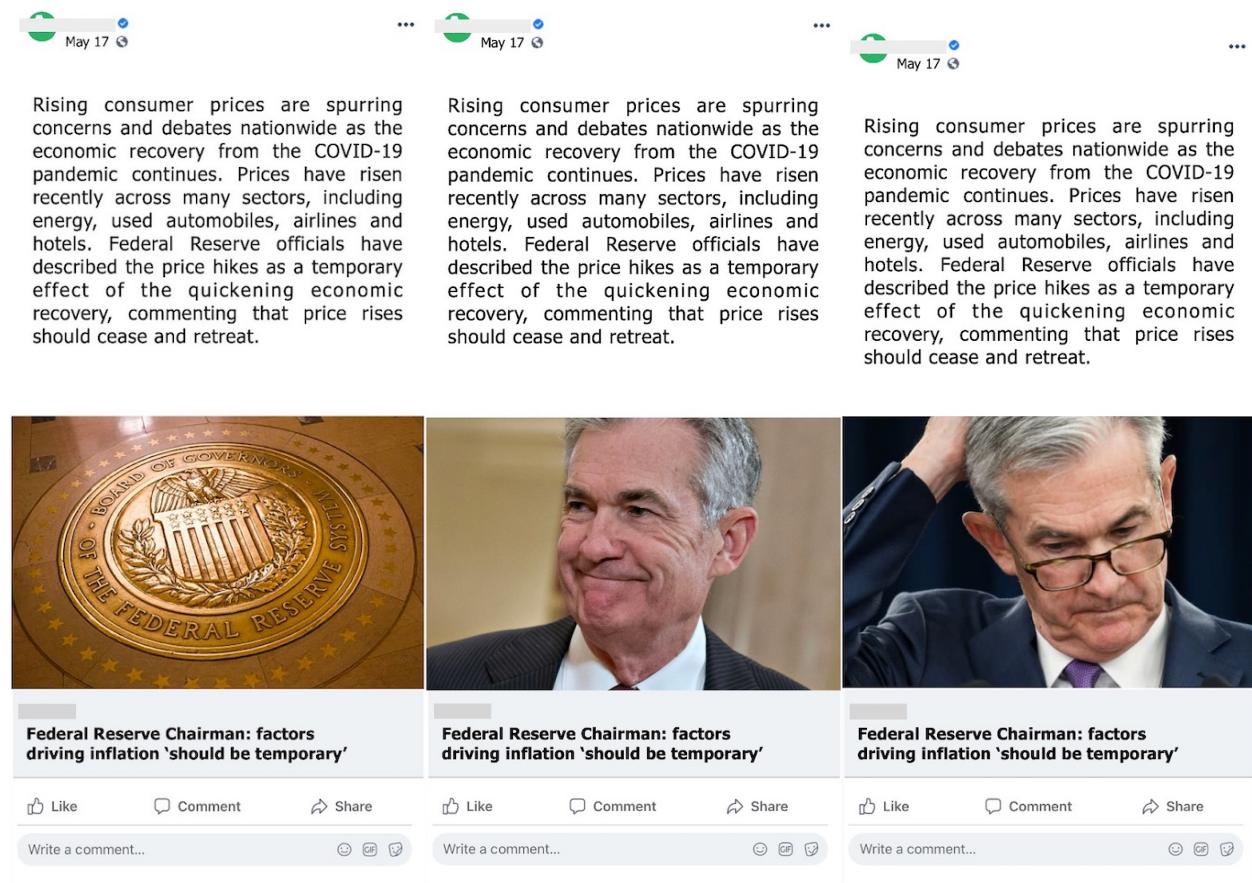


FIGURE (A.2.1)

Treatments for "Inflation" issue:
Non partisan, Dem-leaning, and Rep-leaning images.

The Figure shows the treatments (news previews) for the news issue "Inflation", related to the Economy topic. Panel A (left) shows the treatment with non partisan leading image. Panel B (centre) shows the treatment featuring the Democrat-leaning leading image. Panel C (right) shows the treatment featuring the Republican-leaning leading image.

⁴⁹From Allsides.com's Headline Roundup "*The Politicization of Inflation*", available at: <https://www.allsides.com/story/perspectives-politicization-inflation>

A.2.1.2 Topic: COVID & HEALTH.

Issue: The effectiveness of anti-Covid measures.

Headlines Roundup:⁵⁰

"Opinions range far and wide on the Trump administration's response to the COVID-19 outbreak. Many voices, particularly on the left, criticized the U.S. and White House responses. Others, especially on the right, tended to focus more on China's response to the virus as being worthy of stricter scrutiny. Some minimized the role that the administration was playing, focusing instead on other key actors and decisions."

Other:

"Partisans who harbor high levels of animus towards the other party do not differentiate the United States response to COVID-19 from that of the Trump administration."

(Druckman et al., 2020)

Treatments:

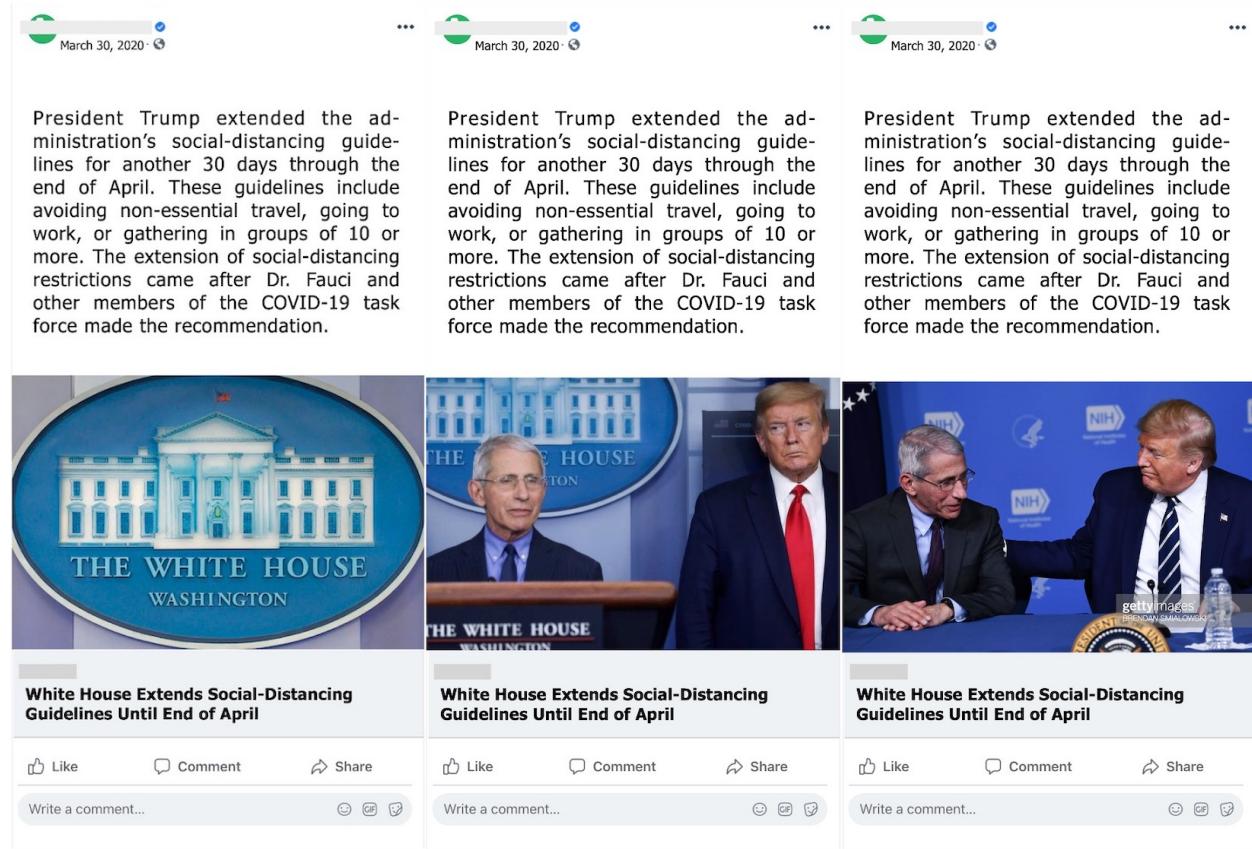


FIGURE (A.2.2)
Treatments for “Covid & Health” issue:
Non partisan, Dem-leaning, and Rep-leaning images.

The Figure shows the treatments (news previews) for the news issue “Covid measures”, related to the Covid & Health topic. Panel A (left) shows the treatment with non partisan leading image. Panel B (centre) shows the treatment featuring the Democrat-leaning leading image. Panel C (right) shows the treatment featuring the Republican-leaning leading image.

⁵⁰From Allsides.com’s Headline Roundup “Trump and the Politics of Coronavirus”, available at: <https://www.allsides.com/story/opinions-trump-and-politics-coronavirus>

A.2.1.3 Topic: POLITICS.

Issue: Renewal of the US-Iran nuclear deal.

Headlines Roundup:⁵¹

"U.S. President Joe Biden is intent on restoring the 2015 nuclear agreement with Iran, and with talks resuming in Vienna on Thursday after a weeklong break, his chief negotiator, Robert Malley, is beginning to develop a road map on how to get there. According to sources close to European and U.S. negotiators, Malley is expected to offer Tehran a Goldilocks-style deal: just enough sanctions relief so Iran will return to the pact but not so much that it would leave Biden vulnerable to attacks from hard-liners at home "

Treatments:

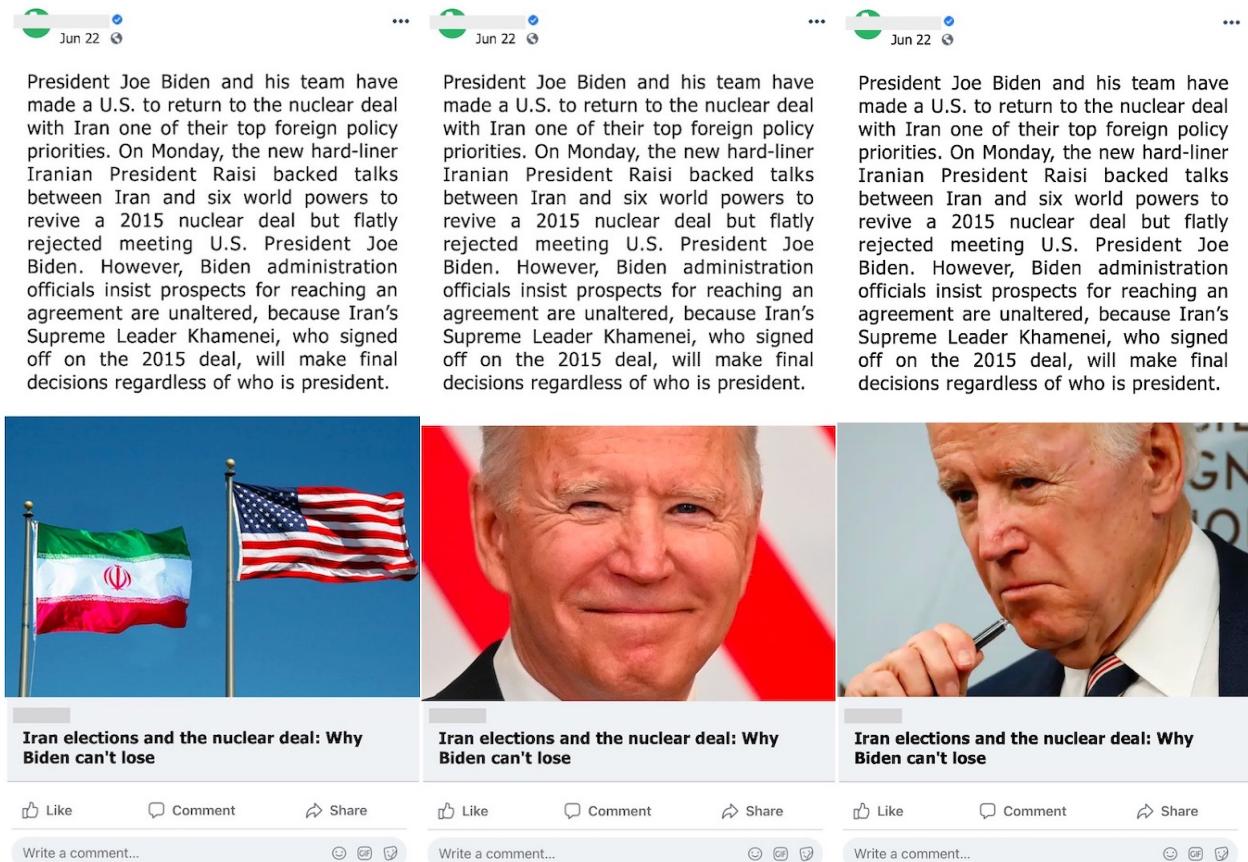


FIGURE (A.2.3)

Treatments for "Iran deal" issue:

Non partisan, Dem-leaning, and Rep-leaning images.

The Figure shows the treatments (news previews) for the news issue "Iran deal", related to the Politics topic. Panel A (left) shows the treatment with non partisan leading image. Panel B (centre) shows the treatment featuring the Democrat-leaning leading image. Panel C (right) shows the treatment featuring the Republican-leaning leading image.

⁵¹From Allsides.com's "U.S. Mounts All-Out Effort to Save Iran Nuclear Deal", available at: <https://www.allsides.com/news/2021-04-15-1349/us-mounts-all-out-effort-save-iran-nuclear-deal>

A.2.1.4 Topic: SOCIETY.

Issue: Juneteenth becomes a Federal holiday.

Headlines Roundup:⁵²

"Most of the opinions about Juneteenth this year were framed around the day becoming an official holiday. Opinions were more common from left- and center-rated outlets. Many left-rated voices celebrated the decision; many also called it a "hollow victory" and grouped it with other "symbolic gestures that are presented as progress without any accompanying economic or structural change." Some right-rated voices criticized that narrative and its proponents, arguing that "there is no concession or show of good faith that will ever placate their ever-increasing litany of demands."

News outlets from the right tended to lead the news on the bipartisan bill by portraying the Republican co-sponsor (John Cornyn), while Democrat-leaning ones tended to portray the Democrat co-sponsor (Ed Markey).

Treatments:

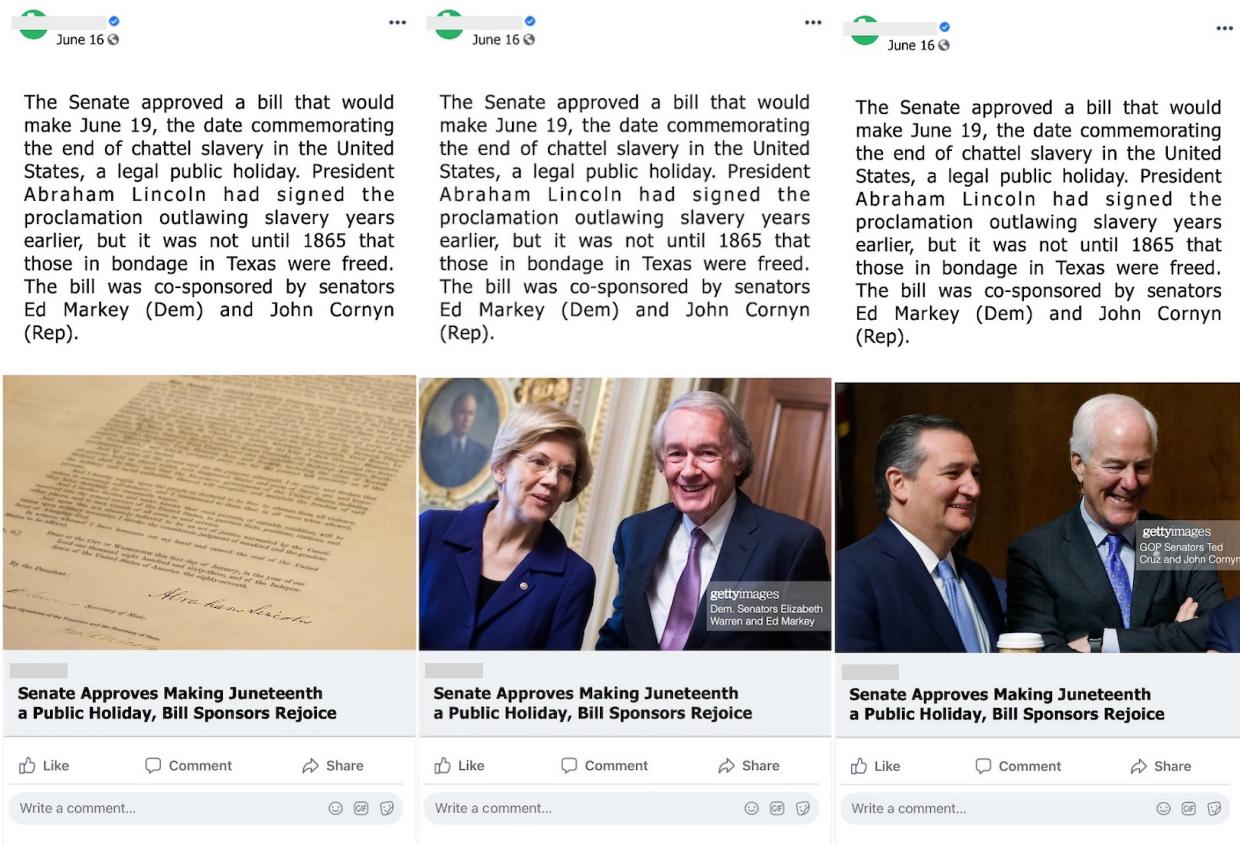


FIGURE (A.2.4)
Treatments for “Juneteenth” issue:
Non partisan, Dem-leaning, and Rep-leaning images.

Notes: The Figure shows the treatments (news previews) for the news issue “Juneteenth”, related to the Society topic. Panel A (left) shows the treatment with non partisan leading image. Panel B (centre) shows the treatment featuring the Democrat-leaning leading image. Panel C (right) shows the treatment featuring the Republican-leaning leading image.

⁵²From Allsides.com’s “Juneteenth 2021”, available at: <https://www.allsides.com/story/perspectives-juneteenth-2021>

A.2.1.5 Topic: SECURITY.

Issue: Police budget cuts.

Headlines Roundup:⁵³

"Some left-rated voices advocated for addressing systemic issues and reforming communities by reallocating significant funds from law enforcement to housing and education budgets. Several also called for an end to mass incarceration, police militarization, and police in schools. Some voices from the right argued that police systems should remain intact, pointing to possible correlations between cities with progressive law enforcement policies and rising crime rates. Many voices from all sides of the spectrum advocated for some form of police reform or reduced funding."

Treatments:

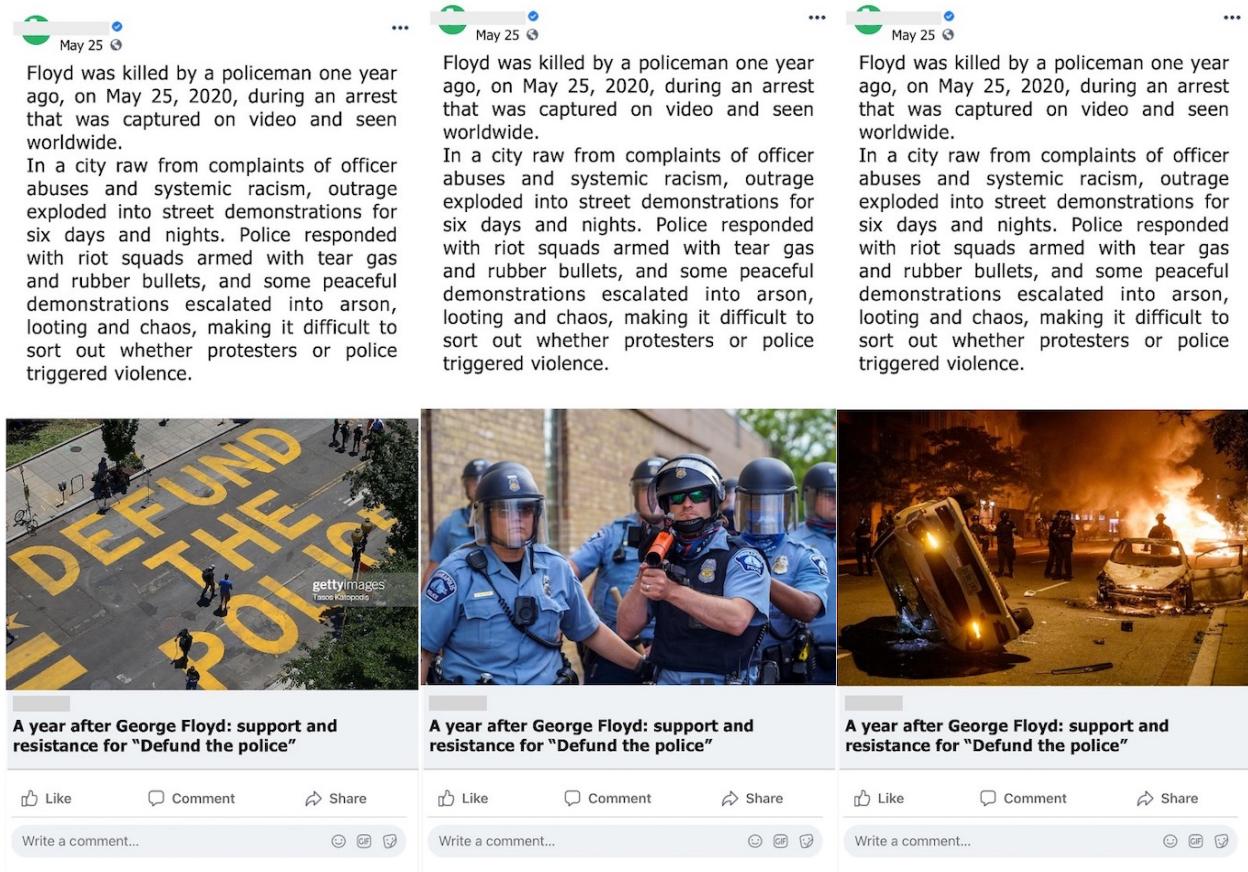


FIGURE (A.2.5)
Treatments for "Police defund" issue:
Non partisan, Dem-leaning, and Rep-leaning images.

Notes: The Figure shows the treatments (news previews) for the news issue "Police defund", related to the topic Security. Panel A (left) shows the treatment with non partisan leading image. Panel B (centre) shows the treatment featuring the Democrat-leaning leading image. Panel C (right) shows the treatment featuring the Republican-leaning leading image.

⁵³From Allsides.com's "Defunding the Police", available at: <https://www.allsides.com/story/perspectives-defunding-police>

TABLE (A.2.1)
Survey Experiment Summary Statistics

		Mean	Sd	Min	Max
Age bracket:	- 18-34	.2273185	.4192057	0	1
	- 35-44	.2565524	.4368403	0	1
	- 45-54	.2525202	.4345675	0	1
	- 55-65	.2636089	.4407007	0	1
Ethnicity:	- Caucasian	.8089718	.3932103	0	1
	- African-American	.0927419	.2901436	0	1
	- Latin American	.0645161	.245732	0	1
	- Asiatic	.0579637	.2337337	0	1
	- Native American	.015625	.1240509	0	1
Schooling < 8 yrs.		.0095766	.097415	0	1
Party affiliation:	- Democrat	.3886089	.487557	0	1
	- Independent	.3069556	.4613471	0	1
	- Republican	.3044355	.4602839	0	1
Politics interest:	- Very low	.0922379	.2894344	0	1
	- Low	.1673387	.3733721	0	1
	- Medium	.3447581	.4754091	0	1
	- High	.2620968	.4398859	0	1
	- Very high	.1335685	.3402739	0	1
Political opinion	(Liberal/Conservative)	4.048387	1.723639	1	7
Gets news from:	- Fox News	1.144153	1.151808	0	3
	- CNN	1.316028	1.143477	0	3
	- Breitbart	.3513105	.7434829	0	3
	- NYT	1.012097	1.058479	0	3
	- MSNBC	1.020665	1.041294	0	3
	- NYPost	.7923387	.9480834	0	3
Main info. source:	- Newspapers	.1789315	.3833915	0	1
	- Radio	.0453629	.2081513	0	1
	- Socials	.1355847	.3424333	0	1
	- TV	.5146169	.4999123	0	1
Clicks in introduction		1.628024	1.339399	1	28
Low screen resolution		.2011089	.4009303	0	1
<i>Defund Police</i>	(Ex ante opinion)	.1334203	44.36863	-100	100
"	(Post treatment opinion)	1.272364	44.13409	-100	100
<i>Iran deal</i>	(Ex ante opinion)	-3.100806	28.06095	-50	50
"	(Post treatment opinion)	-1.830141	27.76713	-50	50
<i>Inflation</i>	(Ex ante opinion)	-2.071069	28.96654	-50	50
"	(Post treatment opinion)	-.4188508	28.7	-50	50
<i>Covid measures</i>	(Ex ante opinion)	4.081653	33.6061	-50	50
"	(Post treatment opinion)	6.431452	32.87022	-50	50
<i>Juneteenth</i>	(Ex ante opinion)	5.163306	37.67628	-50	50
"	(Post treatment opinion)	7.272177	37.56104	-50	50
<i>Defund police</i> issue has low salience		.2520161	.4342799	0	1
<i>Iran deal</i> issue has low salience		.2535282	.4351403	0	1
<i>Inflation</i> issue has low salience		.2510081	.4337025	0	1
<i>Covid measures</i> issue has low salience		.2530242	.4348543	0	1
<i>Juneteenth</i> issue has low subjective salience		.2681452	.4431053	0	1
Familiarity with <i>Defund police</i> issue		2.389113	.7378551	0	3
Familiarity with <i>Inflation</i> issue		1.537802	.9781178	0	3
Familiarity with <i>Iran deal</i> issue		1.307964	1.000462	0	3
Familiarity with <i>Covid measures</i> issue		2.160786	.8902761	0	3
Familiarity with <i>Juneteenth</i> issue		1.995968	.8743466	0	3
Observations		1984			

TABLE (A.2.2)
Impact of Leading Images On News-Readers' Opinion
(Specification without controls)

Dependent variable:	(1)	(2)	(3)	(4)	(5)
	Opinion on “Defund Police”	Opinion on “Iran deal”	Opinion on “Inflation”	Opinion on “Covid measures”	Opinion on “Juneteenth”
	(Budget cut in -100 +100)	(Confidence, in -50+50)	(Confidence, in -50+50)	(Dissatisfaction, in -50+50)	(Policy support, in -50+50)
Neutral images (N)	-0.517 (0.672)	-0.541 (0.769)	0.903 (0.486)	-7.475 (0.317)	8.740 (0.653)
Democrat images (D-N)	1.379 (0.465) [0.059]	-0.257 (1.246) [0.850]	-0.494 (1.359) [0.740]	0.866 (0.314) [0.070]	-0.737 (1.109) [0.554]
Republican images (R-N)	-1.516 (0.519) [0.06]	-1.832 (1.180) [0.22]	-2.239 (0.804) [0.07]	-0.159 (0.631) [0.82]	-0.426 (0.448) [0.41]
Democrat-Republican (D-R)	2.895 (0.422) [0.006]	1.575 (0.845) [0.159]	1.745 (1.014) [0.184]	1.025 (0.799) [0.290]	-0.311 (1.461) [0.845]
Observations	1574	1608	1625	1595	1551
Controls:	N	N	N	N	N

Notes: The Table presents OLS estimates of the effect of the Democrat-leaning (D), neutral (N), and Republican-leaning (R) news-leading images on respondents' opinion after exposure to the news (column headers indicate the relevant news issue). The dependent variable for the “Defund Police” issue ranges in [-100,+100], while all others range in [-50+50]. Variables are adjusted so that the highest value in the range always corresponds to Democrats' ideological position (hence positive coefficients indicate a pro-Democratic opinion shift, and vice versa). The specifications only control for the opinion expressed on the issue prior to treatment exposure, and no other covariates. Round parentheses contain robust standard errors; square brackets contain the p-values for two-sided tests of equality between coefficients (tested pairs indicated on the left) using robust standard errors.

TABLE (A.2.3)
Balance of observable characteristics across treatment branches, “Defund police” news issue

Variables:	Republican		Neutral		Democrat		Normalized difference:		
	Mean	St. err.	Mean	St. err.	Mean	St. err.	(R-N)	(N-D)	(D-R)
Age bracket:									
– 18-34	0.006	(0.019)	-0.009	(0.018)	0.003	(0.018)	0.037	-0.028	-0.009
– 35-44	0.007	(0.020)	0.015	(0.020)	-0.021	(0.018)	-0.019	0.084	-0.064
– 45-54	-0.003	(0.019)	-0.010	(0.019)	0.013	(0.019)	0.014	-0.051	0.037
– 55-65	-0.010	(0.019)	0.003	(0.020)	0.006	(0.019)	-0.029	-0.006	0.035
Ethnicity:									
– Caucasian	0.011	(0.017)	-0.032	(0.018)	0.021	(0.016)	0.107	-0.132	0.025
– African-American	-0.008	(0.013)	0.028	(0.014)	-0.019	(0.012)	-0.117	0.157	-0.041
– Latin American	0.005	(0.011)	0.007	(0.011)	-0.013	(0.010)	-0.008	0.084	-0.075
– Asiatic	0.001	(0.011)	-0.002	(0.010)	0.001	(0.010)	0.015	-0.012	-0.003
– Native American	-0.006	(0.004)	0.010	(0.007)	-0.004	(0.004)	-0.123	0.107	0.017
Schooling < 8 yrs.	-0.003	(0.004)	-0.001	(0.004)	0.004	(0.005)	-0.018	-0.050	0.068
Party affiliation:									
– Democrat	-0.007	(0.022)	0.054	(0.022)	-0.047	(0.021)	-0.125	0.209	-0.084
– Independent	0.006	(0.020)	-0.048	(0.019)	0.042	(0.021)	0.119	-0.196	0.077
– Republican	0.001	(0.020)	-0.007	(0.020)	0.005	(0.020)	0.017	-0.026	0.009
Politics interest:									
– Very low	0.017	(0.014)	-0.009	(0.012)	-0.008	(0.012)	0.087	-0.002	-0.085
– Low	0.011	(0.017)	-0.002	(0.017)	-0.008	(0.016)	0.032	0.018	-0.050
– Medium	-0.013	(0.021)	-0.004	(0.021)	0.016	(0.021)	-0.019	-0.043	0.062
– High	-0.018	(0.019)	0.017	(0.020)	0.001	(0.019)	-0.078	0.036	0.042
– Very high	0.003	(0.015)	-0.002	(0.015)	-0.001	(0.014)	0.015	-0.005	-0.011
Conservative-Liberal score	-0.075	(0.076)	-0.066	(0.075)	0.137	(0.074)	-0.005	-0.119	0.124
Gets news from:									
– Fox News	-0.063	(0.050)	0.005	(0.052)	0.055	(0.050)	-0.060	-0.043	0.104
– CNN	0.062	(0.050)	-0.004	(0.051)	-0.056	(0.050)	0.058	0.045	-0.104
– Breitbart	0.032	(0.032)	-0.029	(0.031)	-0.002	(0.031)	0.085	-0.039	-0.047
– NYT	0.060	(0.046)	-0.015	(0.046)	-0.043	(0.046)	0.072	0.026	-0.099
– MSNBC	0.033	(0.046)	-0.014	(0.045)	-0.018	(0.045)	0.045	0.003	-0.048
– NYPost	0.015	(0.041)	-0.016	(0.041)	0.001	(0.041)	0.033	-0.019	-0.014
Main info. source:									
– Newspapers	0.020	(0.018)	-0.006	(0.017)	-0.014	(0.016)	0.067	0.023	-0.090
– Radio	0.006	(0.010)	-0.014	(0.007)	0.008	(0.010)	0.105	-0.113	0.008
– Socials	0.014	(0.016)	-0.024	(0.014)	0.010	(0.015)	0.113	-0.101	-0.012
– TV	-0.058	(0.022)	0.038	(0.022)	0.019	(0.022)	-0.192	0.038	0.154
Clicks in introduction	-0.033	(0.056)	0.013	(0.072)	0.019	(0.055)	-0.031	-0.004	0.041
Low screen resolution	-0.005	(0.017)	0.019	(0.018)	-0.013	(0.017)	-0.059	0.080	-0.020
Topic of low subjective salience	-0.015	(0.019)	-0.017	(0.018)	0.030	(0.019)	0.005	-0.108	0.103
Topic familiarity:									
– Low	-0.007	(0.005)	0.004	(0.007)	0.002	(0.006)	-0.081	0.015	0.066
– Mid-Low	0.004	(0.013)	0.013	(0.013)	-0.017	(0.011)	-0.032	0.108	-0.077
– Mid-High	-0.002	(0.021)	-0.013	(0.021)	0.015	(0.021)	0.023	-0.058	0.035
– High	0.005	(0.022)	-0.005	(0.022)	-0.000	(0.022)	0.019	-0.009	-0.010
Topic ex ante opinion	0.904	(2.004)	0.662	(1.908)	-1.517	(1.785)	0.005	0.051	-0.056
N of observations:	510		523		532				

Notes: The table presents the means and standard errors for each covariate specified, and the standardized difference between treatment groups for the “Defund Police” news issue to assess balance. Treatment branches are marked in column headers, with “Republican” (“Democrat”) indicating being exposed to news on the issue lead by Republican-leaning (Democrat-leaning) images, and “Neutral” indicating non-partisan leading images.

TABLE (A.2.4)
Balance of observable characteristics across treatment branches, “Iran deal” news issue

Variables:	Republican		Neutral		Democrat		Normalized difference:		
	Mean	St. err.	Mean	St. err.	Mean	St. err.	(R-N)	(N-D)	(D-R)
Age bracket:									
– 18-34	0.010	(0.018)	0.003	(0.018)	-0.013	(0.017)	0.016	0.039	-0.055
– 35-44	0.004	(0.019)	0.001	(0.019)	-0.005	(0.019)	0.006	0.014	-0.020
– 45-54	-0.008	(0.019)	0.012	(0.019)	-0.004	(0.019)	-0.045	0.035	0.010
– 55-65	-0.006	(0.019)	-0.016	(0.019)	0.022	(0.020)	0.023	-0.083	0.060
Ethnicity:									
– Caucasian	-0.020	(0.018)	0.004	(0.017)	0.016	(0.016)	-0.058	-0.033	0.091
– African-American	-0.013	(0.012)	0.008	(0.013)	0.005	(0.013)	-0.070	0.008	0.062
– Latin American	0.016	(0.012)	-0.005	(0.010)	-0.011	(0.010)	0.083	0.030	-0.112
– Asiatic	0.017	(0.012)	-0.017	(0.009)	-0.001	(0.010)	0.143	-0.071	-0.072
– Native American	-0.003	(0.005)	0.001	(0.006)	0.003	(0.006)	-0.030	-0.016	0.046
Schooling < 8 yrs.	-0.009	(0.002)	0.001	(0.005)	0.008	(0.006)	-0.116	-0.063	0.168
Party affiliation:									
– Democrat	-0.003	(0.021)	0.003	(0.021)	0.000	(0.021)	-0.012	0.005	0.008
– Independent	-0.013	(0.020)	-0.003	(0.020)	0.016	(0.020)	-0.022	-0.041	0.063
– Republican	0.016	(0.020)	0.000	(0.020)	-0.017	(0.020)	0.035	0.037	-0.072
Politics interest:									
– Very low	0.008	(0.013)	0.000	(0.012)	-0.008	(0.012)	0.027	0.029	-0.056
– Low	-0.031	(0.015)	0.025	(0.017)	0.006	(0.017)	-0.149	0.047	0.102
– Medium	0.037	(0.021)	-0.039	(0.020)	0.003	(0.021)	0.160	-0.089	-0.071
– High	0.015	(0.019)	-0.006	(0.019)	-0.009	(0.019)	0.049	0.005	-0.054
– Very high	-0.029	(0.013)	0.021	(0.015)	0.008	(0.015)	-0.152	0.038	0.114
Conservative-Liberal score	0.115	(0.074)	0.015	(0.074)	-0.131	(0.074)	0.058	0.086	-0.144
Gets news from:									
– Fox News	0.002	(0.050)	-0.006	(0.050)	0.005	(0.051)	0.007	-0.010	0.003
– CNN	-0.019	(0.050)	-0.003	(0.049)	0.022	(0.049)	-0.014	-0.022	0.036
– Breitbart	-0.006	(0.031)	0.036	(0.032)	-0.030	(0.029)	-0.057	0.092	-0.034
– NYT	-0.031	(0.045)	-0.014	(0.045)	0.045	(0.046)	-0.016	-0.056	0.071
– MSNBC	-0.017	(0.045)	-0.006	(0.046)	0.023	(0.044)	-0.011	-0.028	0.039
– NYPost	-0.029	(0.039)	-0.009	(0.040)	0.039	(0.041)	-0.021	-0.051	0.073
Main info. source:									
– Newspapers	-0.002	(0.017)	0.013	(0.017)	-0.011	(0.016)	-0.037	0.062	-0.026
– Radio	-0.002	(0.009)	-0.000	(0.009)	0.002	(0.009)	-0.008	-0.012	0.020
– Socials	0.002	(0.015)	-0.011	(0.014)	0.009	(0.015)	0.040	-0.060	0.020
– TV	0.010	(0.022)	-0.023	(0.022)	0.013	(0.022)	0.067	-0.074	0.006
Clicks in introduction	-0.101	(0.048)	0.085	(0.074)	0.016	(0.058)	-0.130	0.046	0.095
Low screen resolution	0.005	(0.017)	-0.005	(0.017)	-0.000	(0.017)	0.025	-0.011	-0.014
Topic of low subjective salience	-0.016	(0.018)	-0.009	(0.018)	0.026	(0.019)	-0.016	-0.081	0.096
Topic familiarity:									
– Low	0.009	(0.019)	-0.013	(0.019)	0.004	(0.019)	0.049	-0.038	-0.011
– Mid-Low	-0.011	(0.020)	0.021	(0.021)	-0.010	(0.020)	-0.069	0.067	0.002
– Mid-High	0.014	(0.020)	-0.026	(0.019)	0.012	(0.020)	0.090	-0.083	-0.006
– High	-0.012	(0.014)	0.017	(0.015)	-0.005	(0.014)	-0.088	0.067	0.021
Topic ex ante opinion	0.338	(1.189)	-0.202	(1.203)	-0.137	(1.192)	0.020	-0.002	-0.017
N of observations:	534		536		529				

Notes: The table presents the means and standard errors for each covariate specified, and the standardized difference between treatment groups for the “Iran deal” news issue to assess balance. Treatment branches are marked in column headers, with “Republican” (“Democrat”) indicating being exposed to news on the issue lead by Republican-leaning (Democrat-leaning) images, and “Neutral” indicating non-partisan leading images.

TABLE (A.2.5)
Balance of observable characteristics across treatment branches, “Inflation” news issue

Variables:	Republican		Neutral		Democrat		Normalized difference:		
	Mean	St. err.	Mean	St. err.	Mean	St. err.	(R-N)	(N-D)	(D-R)
Age bracket:									
– 18-34	-0.025	(0.017)	0.011	(0.018)	0.015	(0.018)	-0.089	-0.011	0.099
– 35-44	-0.003	(0.019)	-0.004	(0.019)	0.007	(0.019)	0.001	-0.024	0.023
– 45-54	0.021	(0.019)	-0.005	(0.019)	-0.017	(0.018)	0.059	0.028	-0.088
– 55-65	0.007	(0.019)	-0.002	(0.019)	-0.005	(0.019)	0.021	0.006	-0.027
Ethnicity:									
– Caucasic	-0.019	(0.017)	-0.003	(0.017)	0.022	(0.016)	-0.040	-0.063	0.103
– African-American	0.016	(0.013)	-0.007	(0.012)	-0.010	(0.012)	0.078	0.010	-0.088
– Latin American	0.013	(0.011)	0.001	(0.010)	-0.014	(0.009)	0.048	0.063	-0.111
– Asiatic	0.007	(0.011)	0.006	(0.011)	-0.012	(0.009)	0.004	0.079	-0.083
– Native American	0.005	(0.006)	-0.001	(0.005)	-0.004	(0.004)	0.043	0.034	-0.077
Schooling < 8 yrs.	0.004	(0.005)	-0.005	(0.003)	0.001	(0.005)	0.091	-0.062	-0.030
Party affiliation:									
– Democrat	0.013	(0.021)	-0.013	(0.021)	0.000	(0.021)	0.054	-0.028	-0.026
– Independent	-0.010	(0.020)	0.025	(0.020)	-0.015	(0.020)	-0.076	0.085	-0.009
– Republican	-0.003	(0.020)	-0.012	(0.020)	0.014	(0.020)	0.020	-0.056	0.036
Politics interest:									
– Very low	-0.009	(0.012)	0.009	(0.013)	0.000	(0.012)	-0.062	0.028	0.034
– Low	-0.007	(0.016)	0.001	(0.016)	0.006	(0.016)	-0.021	-0.015	0.036
– Medium	-0.015	(0.020)	0.041	(0.021)	-0.026	(0.020)	-0.117	0.140	-0.023
– High	0.020	(0.019)	-0.038	(0.018)	0.018	(0.019)	0.133	-0.127	-0.006
– Very high	0.011	(0.015)	-0.013	(0.014)	0.002	(0.015)	0.071	-0.043	-0.028
Conservative-Liberal score	0.001	(0.074)	-0.010	(0.073)	0.009	(0.074)	0.006	-0.011	0.005
Gets news from:									
– Fox News	-0.010	(0.050)	-0.015	(0.049)	0.025	(0.050)	0.005	-0.036	0.030
– CNN	0.014	(0.049)	-0.025	(0.049)	0.011	(0.050)	0.034	-0.031	-0.003
– Breitbart	0.015	(0.032)	-0.018	(0.029)	0.003	(0.030)	0.046	-0.030	-0.017
– NYT	0.011	(0.046)	-0.028	(0.046)	0.017	(0.045)	0.037	-0.043	0.006
– MSNBC	0.047	(0.046)	-0.084	(0.044)	0.037	(0.045)	0.125	-0.118	-0.009
– NYPost	0.002	(0.041)	-0.014	(0.040)	0.012	(0.040)	0.017	-0.028	0.010
Main info. source:									
– Newspapers	0.013	(0.017)	0.006	(0.017)	-0.020	(0.016)	0.017	0.069	-0.086
– Radio	0.007	(0.009)	-0.001	(0.009)	-0.006	(0.008)	0.041	0.026	-0.067
– Socials	0.003	(0.015)	0.009	(0.015)	-0.012	(0.014)	-0.016	0.061	-0.045
– TV	0.007	(0.021)	-0.039	(0.022)	0.032	(0.022)	0.090	-0.143	0.052
Clicks in introduction	-0.076	(0.049)	0.080	(0.075)	-0.004	(0.054)	-0.107	0.056	0.060
Low screen resolution	0.013	(0.017)	-0.010	(0.017)	-0.003	(0.017)	0.059	-0.020	-0.039
Topic of low subjective salience	0.009	(0.019)	0.014	(0.019)	-0.023	(0.018)	-0.012	0.086	-0.074
Topic familiarity:									
– Low	-0.035	(0.015)	0.043	(0.018)	-0.009	(0.016)	-0.204	0.132	0.072
– Mid-Low	0.018	(0.020)	-0.030	(0.019)	0.011	(0.020)	0.107	-0.090	-0.016
– Mid-High	0.002	(0.021)	0.001	(0.021)	-0.002	(0.021)	0.002	0.007	-0.008
– High	0.014	(0.016)	-0.015	(0.015)	0.000	(0.016)	0.079	-0.041	-0.038
Topic ex ante opinion	-0.400	(1.230)	-1.787	(1.234)	2.217	(1.241)	0.048	-0.140	0.091
N of observations:	545		538		532				

Notes: The table presents the means and standard errors for each covariate specified, and the standardized difference between treatment groups for the “Inflation” news issue to assess balance. Treatment branches are marked in column headers, with “Republican” (“Democrat”) indicating being exposed to news on the issue lead by Republican-leaning (Democrat-leaning) images, and “Neutral” indicating non-partisan leading images.

TABLE (A.2.6)
Balance of observable characteristics across treatment branches, “Covid measures” news issue

Variables:	Republican		Neutral		Democrat		Normalized difference:		
	Mean	St. err.	Mean	St. err.	Mean	St. err.	(R-N)	(N-D)	(D-R)
Age bracket:									
– 18-34	0.013	(0.018)	0.000	(0.018)	-0.013	(0.017)	0.030	0.031	-0.062
– 35-44	0.026	(0.020)	-0.032	(0.018)	0.006	(0.019)	0.132	-0.088	-0.044
– 45-54	-0.011	(0.018)	0.017	(0.019)	-0.005	(0.019)	-0.065	0.049	0.016
– 55-65	-0.027	(0.019)	0.015	(0.020)	0.012	(0.020)	-0.094	0.008	0.086
Ethnicity:									
– Caucasic	-0.010	(0.017)	-0.006	(0.017)	0.016	(0.016)	-0.009	-0.056	0.065
– African-American	0.007	(0.013)	0.005	(0.013)	-0.012	(0.012)	0.006	0.061	-0.066
– Latin American	-0.014	(0.009)	0.016	(0.012)	-0.001	(0.010)	-0.126	0.068	0.059
– Asiatic	0.002	(0.010)	0.003	(0.011)	-0.004	(0.010)	-0.005	0.030	-0.025
– Native American	0.002	(0.006)	-0.003	(0.005)	0.000	(0.006)	0.043	-0.028	-0.015
Schooling < 8 yrs.	-0.009	(0.002)	0.003	(0.005)	0.006	(0.006)	-0.133	-0.028	0.156
Party affiliation:									
– Democrat	0.005	(0.021)	-0.018	(0.021)	0.013	(0.021)	0.046	-0.062	0.016
– Independent	0.022	(0.020)	-0.006	(0.020)	-0.017	(0.020)	0.060	0.024	-0.084
– Republican	-0.027	(0.019)	0.023	(0.021)	0.004	(0.020)	-0.109	0.041	0.068
Politics interest:									
– Very low	0.021	(0.014)	-0.006	(0.012)	-0.015	(0.011)	0.089	0.035	-0.124
– Low	-0.021	(0.015)	0.015	(0.017)	0.007	(0.016)	-0.097	0.021	0.075
– Medium	0.006	(0.021)	-0.011	(0.021)	0.005	(0.021)	0.037	-0.034	-0.003
– High	-0.008	(0.019)	-0.007	(0.019)	0.015	(0.019)	-0.004	-0.049	0.053
– Very high	0.002	(0.015)	0.009	(0.015)	-0.011	(0.014)	-0.019	0.060	-0.040
Conservative-Liberal score	0.001	(0.074)	0.014	(0.075)	-0.014	(0.073)	-0.008	0.017	-0.009
Gets news from:									
– Fox News	0.004	(0.050)	-0.051	(0.050)	0.045	(0.050)	0.048	-0.084	0.036
– CNN	0.036	(0.049)	-0.050	(0.050)	0.013	(0.050)	0.076	-0.054	-0.021
– Breitbart	0.013	(0.032)	-0.007	(0.031)	-0.007	(0.029)	0.028	0.000	-0.028
– NYT	-0.001	(0.045)	-0.019	(0.047)	0.019	(0.046)	0.017	-0.036	0.020
– MSNBC	0.019	(0.045)	-0.007	(0.047)	-0.011	(0.045)	0.025	0.004	-0.029
– NYPost	0.028	(0.041)	-0.017	(0.040)	-0.011	(0.040)	0.048	-0.007	-0.042
Main info. source:									
– Newspapers	-0.009	(0.016)	0.012	(0.017)	-0.002	(0.017)	-0.054	0.036	0.018
– Radio	0.002	(0.009)	-0.001	(0.009)	-0.002	(0.008)	0.015	0.005	-0.020
– Socials	0.006	(0.015)	0.003	(0.015)	-0.009	(0.014)	0.008	0.038	-0.045
– TV	0.003	(0.022)	0.001	(0.022)	-0.004	(0.022)	0.005	0.011	-0.015
Clicks in introduction	-0.034	(0.051)	-0.030	(0.055)	0.063	(0.077)	-0.003	-0.061	0.065
Low screen resolution	-0.023	(0.016)	0.013	(0.018)	0.010	(0.017)	-0.093	0.008	0.086
Topic of low subjective salience	0.023	(0.019)	-0.009	(0.018)	-0.014	(0.018)	0.073	0.013	-0.086
Topic familiarity:									
– Low	0.014	(0.011)	-0.006	(0.010)	-0.009	(0.009)	0.083	0.014	-0.097
– Mid-Low	-0.005	(0.014)	0.000	(0.015)	0.004	(0.015)	-0.014	-0.013	0.026
– Mid-High	0.001	(0.021)	-0.018	(0.021)	0.017	(0.021)	0.040	-0.071	0.032
– High	-0.011	(0.022)	0.023	(0.022)	-0.012	(0.021)	-0.069	0.072	-0.003
Topic ex ante opinion	0.945	(1.481)	-1.146	(1.478)	0.176	(1.447)	0.062	-0.039	-0.023
N of observations:	531		520		533				

Notes: The table presents the means and standard errors for each covariate specified, and the standardized difference between treatment groups for the “Covid measures” news issue to assess balance. Treatment branches are marked in column headers, with “Republican” (“Democrat”) indicating being exposed to news on the issue lead by Republican-leaning (Democrat-leaning) images, and “Neutral” indicating non-partisan leading images.

TABLE (A.2.7)
Balance of observable characteristics across treatment branches, “Juneteenth” news issue

Variables:	Republican		Neutral		Democrat		Normalized difference:		
	Mean	St. err.	Mean	St. err.	Mean	St. err.	(R-N)	(N-D)	(D-R)
Age bracket:									
– 18-34	-0.005	(0.018)	-0.011	(0.018)	0.015	(0.019)	0.016	-0.064	0.048
– 35-44	0.019	(0.020)	-0.001	(0.020)	-0.018	(0.019)	0.045	0.041	-0.085
– 45-54	-0.018	(0.018)	0.016	(0.020)	0.002	(0.019)	-0.079	0.032	0.047
– 55-65	0.003	(0.020)	-0.004	(0.020)	0.001	(0.020)	0.017	-0.011	-0.006
Ethnicity:									
– Caucasian	0.000	(0.017)	0.012	(0.017)	-0.012	(0.018)	-0.029	0.059	-0.031
– African-American	-0.008	(0.013)	-0.016	(0.012)	0.023	(0.014)	0.028	-0.128	0.101
– Latin American	0.001	(0.011)	-0.006	(0.010)	0.005	(0.011)	0.031	-0.048	0.017
– Asiatic	0.002	(0.010)	-0.003	(0.010)	0.001	(0.010)	0.023	-0.016	-0.007
– Native American	0.002	(0.006)	-0.001	(0.005)	-0.001	(0.005)	0.026	0.005	-0.031
Schooling < 8 yrs.									
	0.004	(0.005)	0.001	(0.005)	-0.005	(0.003)	0.029	0.066	-0.094
Party affiliation:									
– Democrat	-0.005	(0.021)	0.026	(0.022)	-0.021	(0.021)	-0.063	0.096	-0.032
– Independent	-0.005	(0.020)	-0.016	(0.020)	0.021	(0.020)	0.025	-0.081	0.057
– Republican	0.010	(0.020)	-0.010	(0.020)	-0.000	(0.020)	0.044	-0.021	-0.022
Politics interest:									
– Very low	-0.009	(0.012)	0.021	(0.014)	-0.011	(0.012)	-0.100	0.106	-0.006
– Low	-0.026	(0.016)	0.002	(0.017)	0.024	(0.018)	-0.077	-0.056	0.133
– Medium	-0.012	(0.021)	-0.001	(0.021)	0.013	(0.021)	-0.022	-0.029	0.051
– High	0.018	(0.020)	-0.011	(0.019)	-0.007	(0.019)	0.068	-0.009	-0.059
– Very high	0.029	(0.016)	-0.010	(0.015)	-0.019	(0.014)	0.112	0.026	-0.138
Conservative-Liberal score									
	0.054	(0.077)	-0.129	(0.074)	0.070	(0.074)	0.108	-0.119	0.009
Gets news from:									
– Fox News	0.004	(0.052)	-0.096	(0.050)	0.089	(0.050)	0.086	-0.163	0.073
– CNN	-0.015	(0.051)	0.001	(0.051)	0.013	(0.050)	-0.014	-0.011	0.024
– Breitbart	0.066	(0.034)	-0.079	(0.027)	0.010	(0.033)	0.206	-0.129	-0.073
– NYT	0.031	(0.047)	-0.007	(0.048)	-0.024	(0.046)	0.035	0.016	-0.052
– MSNBC	0.009	(0.046)	-0.015	(0.046)	0.005	(0.045)	0.024	-0.020	-0.004
– NYPost	0.079	(0.043)	-0.040	(0.040)	-0.041	(0.040)	0.125	0.002	-0.126
Main info. source:									
– Newspapers	0.031	(0.018)	0.018	(0.018)	-0.049	(0.014)	0.033	0.185	-0.218
– Radio	-0.001	(0.009)	0.006	(0.010)	-0.005	(0.008)	-0.038	0.057	-0.019
– Socials	0.006	(0.016)	-0.007	(0.015)	0.001	(0.015)	0.039	-0.024	-0.015
– TV	-0.012	(0.022)	-0.010	(0.022)	0.022	(0.022)	-0.005	-0.065	0.069
Clicks in introduction									
	0.020	(0.063)	-0.059	(0.049)	0.036	(0.071)	0.061	-0.069	0.010
Low screen resolution									
	0.007	(0.017)	0.007	(0.018)	-0.014	(0.017)	-0.002	0.054	-0.052
Topic of low subjective salience									
	-0.011	(0.019)	0.004	(0.020)	0.007	(0.020)	-0.034	-0.007	0.042
Topic familiarity:									
– Low	-0.008	(0.010)	-0.005	(0.011)	0.013	(0.012)	-0.011	-0.074	0.085
– Mid-Low	-0.025	(0.017)	0.011	(0.018)	0.014	(0.018)	-0.090	-0.008	0.099
– Mid-High	0.015	(0.022)	0.006	(0.022)	-0.020	(0.021)	0.018	0.053	-0.071
– High	0.018	(0.021)	-0.011	(0.021)	-0.008	(0.020)	0.063	-0.008	-0.055
Topic ex ante opinion									
	0.222	(1.646)	1.677	(1.668)	-1.835	(1.632)	-0.039	0.094	-0.055
N of observations:		522		500		520			

Notes: The table presents the means and standard errors for each covariate specified, and the standardized difference between treatment groups for the “Juneteenth” news issue to assess balance. Treatment branches are marked in column headers, with “Republican” (“Democrat”) indicating being exposed to news on the issue lead by Republican-leaning (Democrat-leaning) images, and “Neutral” indicating non-partisan leading images.

TABLE (A.2.8)
Heterogeneous Impact of Leading Images on News-Readers' Opinion:
(by Readers' Political Party Affiliation)

	(1) Defund Police	(2) Iran deal	(3) Inflation	(4) Covid measures	(5) Juneteenth
Dependent variable:					
Post-treatment opinion on topic					
Democrats x Dem-leaning images (D)	3.880 (1.172)	0.984 (1.936)	2.466 (1.735)	-0.341 (2.134)	1.242 (2.878)
Democrats x neutral images (N)	2.197 (1.360)	1.939 (2.006)	2.236 (0.962)	-2.073 (0.386)	0.680 (0.537)
Democrats x Rep-leaning images (R)	2.449 (1.113)	-0.555 (1.621)	-0.103 (1.106)	-2.390 (1.349)	1.116 (0.419)
Independents x Dem-leaning images (D)	0.632 (1.107)	-0.923 (2.505)	-1.865 (1.055)	-0.184 (1.031)	-1.017 (0.229)
Independents x Rep-leaning images (R)	-0.835 (1.493)	-3.764 (1.913)	-1.626 (0.880)	-1.895 (0.769)	-0.781 (1.042)
Republicans x Dem-leaning images (D)	1.249 (2.303)	-3.456 (1.964)	-1.685 (0.673)	-4.874 (0.897)	-1.983 (0.848)
Republicans x neutral images (N)	-1.279 (2.621)	-4.721 (2.115)	-1.277 (0.842)	-5.555 (0.835)	0.022 (1.018)
Republicans x Rep-leaning images (R)	-5.956 (1.369)	-4.939 (1.976)	-5.571 (1.956)	-5.577 (1.454)	-1.155 (1.602)
H0 for equality tests:	P value:	P value:	P value:	P value:	P value:
Dem*(D) - Rep*(R) \leq Dem*(R) - Rep*(D):	0.002	0.030	0.027	0.086	0.564
Dem*(R) = Rep*(D):	0.424	0.091	0.118	0.275	0.010
Dem*(D) - Rep*(R) \leq Dem*(N) - Rep*(N):	0.027	0.613	0.062	0.290	0.315
Observations	1574	1608	1625	1595	1551
Treatment-independent controls	Y	Y	Y	Y	Y

Notes: The Table presents OLS estimates of the effect of the Democrat-leaning (D), neutral (N) and Republican-leaning (R) news-leading images. The dependent variable is respondents' opinion after exposure to the news (column headers indicate the relevant news issue). Treatments are interacted with indicators of the respondent's political affiliation (Democratic, Independent, or Republican), which is registered before treatment. The dependent variable for the "Defund Police" issue ranges between -100 and 100, while all others range in -50+50. Variables are adjusted so that the highest value in the range always corresponds to the Democrats' ideological position (hence the largest of any two coefficients indicates a relatively more pro-Democratic opinion, and vice versa). Treatment-independent controls are the same as in the main specification (with the natural exception of controls for political opinion and party preference). The panel below the regression coefficients reports the P-values for one-sided and two-sided tests of equality between coefficients (null hypotheses are indicated on the left) using robust standard errors. Heteroskedasticity-robust standard errors in parentheses.

TABLE (A.2.9)
Heterogeneous Impact of Leading Images on News-Readers' Opinion:
(by tercile of readers' prior opinion on the issue)

	(1) Police funds	(2) Covid measures	(3) Iran deal	(4) Inflation	(5) Juneteenth
Dependent variable:					
Opinion difference					
Lowest prior opinion x Dem-leaning images (D)	12.554 (3.362)	2.138 (4.662)	0.105 (2.870)	-2.888 (3.720)	-1.965 (3.529)
Lowest prior opinion x neutral images (N)	7.375 (3.375)	0.409 (4.888)	1.544 (2.842)	-3.352 (3.706)	-0.000 (3.380)
Lowest prior opinion x Rep-leaning images (R)	7.854 (3.729)	-3.715 (4.517)	-3.153 (3.000)	-3.577 (3.800)	-0.888 (3.506)
Medium prior opinion x Dem-leaning images (D)	6.624 (2.437)	-2.498 (2.622)	1.108 (1.885)	0.125 (2.395)	1.232 (1.813)
Medium prior opinion x neutral images (N)	7.431 (2.521)	-4.320 (2.692)	-1.083 (1.903)	1.703 (2.499)	0.799 (1.850)
Medium prior opinion x Rep-leaning images (R)	4.628 (2.562)	-4.247 (2.616)	-2.021 (1.946)	-0.662 (2.499)	1.207 (1.713)
Highest prior opinion x Dem-leaning images (D)	4.582 (2.184)	0.288 (2.367)	-0.794 (1.613)	3.552 (2.369)	0.860 (0.924)
Highest prior opinion x neutral images (N)	6.462 (2.366)	-1.969 (2.250)	1.004 (1.519)	5.445 (2.293)	1.481 (0.952)
Constant	-4.536 (17.908)	-23.977 (9.630)	15.283 (9.085)	20.849 (7.958)	17.378 (6.435)
Observations	1436	1491	1510	1505	1414
Treatment-independent controls	Y	Y	Y	Y	Y

Notes: The Table presents the OLS estimates for the effect of the Democrat-leaning (D), neutral (N) and Republican-leaning (R) news-leading images interacted with the terciles of respondents' first opinion on the news issue, i.e. that expressed before the treatment exposure. The dependent variable is respondents' opinion after exposure to the news (column headers indicate the relevant news issue). Treatment-independent controls are the same as in the main specification. Heteroskedasticity-robust standard errors are in parentheses.

TABLE (A.2.10)
 Heterogeneous Impact of Leading Images on News-Readers' Opinion:
 (by level of subjective salience of the issue)

	(1) Police funds	(2) Covid measures	(3) Iran deal	(4) Inflation	(5) Juneteenth
Dependent variable:					
Opinion difference					
Lowest salience x Dem-leaning images (D)	3.707 (2.449)	6.523 (2.437)	-1.574 (1.772)	2.808 (2.303)	-7.221 (2.066)
Lowest salience x neutral images (N)	4.446 (2.898)	5.313 (2.551)	0.314 (1.691)	2.984 (2.299)	-4.130 (1.926)
Lowest salience x Rep-leaning images (R)	0.286 (2.807)	3.319 (2.588)	-2.531 (1.579)	0.674 (2.318)	-5.581 (2.146)
Medium salience x Dem-leaning images (D)	1.781 (2.547)	3.876 (2.451)	3.167 (1.602)	4.352 (2.275)	-2.480 (1.473)
Medium salience x neutral images (N)	0.784 (2.685)	0.711 (2.406)	0.812 (1.677)	3.862 (2.285)	-1.556 (1.695)
Medium salience x Rep-leaning images (R)	0.384 (2.697)	0.682 (2.470)	-0.301 (1.658)	2.668 (2.381)	-4.360 (1.467)
Highest salience x Dem-leaning images (D)	6.260 (2.522)	1.334 (2.557)	1.491 (1.728)	1.248 (2.403)	-0.104 (1.583)
Highest salience x neutral images (N)	3.655 (2.394)	0.187 (2.592)	3.238 (1.783)	4.142 (2.334)	-2.251 (1.427)
Constant	-0.179 (17.686)	-30.384 (8.476)	16.298 (8.858)	14.877 (6.963)	20.526 (5.184)
Observations	1436	1491	1510	1505	1414
Treatment-independent controls	Y	Y	Y	Y	Y

Notes: The Table presents the OLS estimates for the effect of the Democrat-leaning (D), neutral (N) and Republican-leaning (R) news-leading images interacted with indicators for the level of subjective salience assigned by respondents to the news issue (salience is measured before the treatment exposure). The dependent variable is respondents' opinion after exposure to the news (column headers indicate the relevant news issue). All dependent variables are adjusted so that the highest value corresponds to the Democrats' ideological position, hence positive coefficients indicate a pro-Democratic opinion shift. Treatment-independent controls are the same as in the main specification. Heteroskedasticity-robust standard errors are in parentheses.

TABLE (A.2.11)
 Heterogeneous Impact of Leading Images on News-Readers' Opinion:
 (by level of self-reported knowledge of the issue)

	(1) Police funds	(2) Covid measures	(3) Iran deal	(4) Inflation	(5) Juneteenth
Dependent variable:					
Opinion difference					
Lowest knowledge x Dem-leaning images (D)	3.573 (1.842)	5.511 (3.123)	1.508 (1.400)	3.636 (1.692)	0.036 (1.193)
Lowest knowledge x neutral images (N)	3.437 (2.118)	4.474 (3.369)	2.707 (1.403)	2.949 (1.727)	0.844 (1.148)
Lowest knowledge x Rep-leaning images (R)	1.214 (2.145)	3.114 (3.133)	0.216 (1.350)	1.018 (1.830)	-0.172 (1.153)
Highest knowledge x Dem-leaning images (D)	5.080 (1.855)	3.030 (2.090)	2.787 (1.439)	0.813 (1.687)	-0.795 (1.157)
Highest knowledge x neutral images (N)	3.240 (1.803)	0.108 (2.122)	2.239 (1.472)	3.036 (1.574)	-0.279 (1.363)
Constant	-0.409 (17.627)	-31.033 (8.207)	13.792 (8.704)	13.671 (6.948)	16.434 (4.977)
Observations	1436	1491	1510	1505	1414
Treatment-independent controls	Y	Y	Y	Y	Y

Notes: The Table presents the OLS estimates for the effect of the Democrat-leaning (D), neutral (N) and Republican-leaning (R) news-leading images interacted with indicators for two levels of (self-reported) knowledge on the issue prior to the news exposure. The dependent variable is respondents' opinion after treatment exposure (column headers indicate the relevant news issue). All dependent variables are adjusted so that the highest value corresponds to the Democrats' ideological position, hence positive coefficients indicate a pro-Democratic opinion shift. Treatment-independent controls are the same as in the main specification. Heteroskedasticity-robust standard errors are in parentheses.

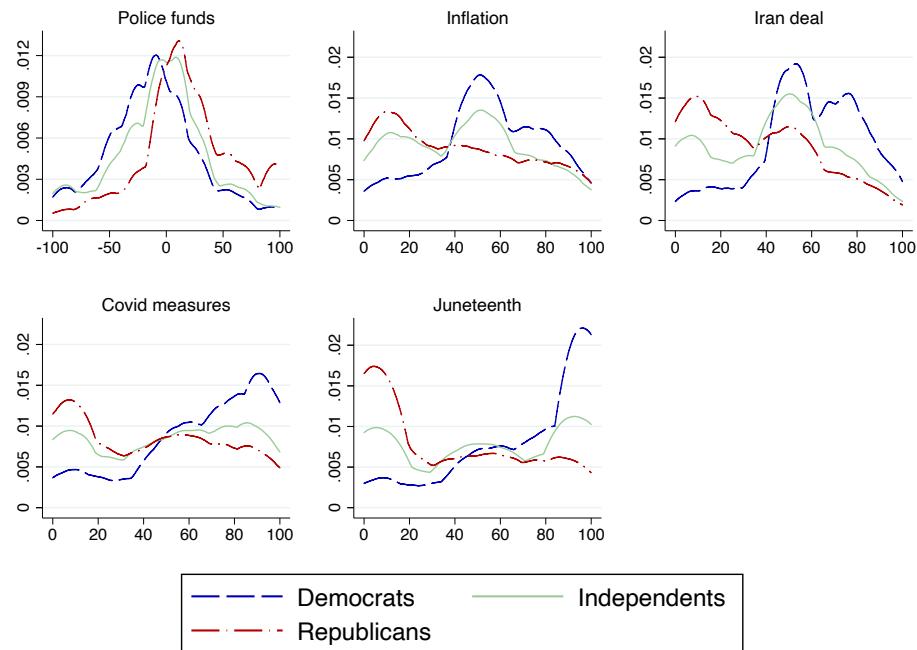


FIGURE (A.2.6)
Densities of first opinion by news issue

Notes: The Figure displays the densities of first opinions (i.e. the opinion expressed before treatment exposure) on the five news issues, dividing respondents by their party affiliation. Parties' modes of opinion are closer in the "Police funds" issue, and most distant in the "Juneteenth" issue, suggesting that the ideological distance between Democrats and Republicans in the sample is smaller in the former and wider in the latter case.

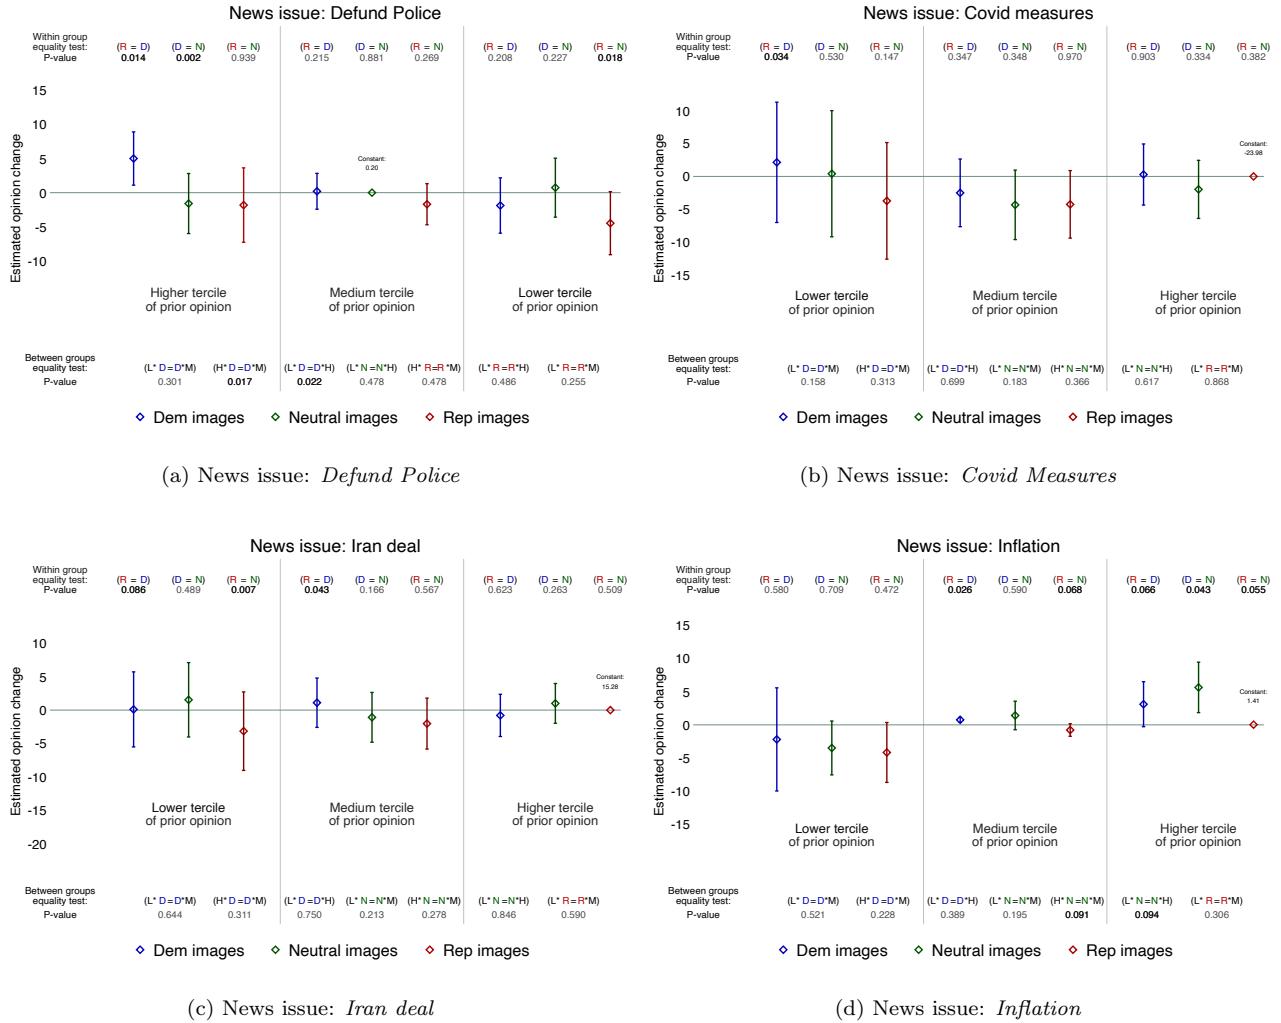


FIGURE (A.2.7)
Heterogeneous effects of images on opinion,
by respondents' tercile of prior opinion on the issue

Notes: The Figure shows OLS estimates of opinion changes after news exposure (news issues indicated below each panel). Treatments are interacted with respondents' tercile of prior opinion on the news issue. Omitted regression category: Respondents in the highest tercile of prior opinion, exposed to Rep-leaning images. Lines indicate 95% CI (heteroskedasticity-robust st. errors). Equality tests on top of each Figure compare coefficients within each party; those at the bottom compare coefficients across parties (tested coefficients indicated in parentheses). All p-values are for two-sided tests of equality, with bold font marking statistical significance at 10 percent level or higher.

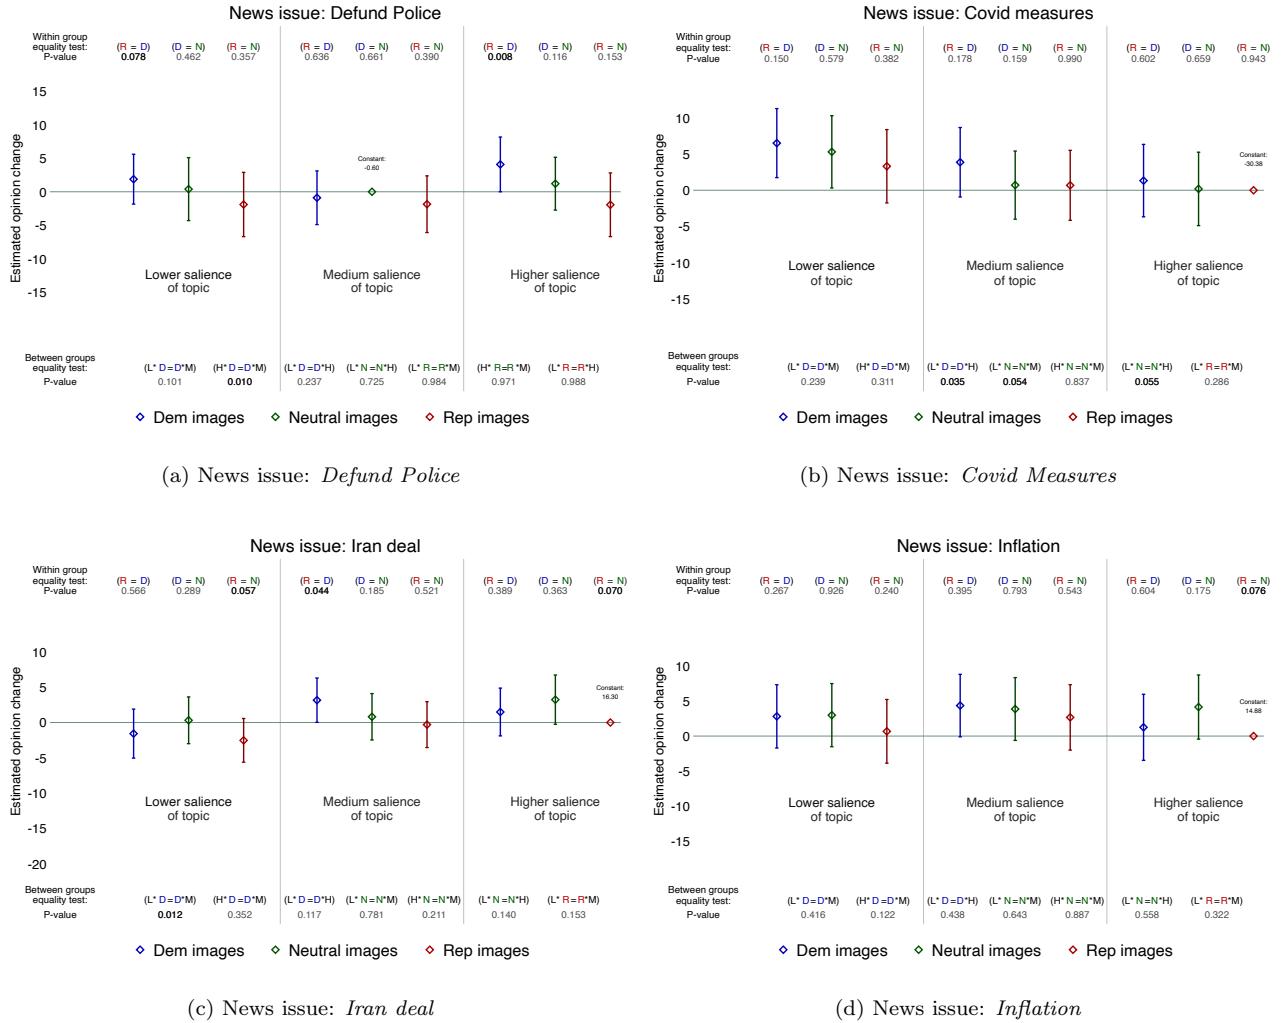


FIGURE (A.2.8)
Heterogeneous effects of images on opinion,
by respondents' subjective salience of the news issue

Notes: The Figure shows OLS estimates of opinion changes after news exposure (news issues indicated below each panel). Treatments are interacted with three indicators for the level of subjective salience respondents assign to the news issue prior to treatment exposure. Omitted regression category: Respondents in the highest level of salience exposed to Rep-leaning images. Lines indicate 95% CI (heteroskedasticity-robust st. errors). Equality tests on top of each Figure compare coefficients within each party; those at the bottom compare coefficients across parties (tested coefficients indicated in parentheses). All p-values are for two-sided tests of equality, with bold font marking statistical significance at 10 percent level or higher.

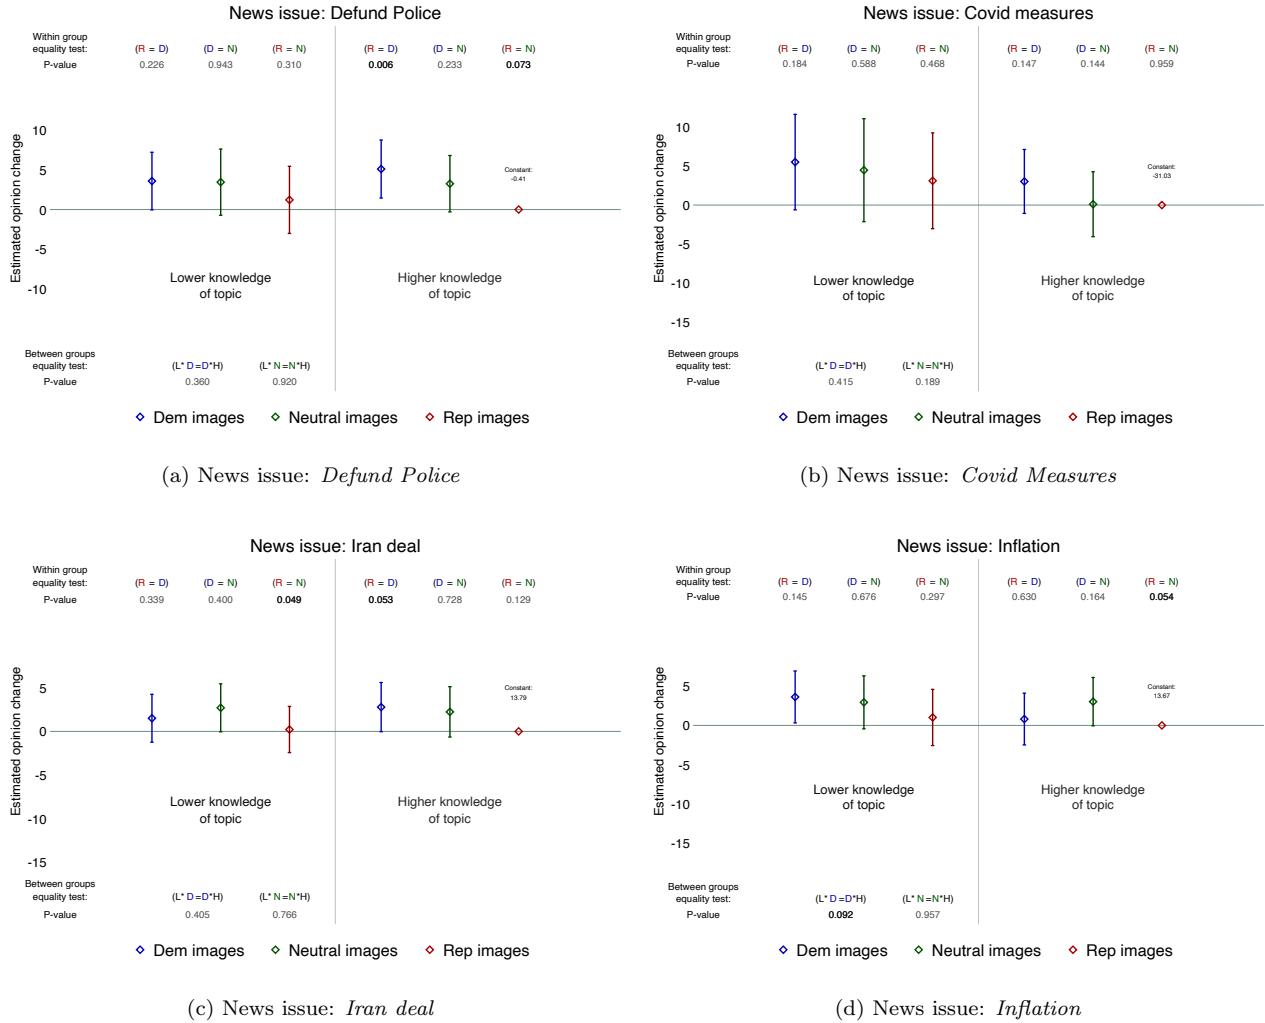


FIGURE (A.2.9)
Heterogeneous effects of images on opinion,
by respondents' knowledge of the news issue

Notes: The Figure shows OLS estimates of opinion changes after news exposure (news issues indicated below each panel). Treatments are interacted with two indicators for respondents' (self reported) level of knowledge on the news issue prior to treatment exposure. Omitted regression category: Respondents in the highest level of salience exposed to Rep-leaning images. Lines indicate 95% CI (heteroskedasticity-robust st. errors). Equality tests on top of each Figure compare coefficients within each party; those at the bottom compare coefficients across parties (tested coefficients indicated in parentheses). All p-values are for two-sided tests of equality, with bold font marking statistical significance at 10 percent level or higher.